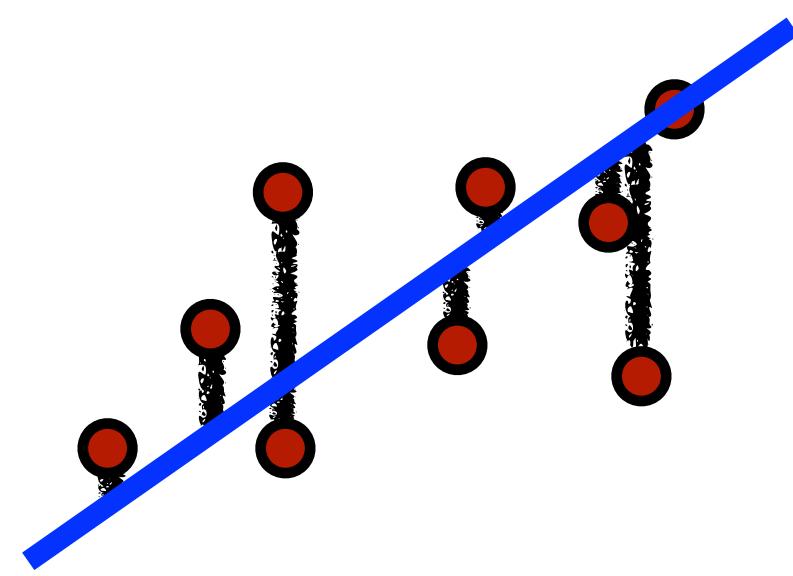


Linear Regression



Regression

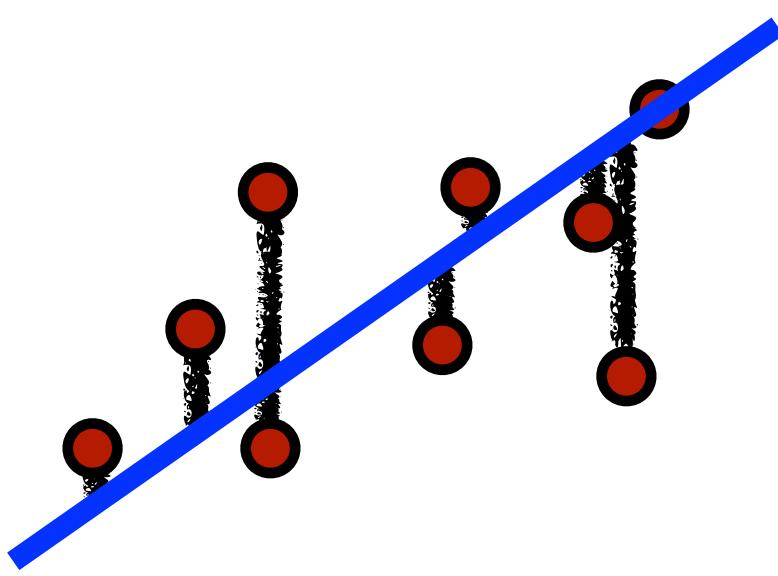


Goal: Learn a mapping from observations (features) to continuous labels given a training set (supervised learning)

Example: Height, Gender, Weight → Shoe Size

- Audio features → Song year
- Processes, memory → Power consumption
- Historical financials → Future stock price
- Many more

Linear Least Squares Regression



Example: Predicting shoe size from height, gender, and weight

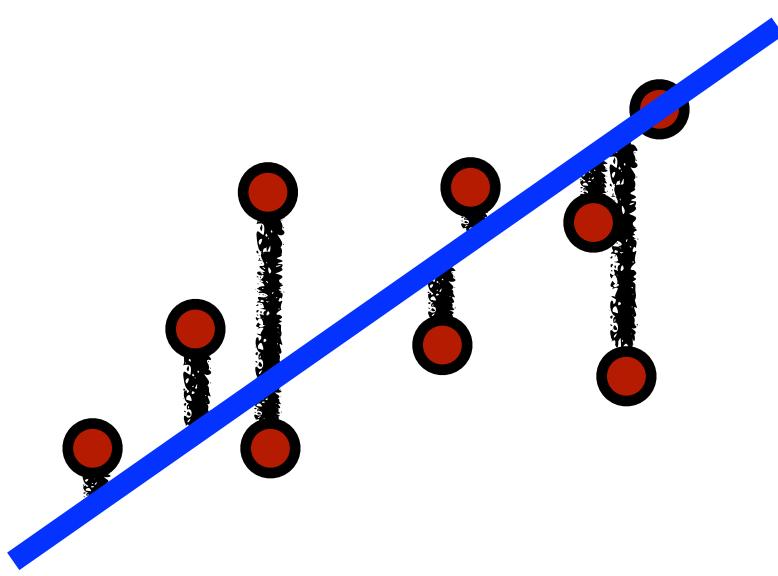
For each observation we have a feature vector, \mathbf{x} , and label, y

$$\mathbf{x}^\top = [x_1 \quad x_2 \quad x_3]$$

We assume a *linear* mapping between features and label:

$$y \approx w_0 + w_1 x_1 + w_2 x_2 + w_3 x_3$$

Linear Least Squares Regression



Example: Predicting shoe size from height, gender, and weight

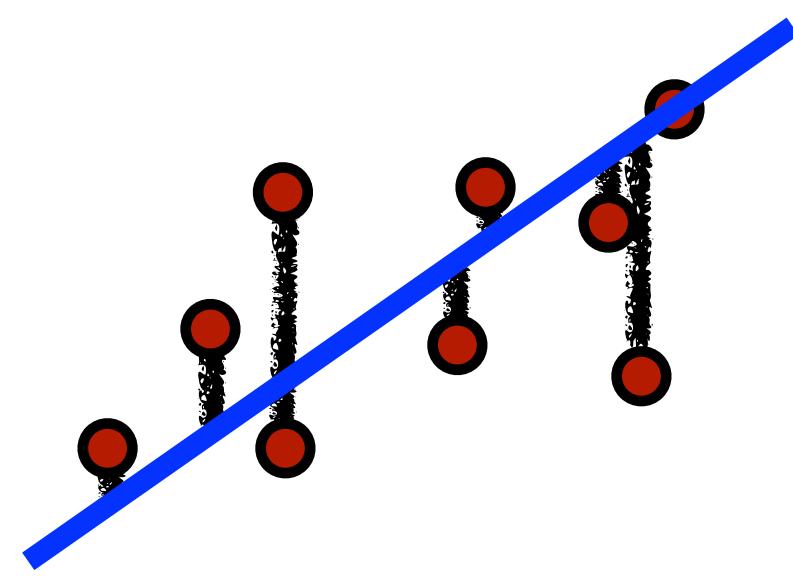
We can augment the feature vector to incorporate offset:

$$\mathbf{x}^\top = [1 \quad x_1 \quad x_2 \quad x_3]$$

We can then rewrite this linear mapping as scalar product:

$$y \approx \hat{y} = \sum_{i=0}^3 w_i x_i = \mathbf{w}^\top \mathbf{x}$$

Why a Linear Mapping?



Simple

Often works well in practice

Can introduce complexity via feature extraction

1D Example

Goal: find the line of best fit

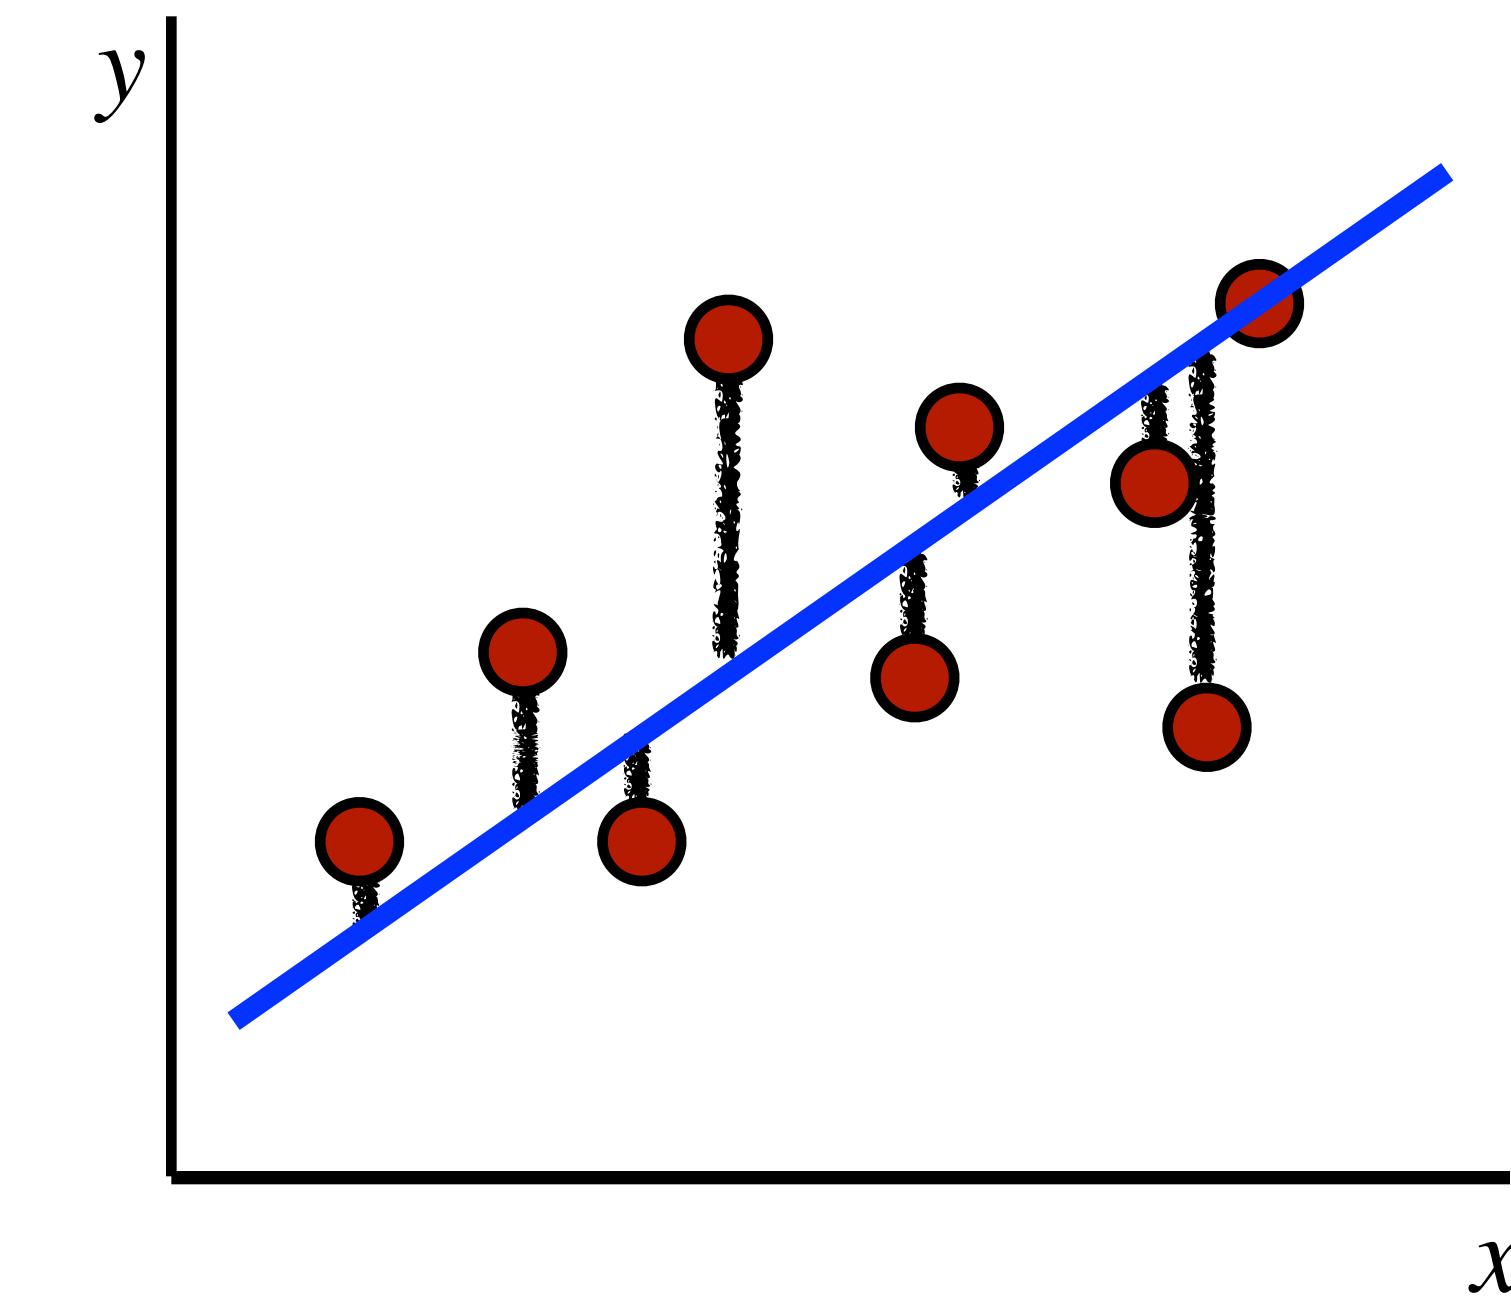
x coordinate: features

y coordinate: labels

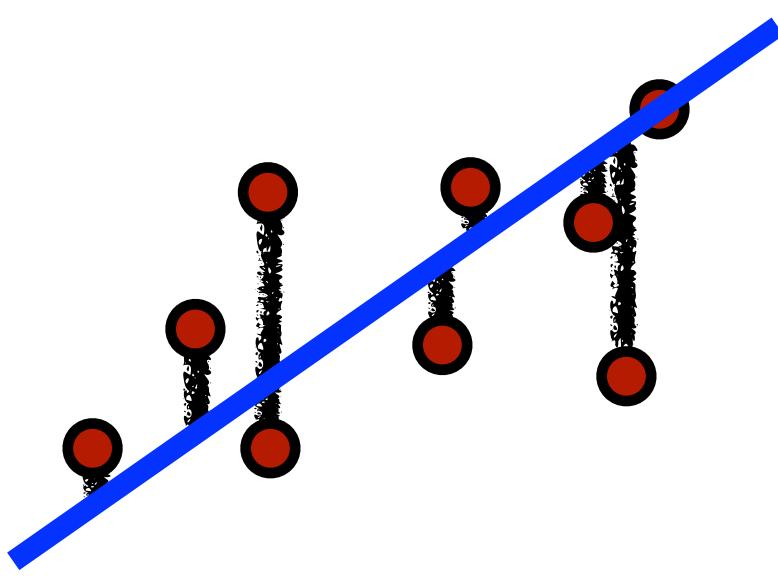
$$y \approx \hat{y} = w_0 + w_1 x$$

Intercept / Offset

Slope



Evaluating Predictions



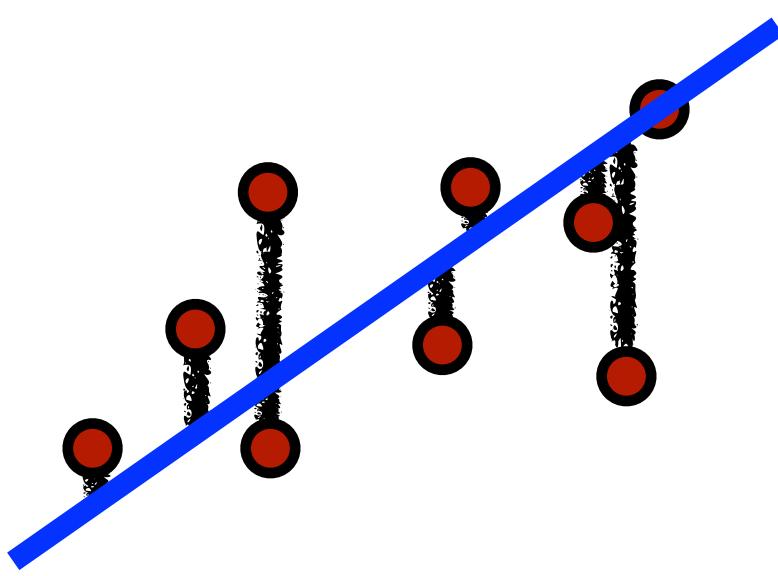
Can measure ‘closeness’ between label and prediction

- Shoe size: better to be off by one size than 5 sizes
- Song year prediction: better to be off by a year than by 20 years

What is an appropriate evaluation metric or ‘loss’ function?

- Absolute loss: $|y - \hat{y}|$
- Squared loss: $(y - \hat{y})^2$ ← Has nice mathematical properties

How Can We Learn Model (\mathbf{w})?



Assume we have n training points, where $\mathbf{x}^{(i)}$ denotes the i th point

Recall two earlier points:

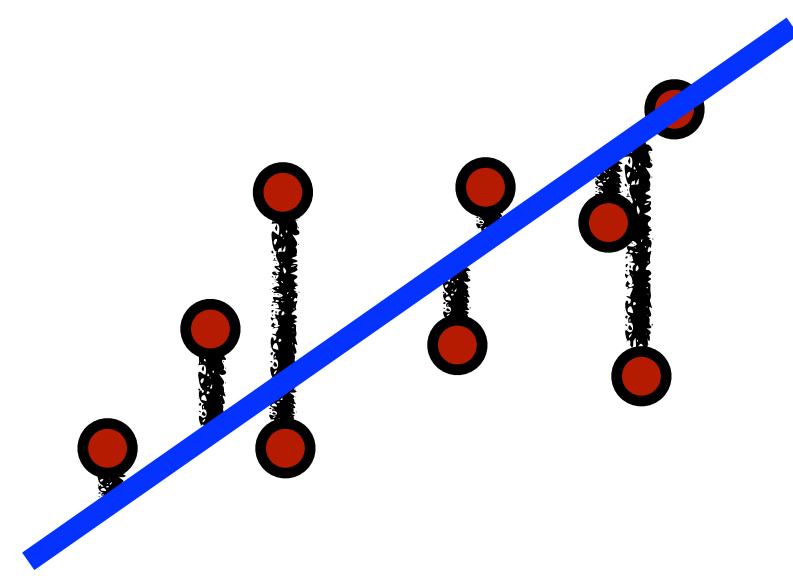
- *Linear assumption:* $\hat{y} = \mathbf{w}^\top \mathbf{x}$
- We use *squared loss*: $(y - \hat{y})^2$

Idea: Find \mathbf{w} that minimizes squared loss over training points:

$$\min_{\mathbf{w}} \sum_{i=1}^n (\underline{\mathbf{w}^\top \mathbf{x}^{(i)}} - \hat{y}^{(i)})^2$$

Given n training points with d features, we define:

- $\mathbf{X} \in \mathbb{R}^{n \times d}$: matrix storing points
- $\mathbf{y} \in \mathbb{R}^n$: real-valued labels
- $\hat{\mathbf{y}} \in \mathbb{R}^n$: predicted labels, where $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$
- $\mathbf{w} \in \mathbb{R}^d$: regression parameters / model to learn



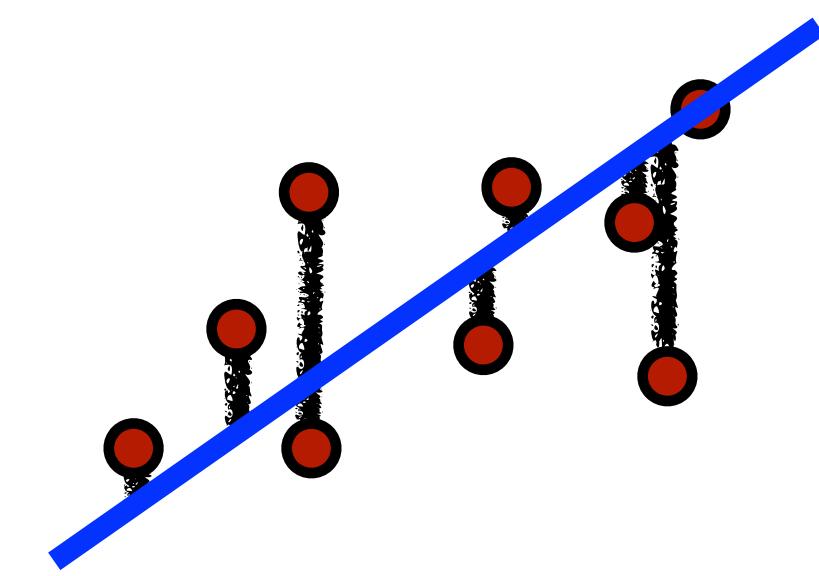
Least Squares Regression: Learn mapping (\mathbf{w}) from features to labels that minimizes residual sum of squares:

$$\min_{\mathbf{w}} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

Equivalent $\min_{\mathbf{w}} \sum_{i=1}^n (\mathbf{w}^\top \mathbf{x}^{(i)} - y^{(i)})^2$ by definition of Euclidean norm

Find solution by setting derivative to zero

$$1D: f(w) = \|\mathbf{w}\mathbf{x} - \mathbf{y}\|_2^2 = \sum_{i=1}^n (w\mathbf{x}^{(i)} - y^{(i)})^2$$



$$\frac{df}{dw}(w) = 2 \underbrace{\sum_{i=1}^n x^{(i)} (w\mathbf{x}^{(i)} - y^{(i)})}_{w\mathbf{x}^\top \mathbf{x} - \mathbf{x}^\top \mathbf{y}} = 0 \iff w\mathbf{x}^\top \mathbf{x} - \mathbf{x}^\top \mathbf{y} = 0$$
$$\iff w = (\mathbf{x}^\top \mathbf{x})^{-1} \mathbf{x}^\top \mathbf{y}$$

Least Squares Regression: Learn mapping (\mathbf{w}) from features to labels that minimizes residual sum of squares:

$$\min_{\mathbf{w}} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

Closed form solution: $\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ (if inverse exists)

Overfitting and Generalization

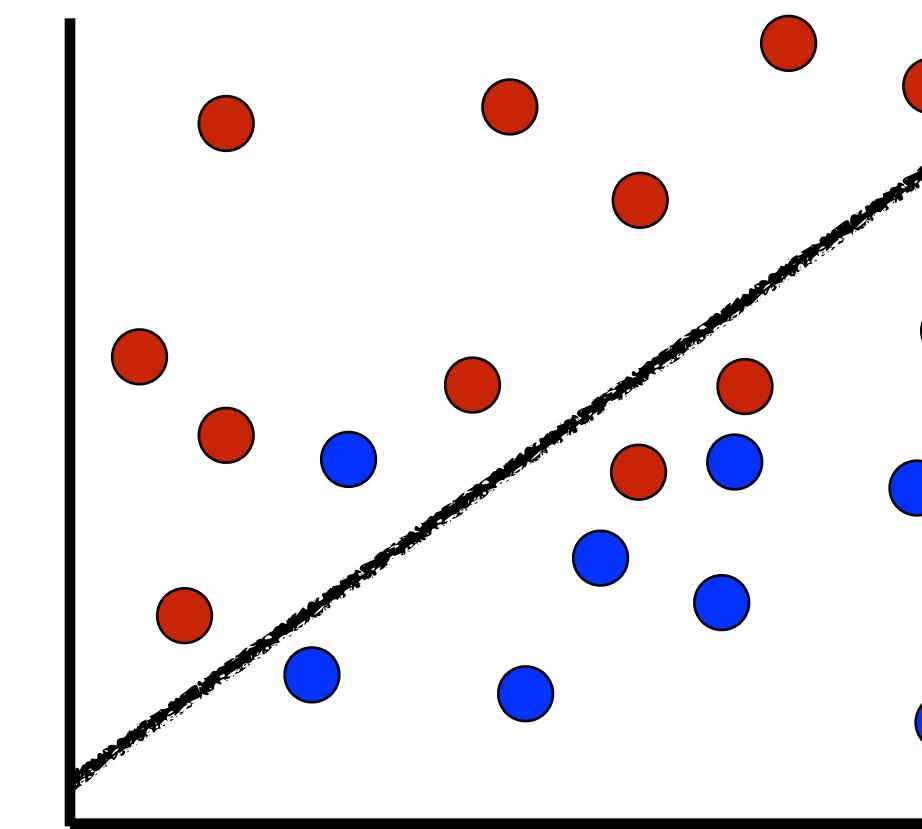
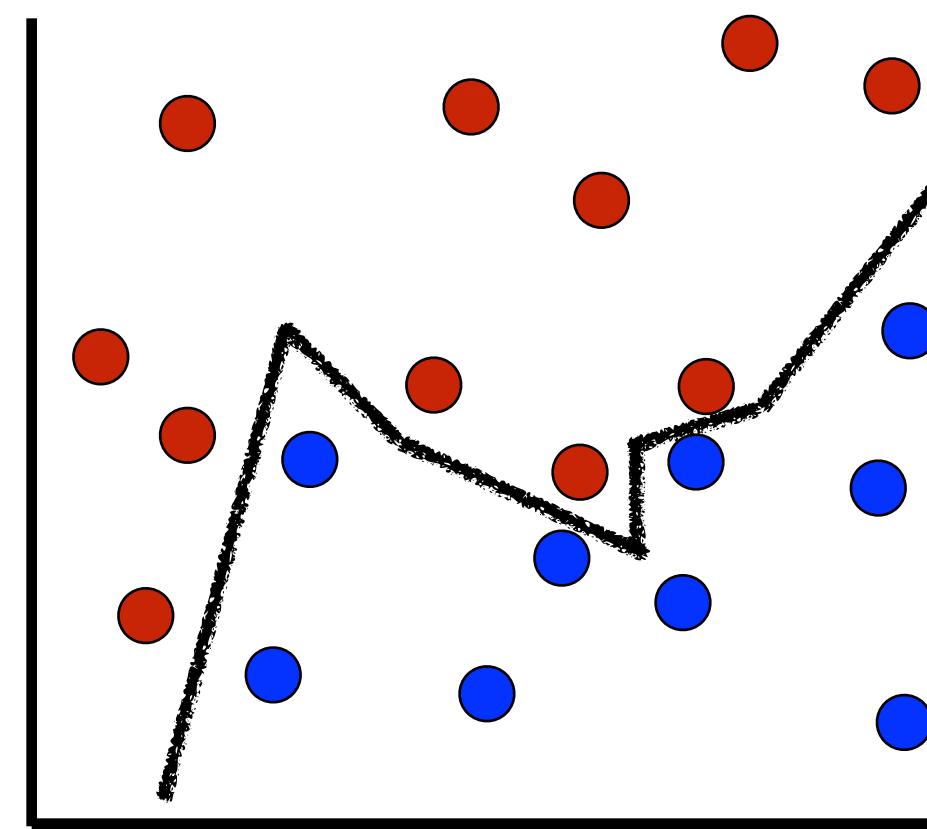
We want good predictions on new data, i.e., 'generalization'

Least squares regression minimizes training error, and could overfit

- Simpler models are more likely to generalize (Occam's razor)

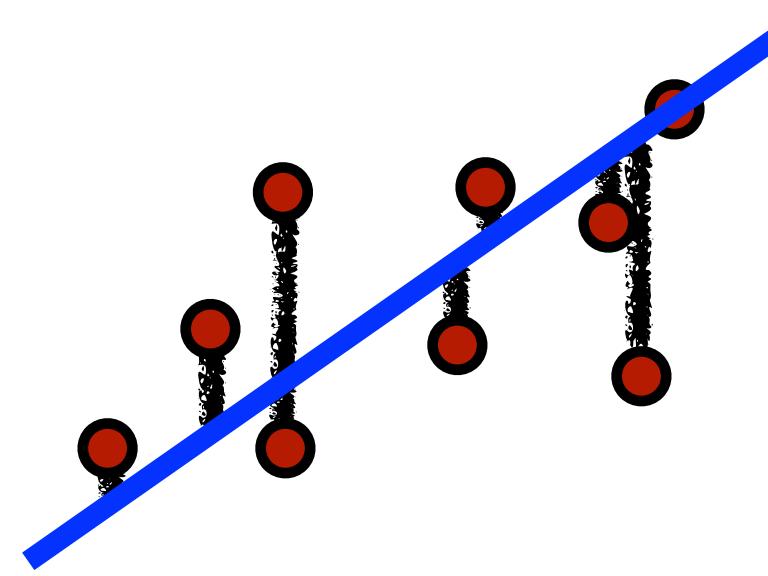
Can we change the problem to penalize for model complexity?

- Intuitively, models with smaller weights are simpler



Given n training points with d features, we define:

- $\mathbf{X} \in \mathbb{R}^{n \times d}$: matrix storing points
- $\mathbf{y} \in \mathbb{R}^n$: real-valued labels
- $\hat{\mathbf{y}} \in \mathbb{R}^n$: predicted labels, where $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$
- $\mathbf{w} \in \mathbb{R}^d$: regression parameters / model to learn



Ridge Regression: Learn mapping (\mathbf{w}) that minimizes residual sum of squares along with a regularization term:

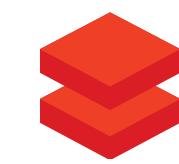
$$\min_{\mathbf{w}} \frac{\text{Training Error}}{||\mathbf{X}\mathbf{w} - \mathbf{y}||_2^2} + \frac{\text{Model Complexity}}{\lambda ||\mathbf{w}||_2^2}$$

Closed-form solution: $\mathbf{w} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_d)^{-1} \mathbf{X}^\top \mathbf{y}$

free parameter trades off
between training error and
model complexity

Millionsong Regression Pipeline

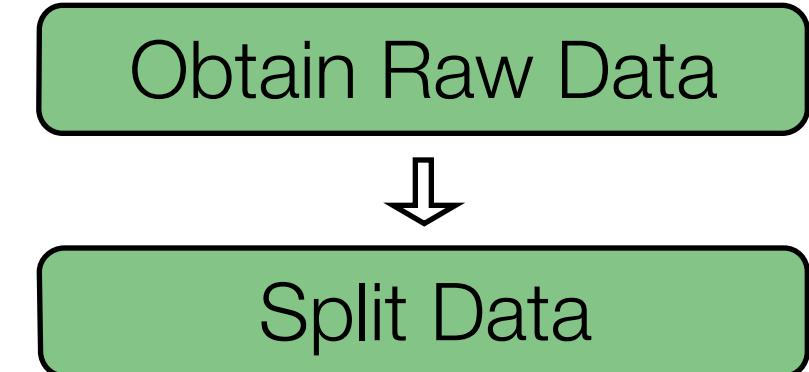
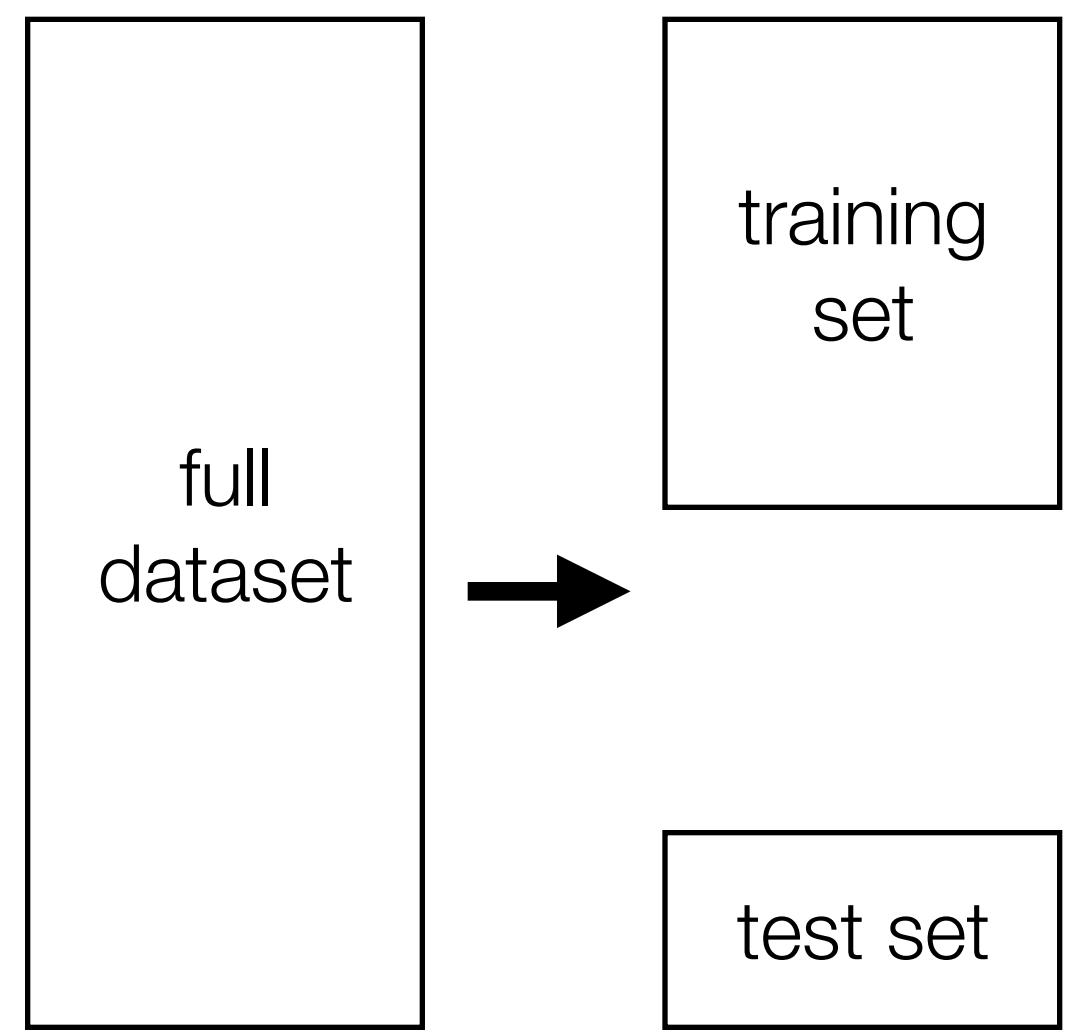


UCLA 
 databricks™

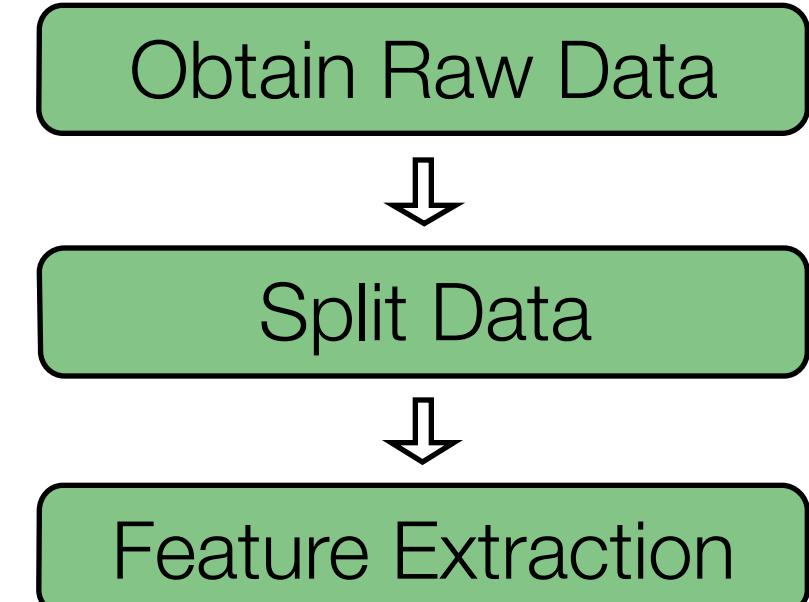
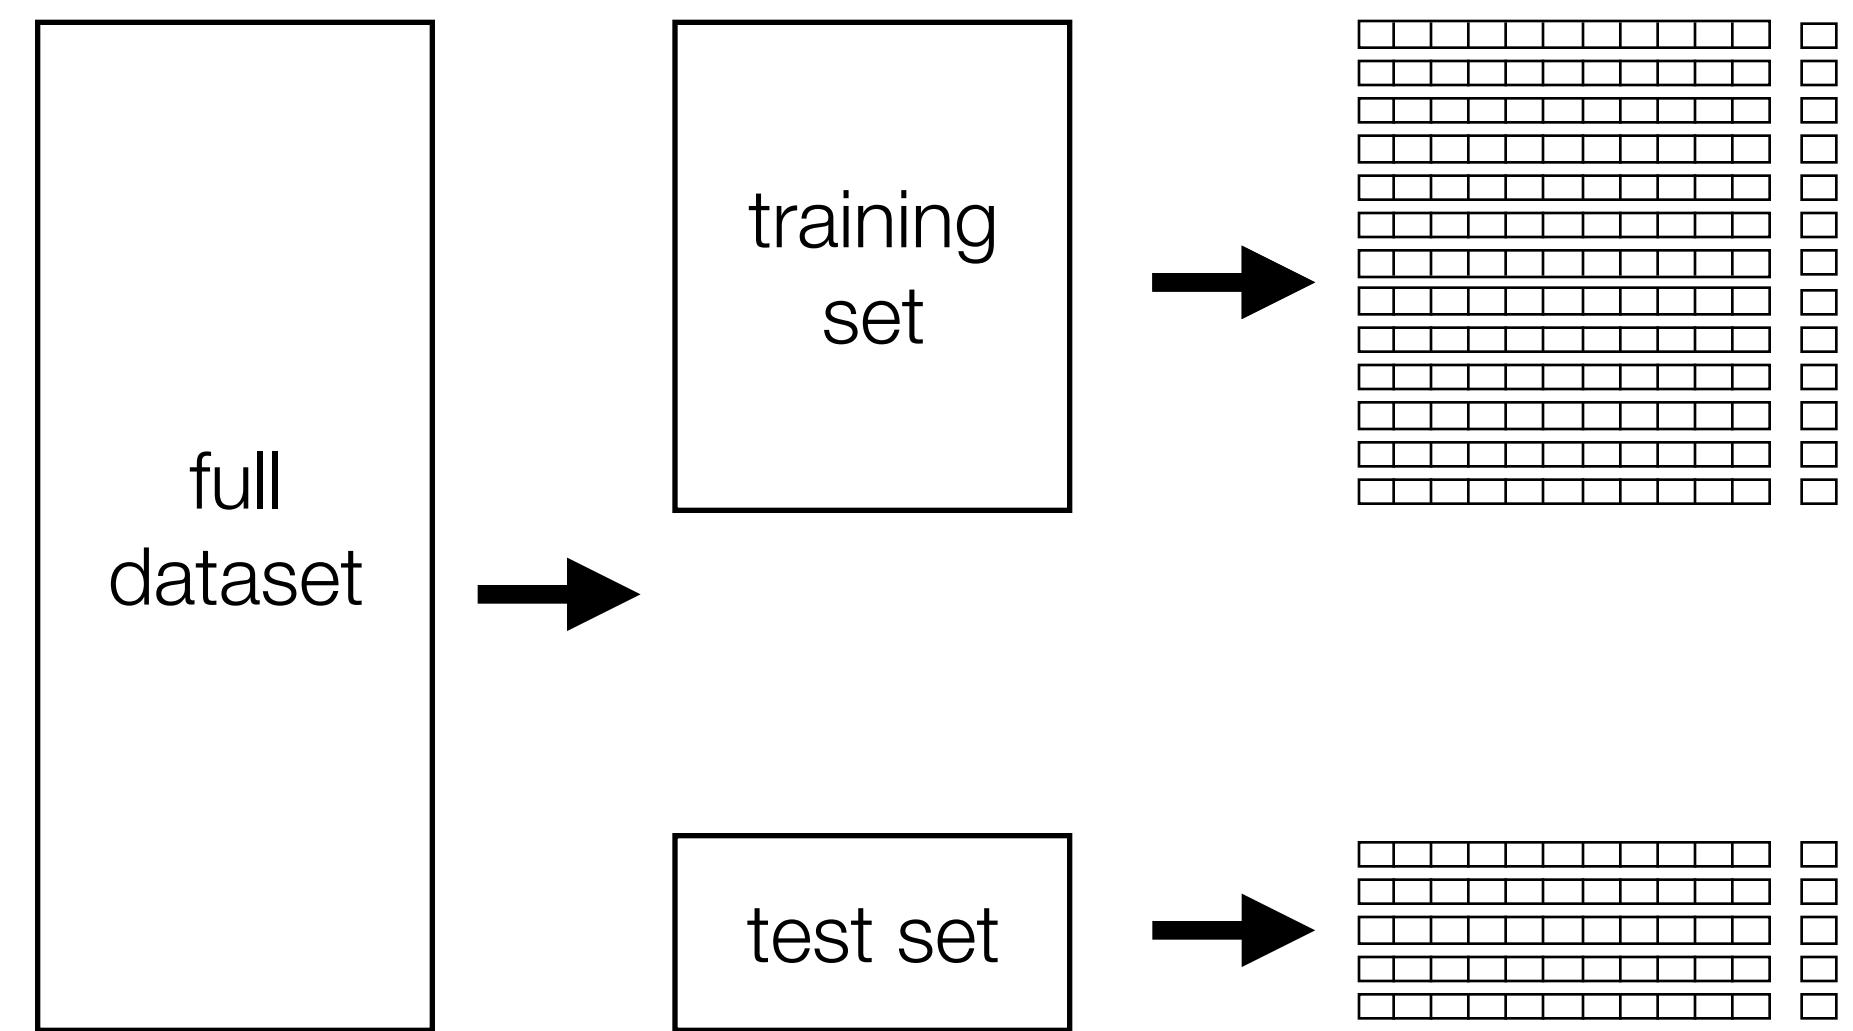


Obtain Raw Data

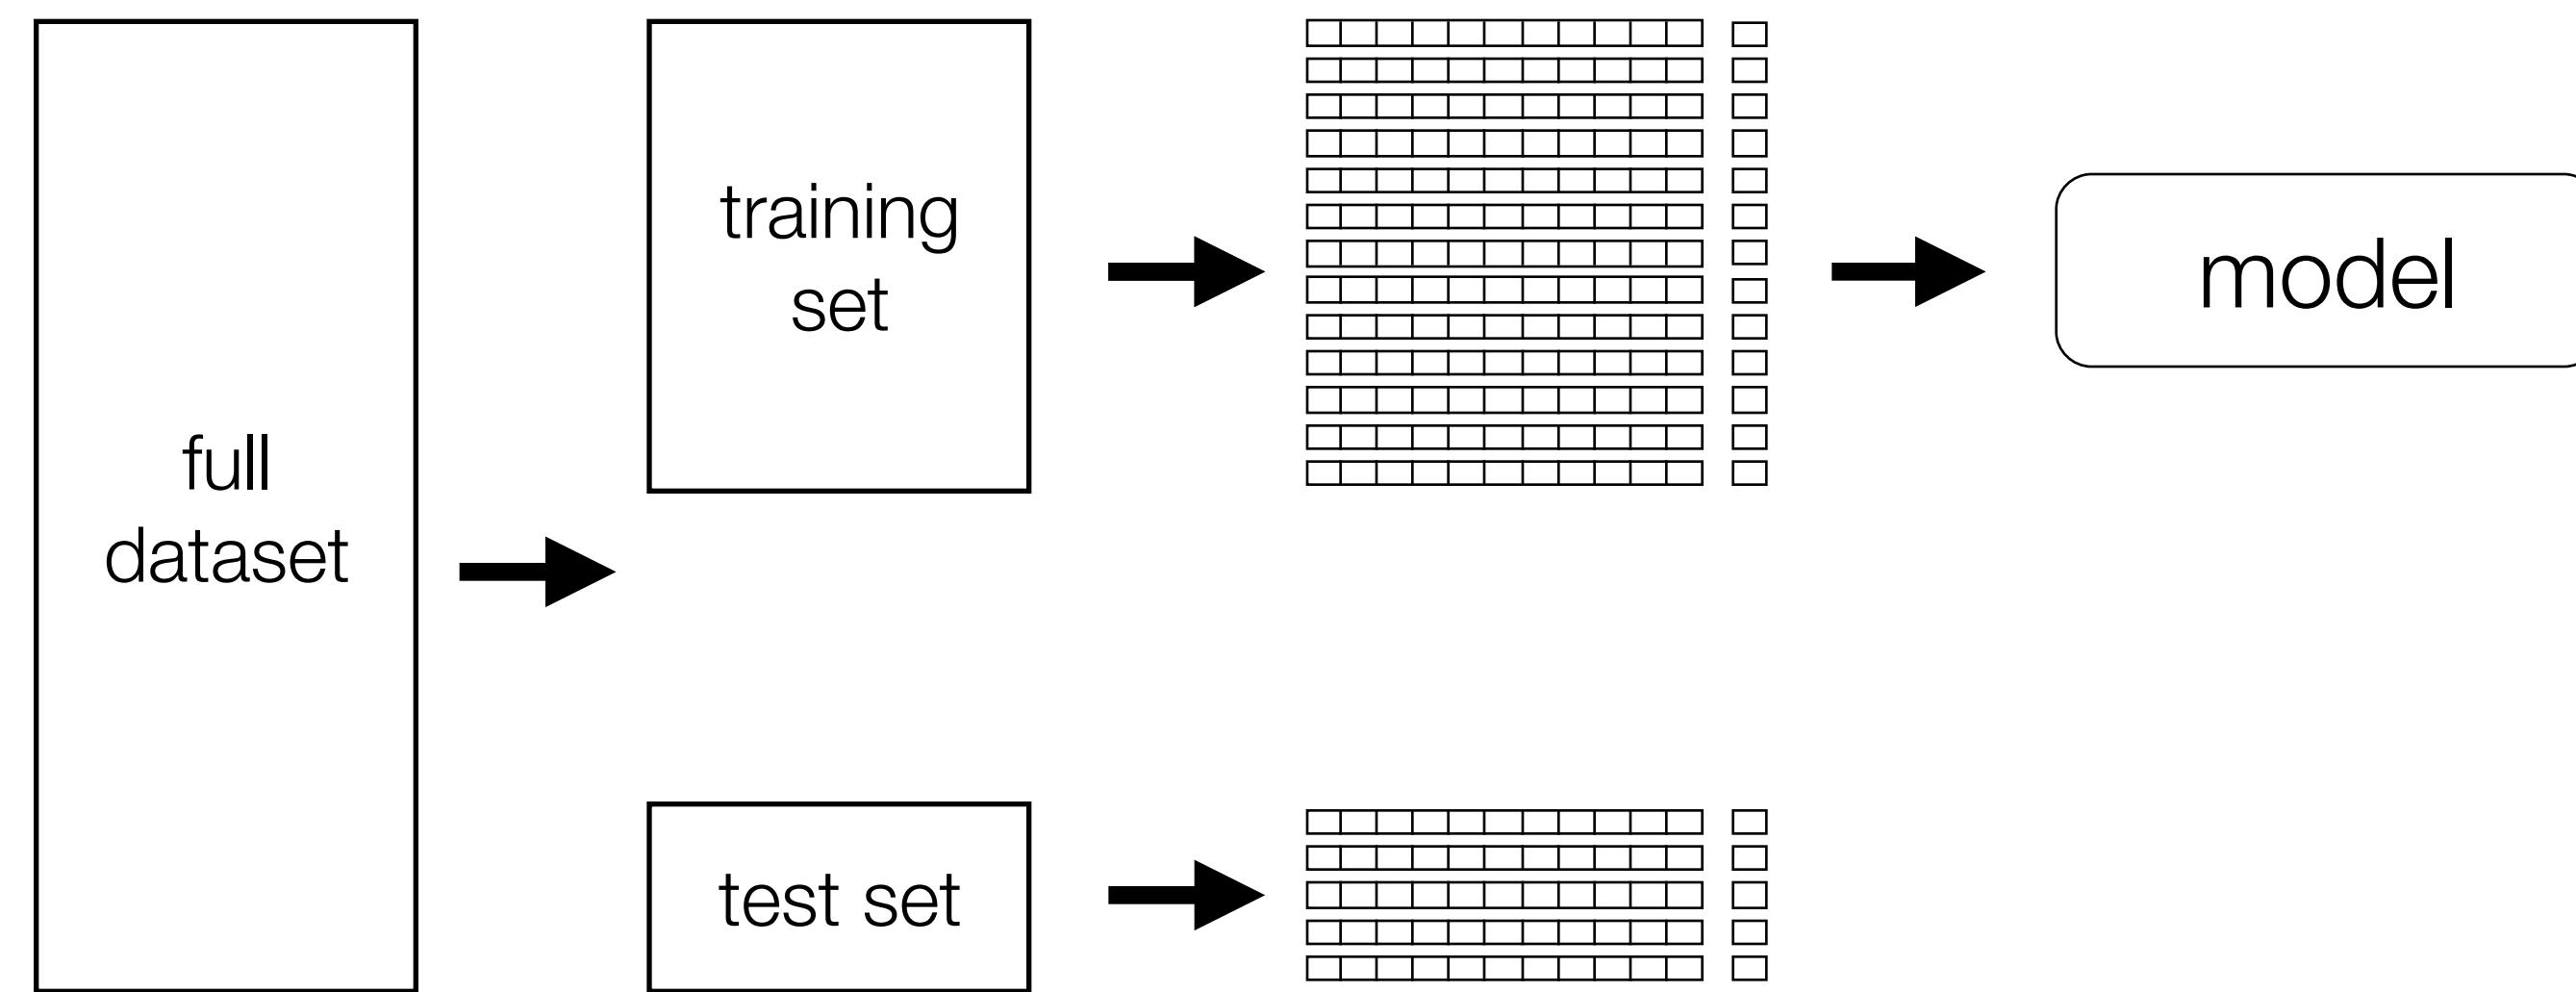
Supervised Learning Pipeline



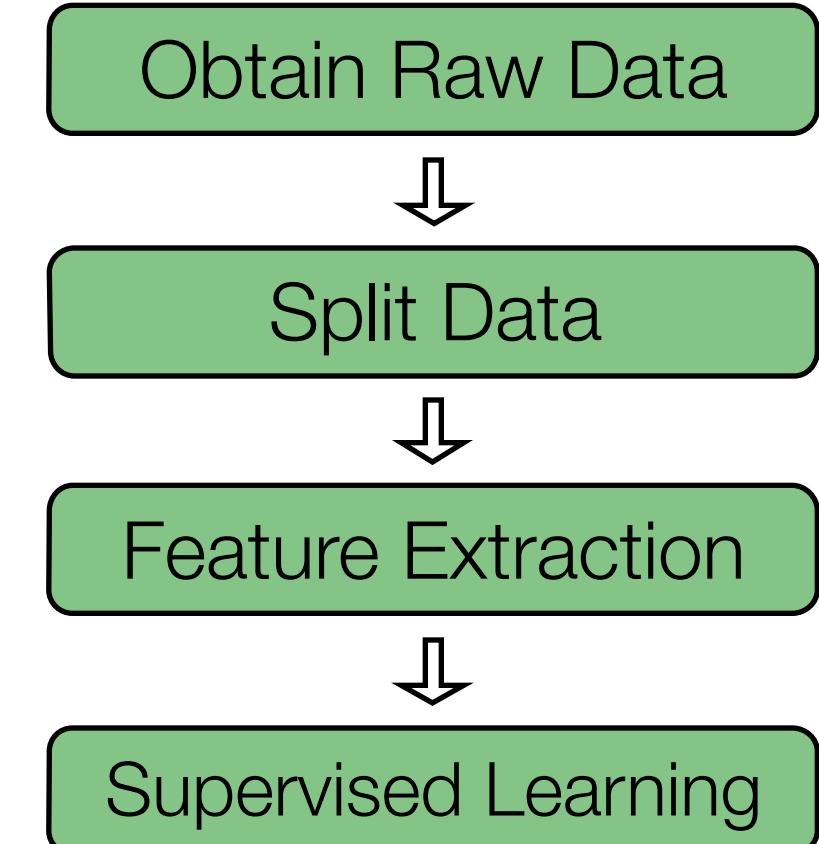
Supervised Learning Pipeline

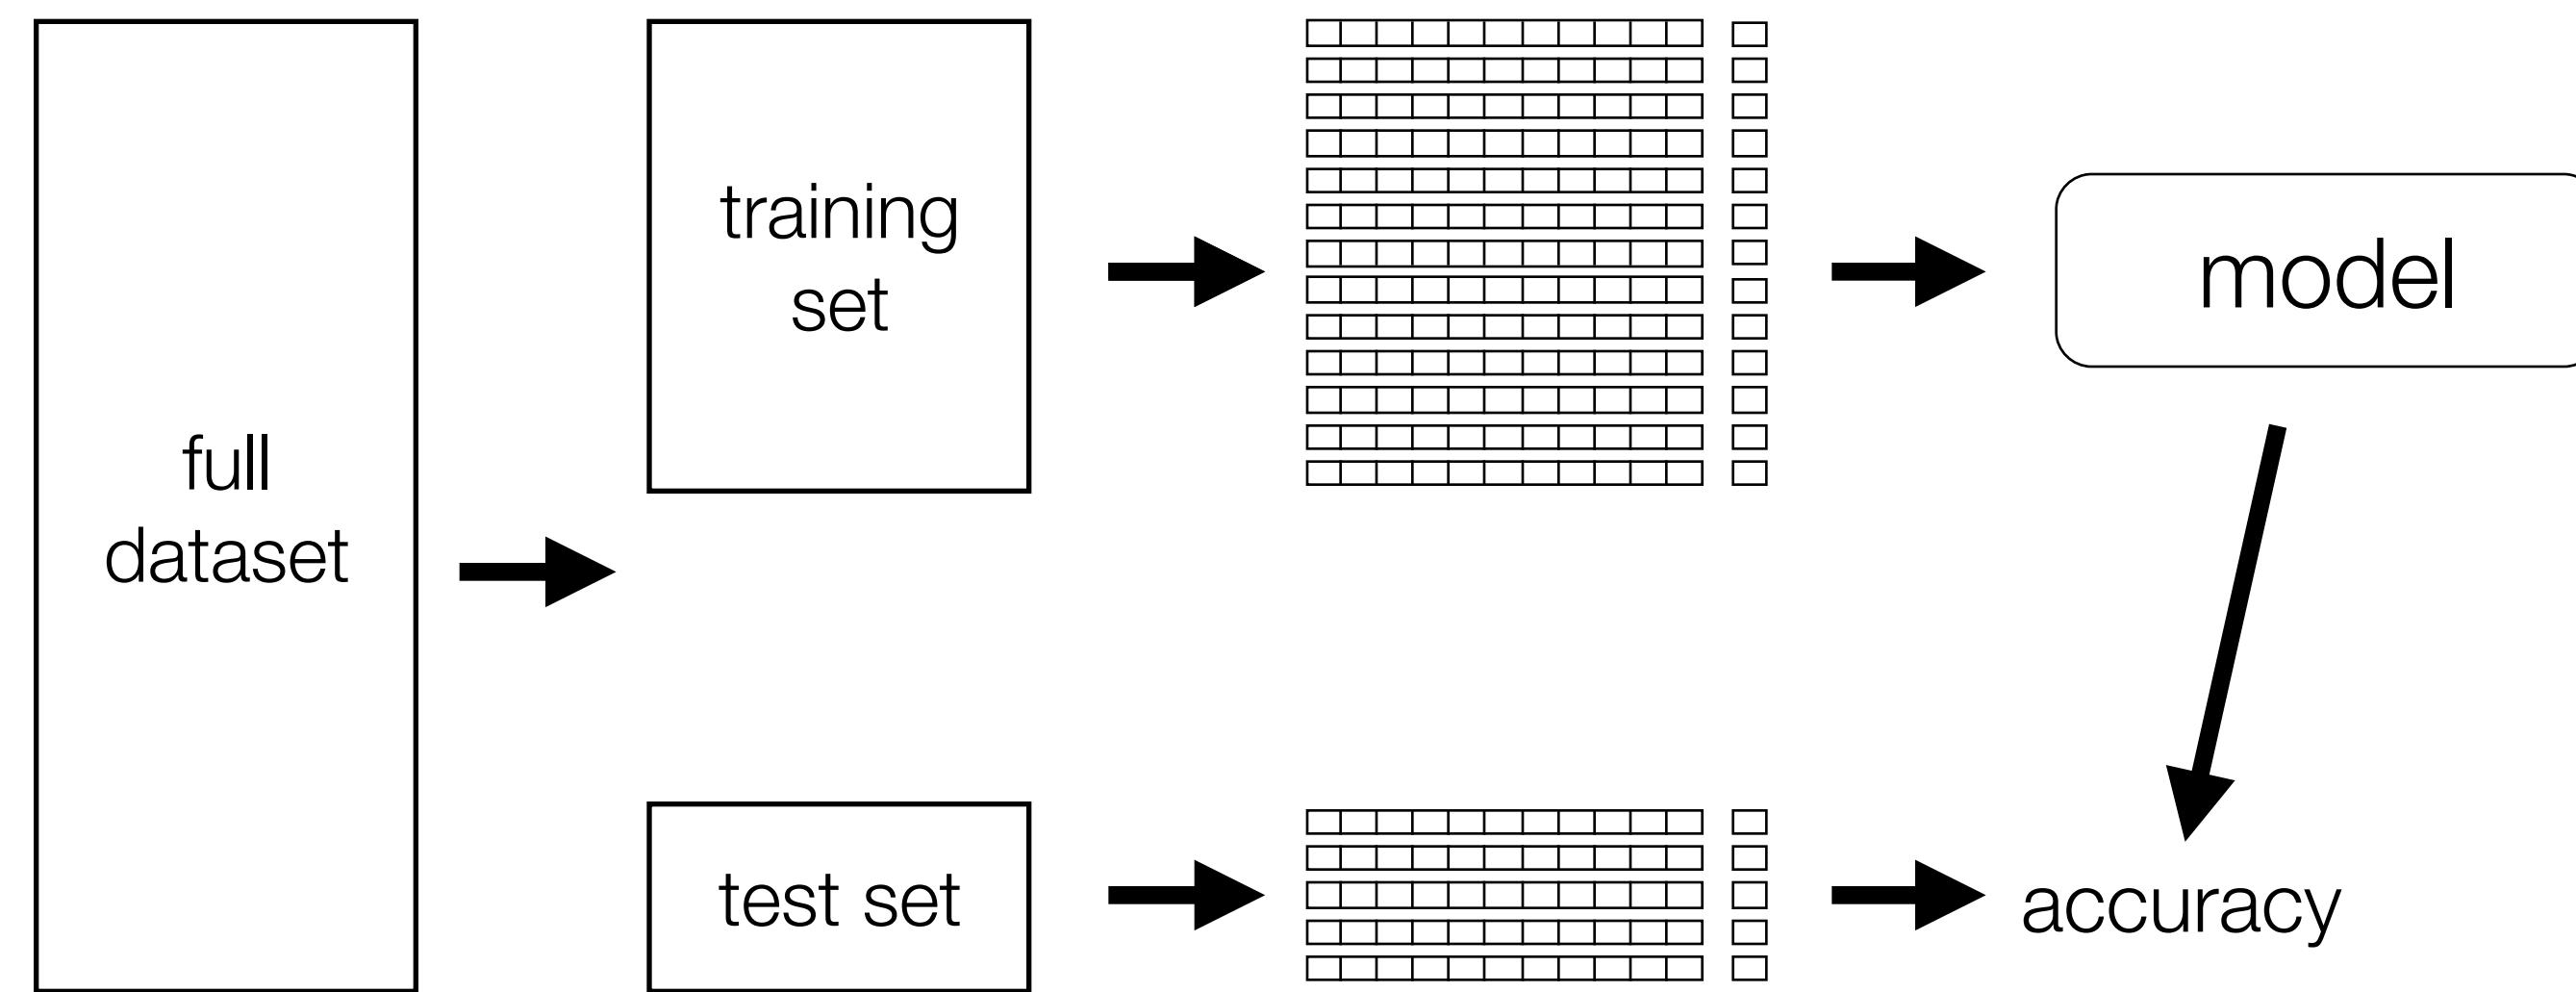


Supervised Learning Pipeline

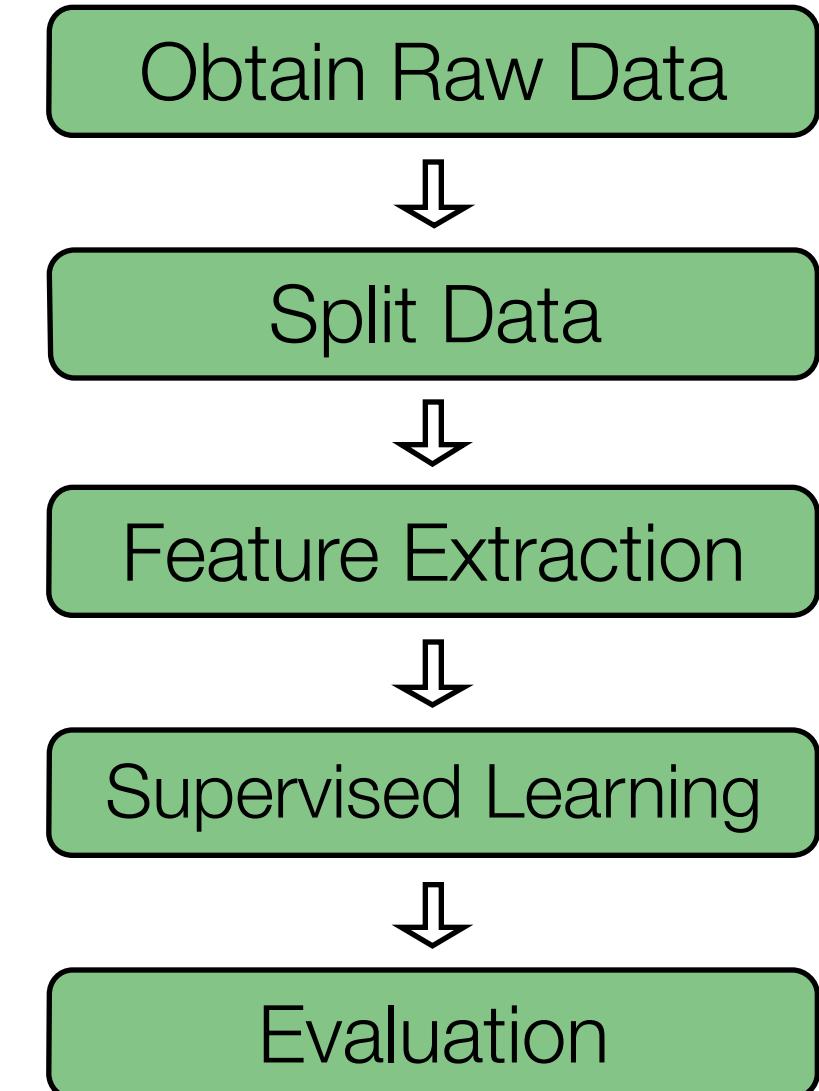


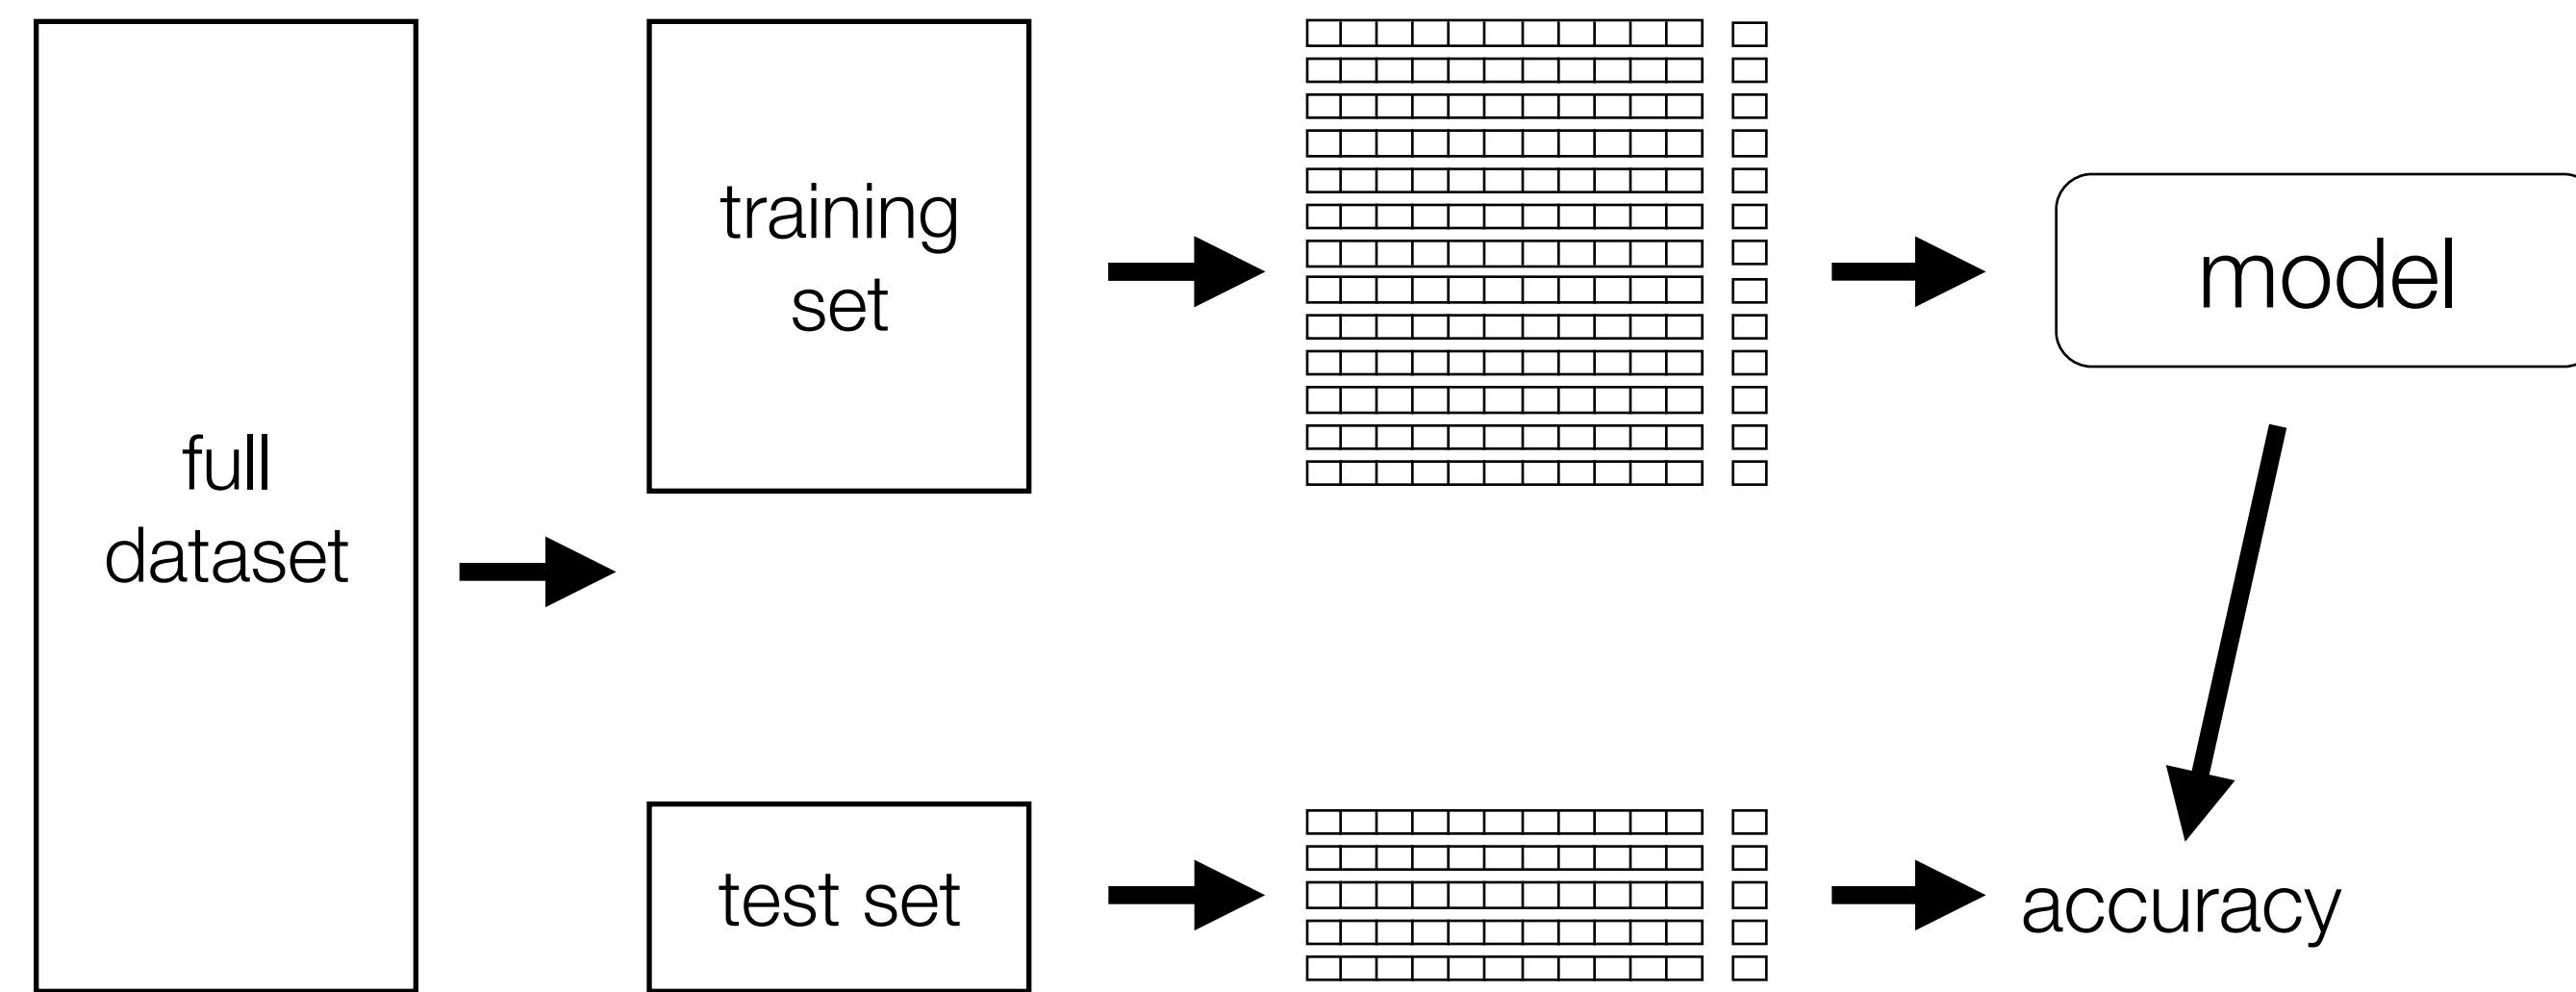
Supervised Learning Pipeline



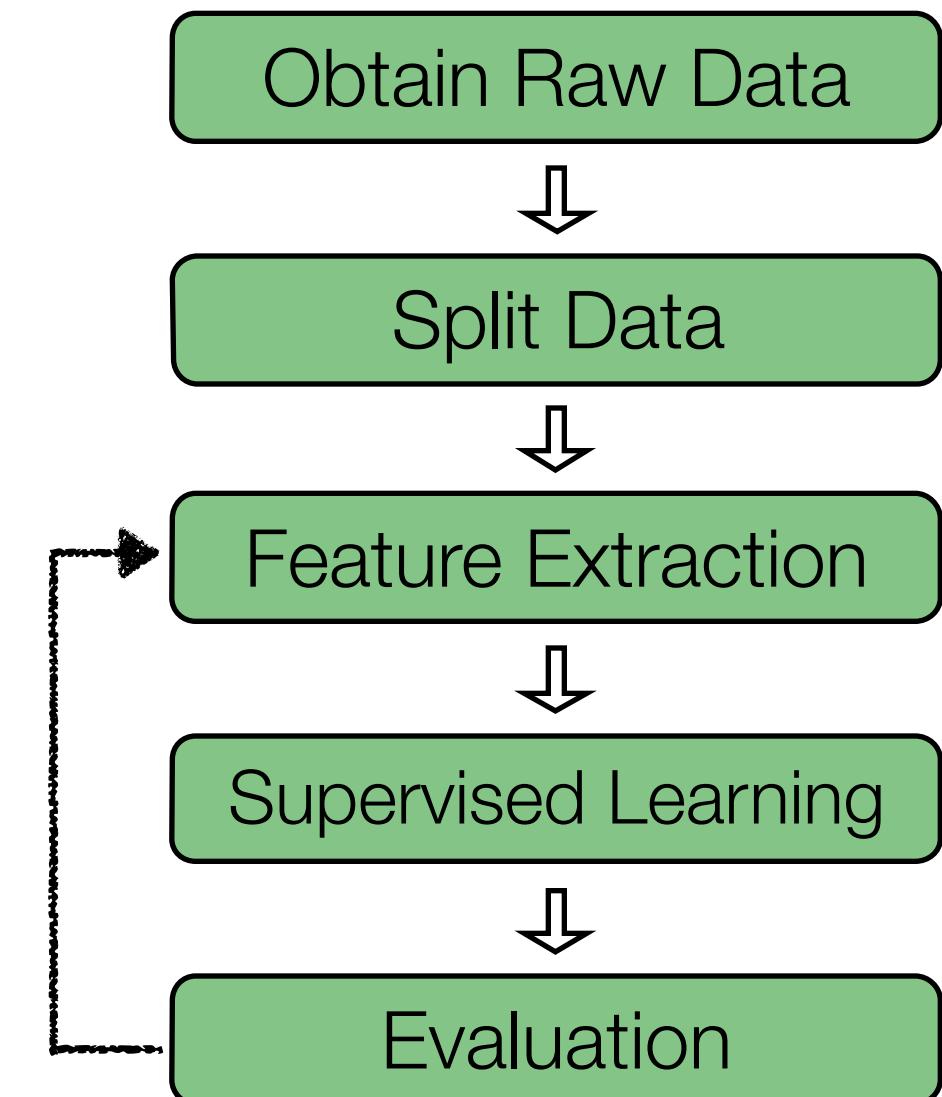


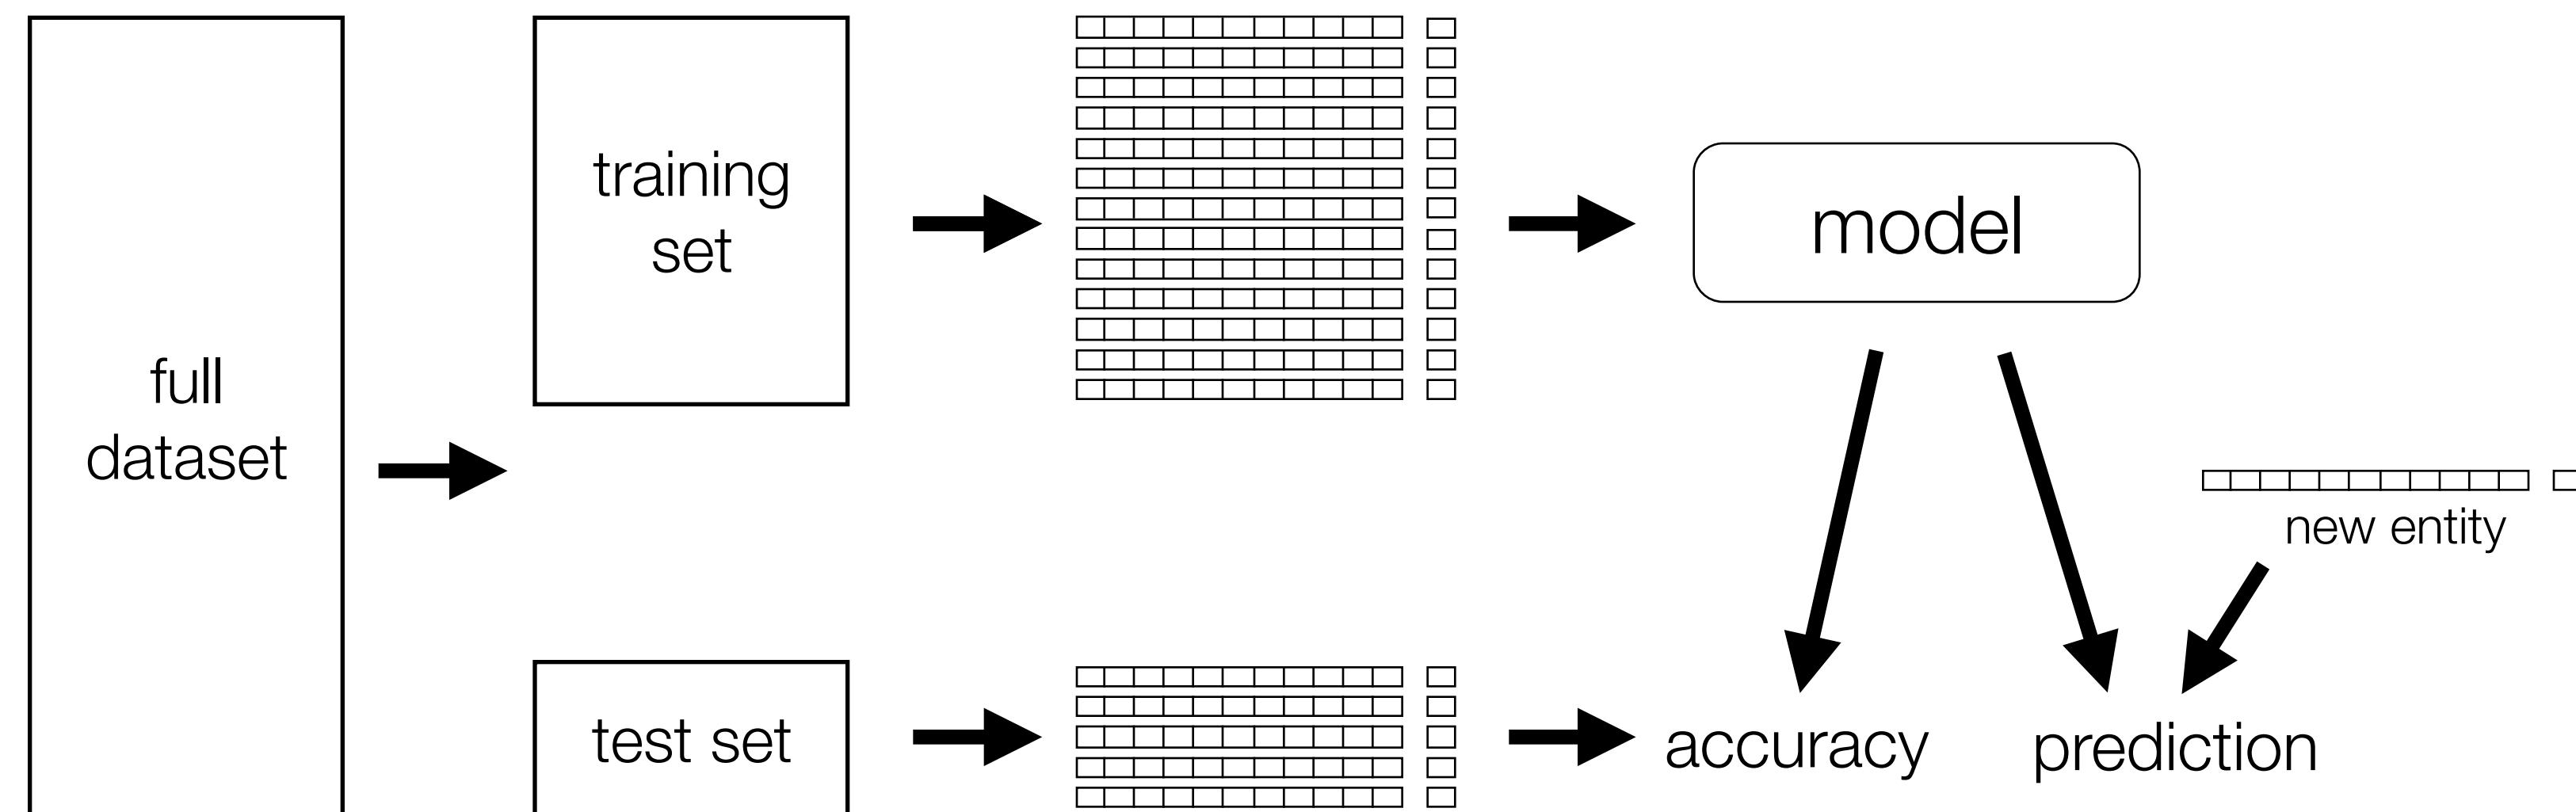
Supervised Learning Pipeline



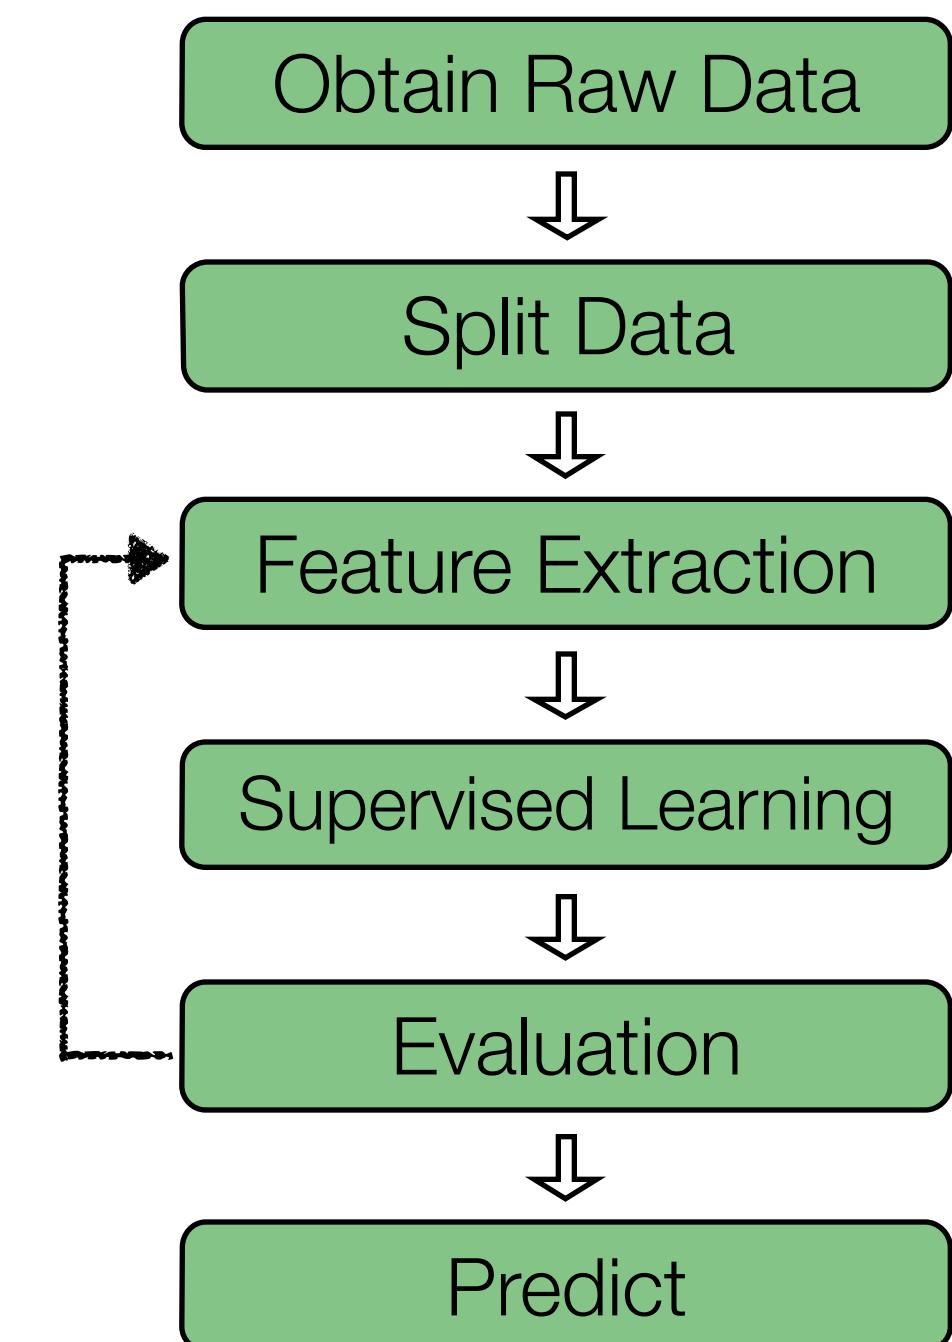


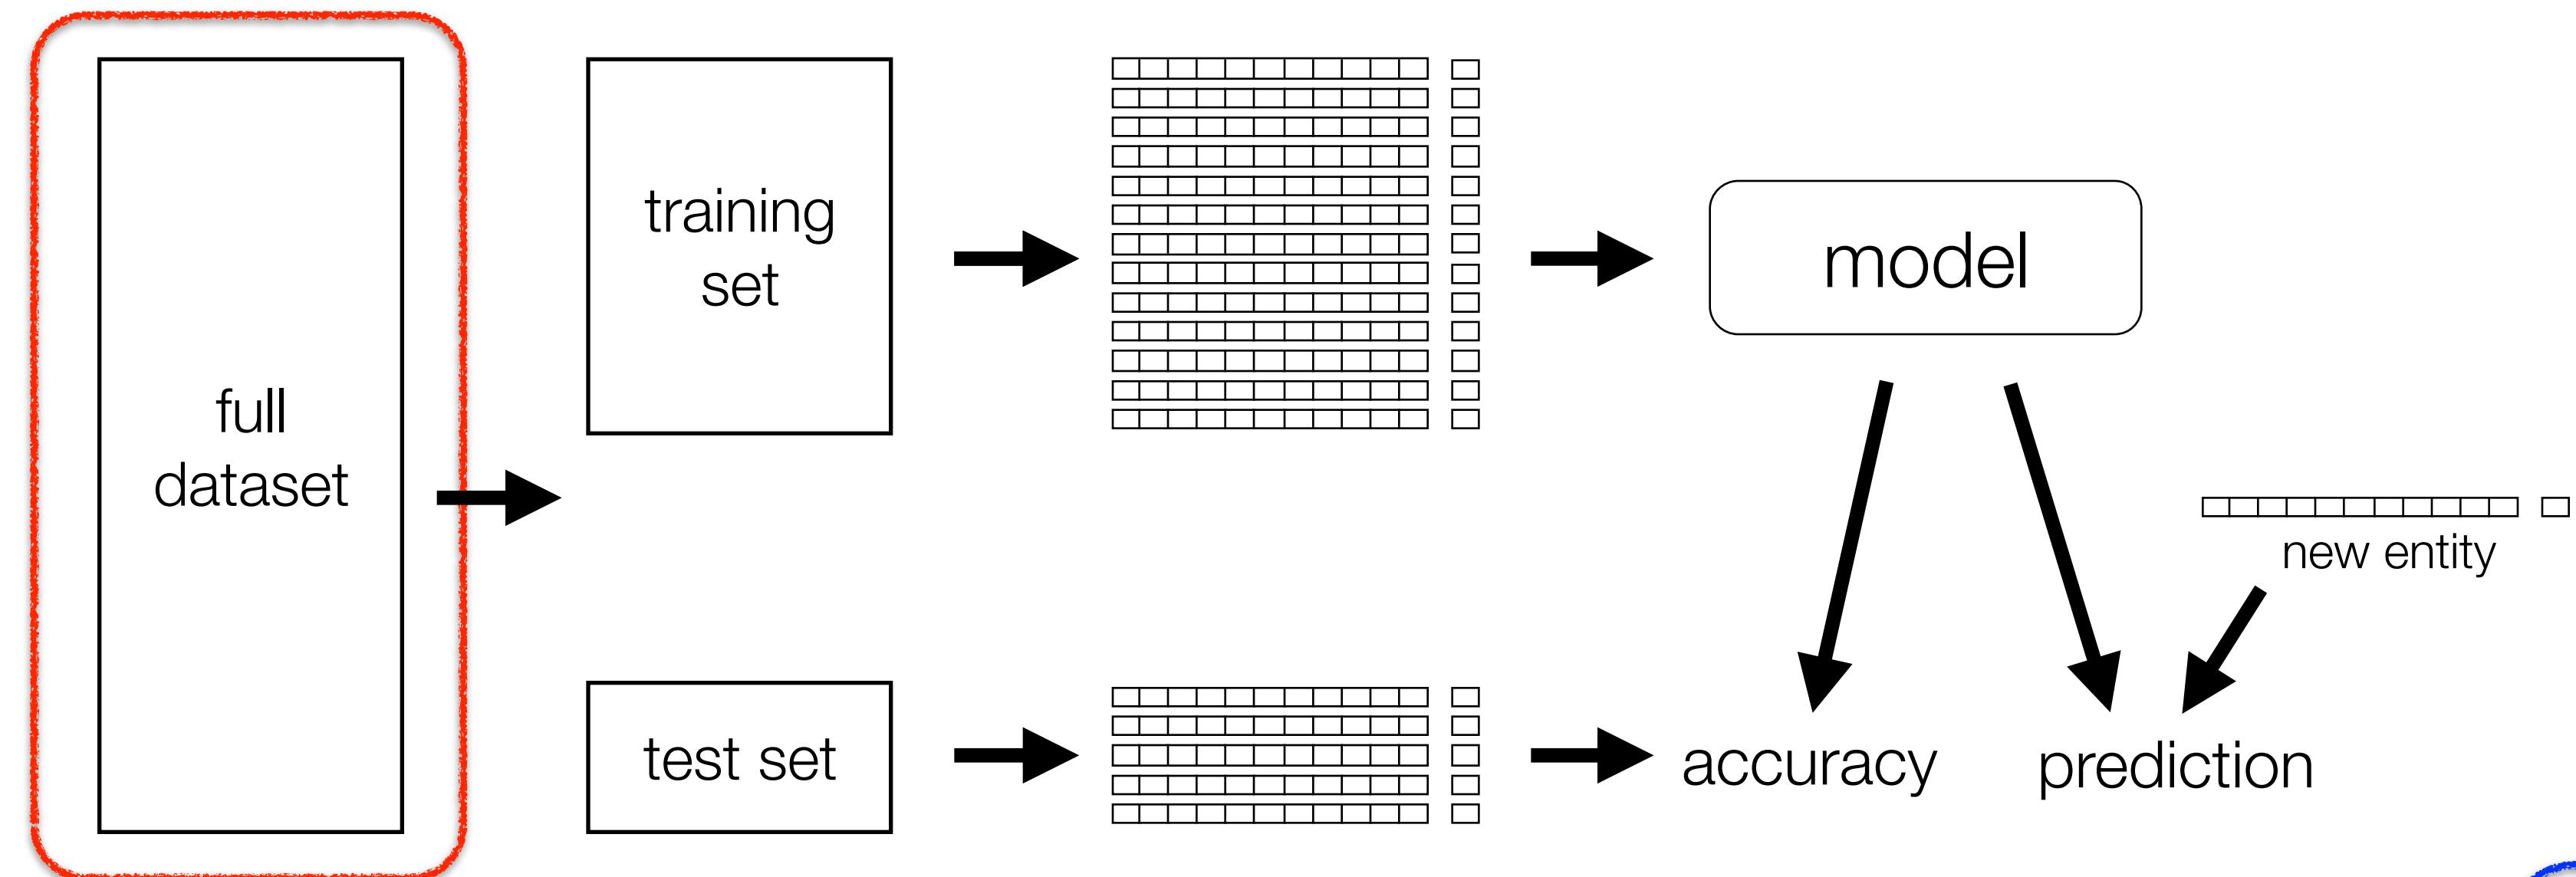
Supervised Learning Pipeline





Supervised Learning Pipeline

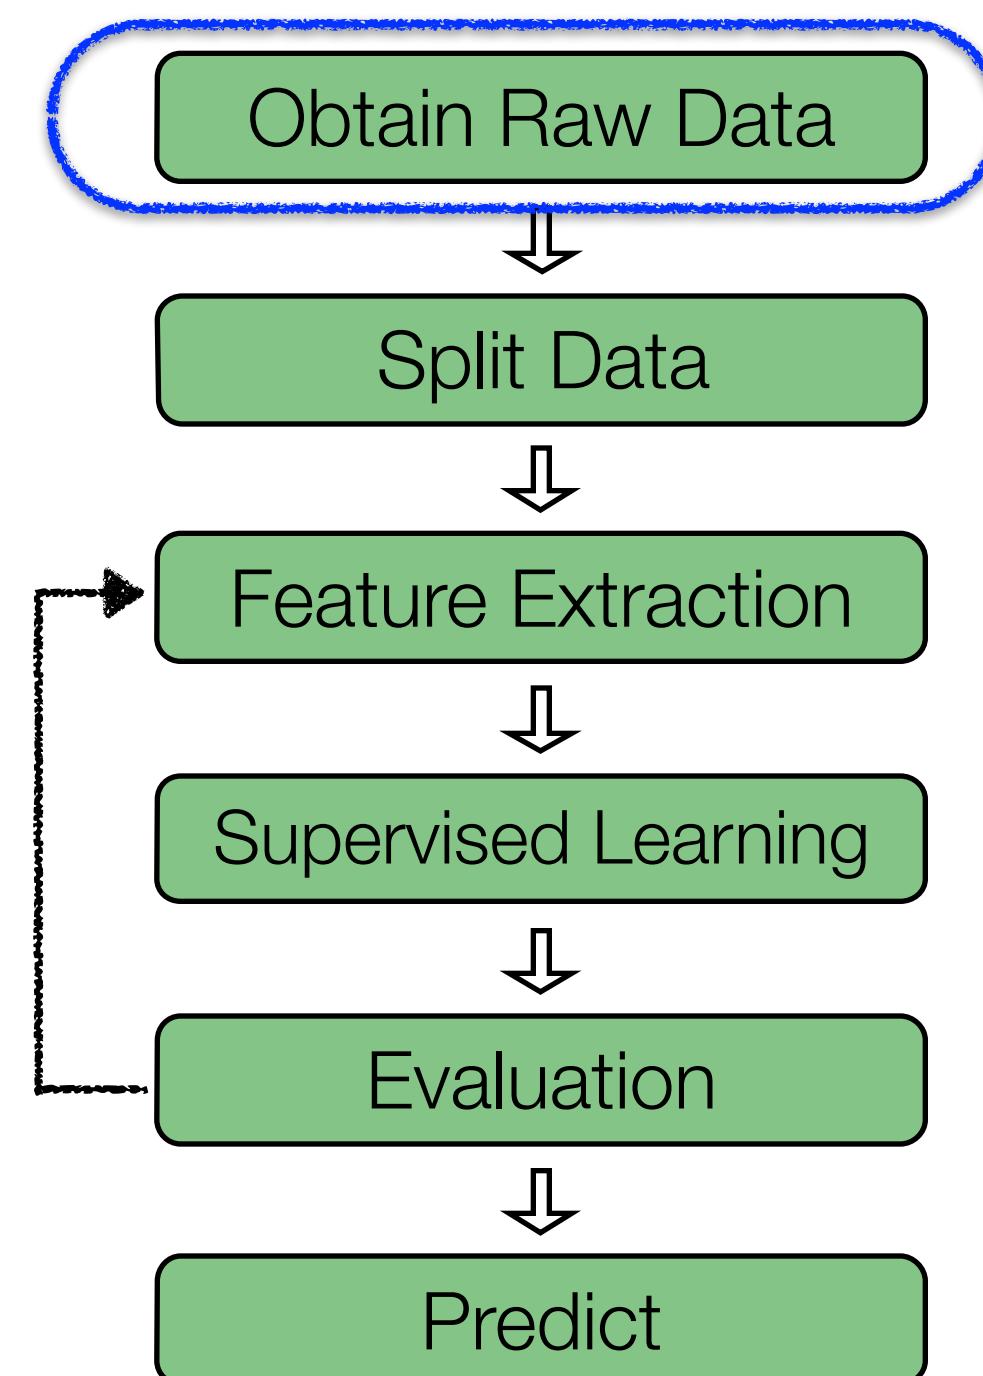


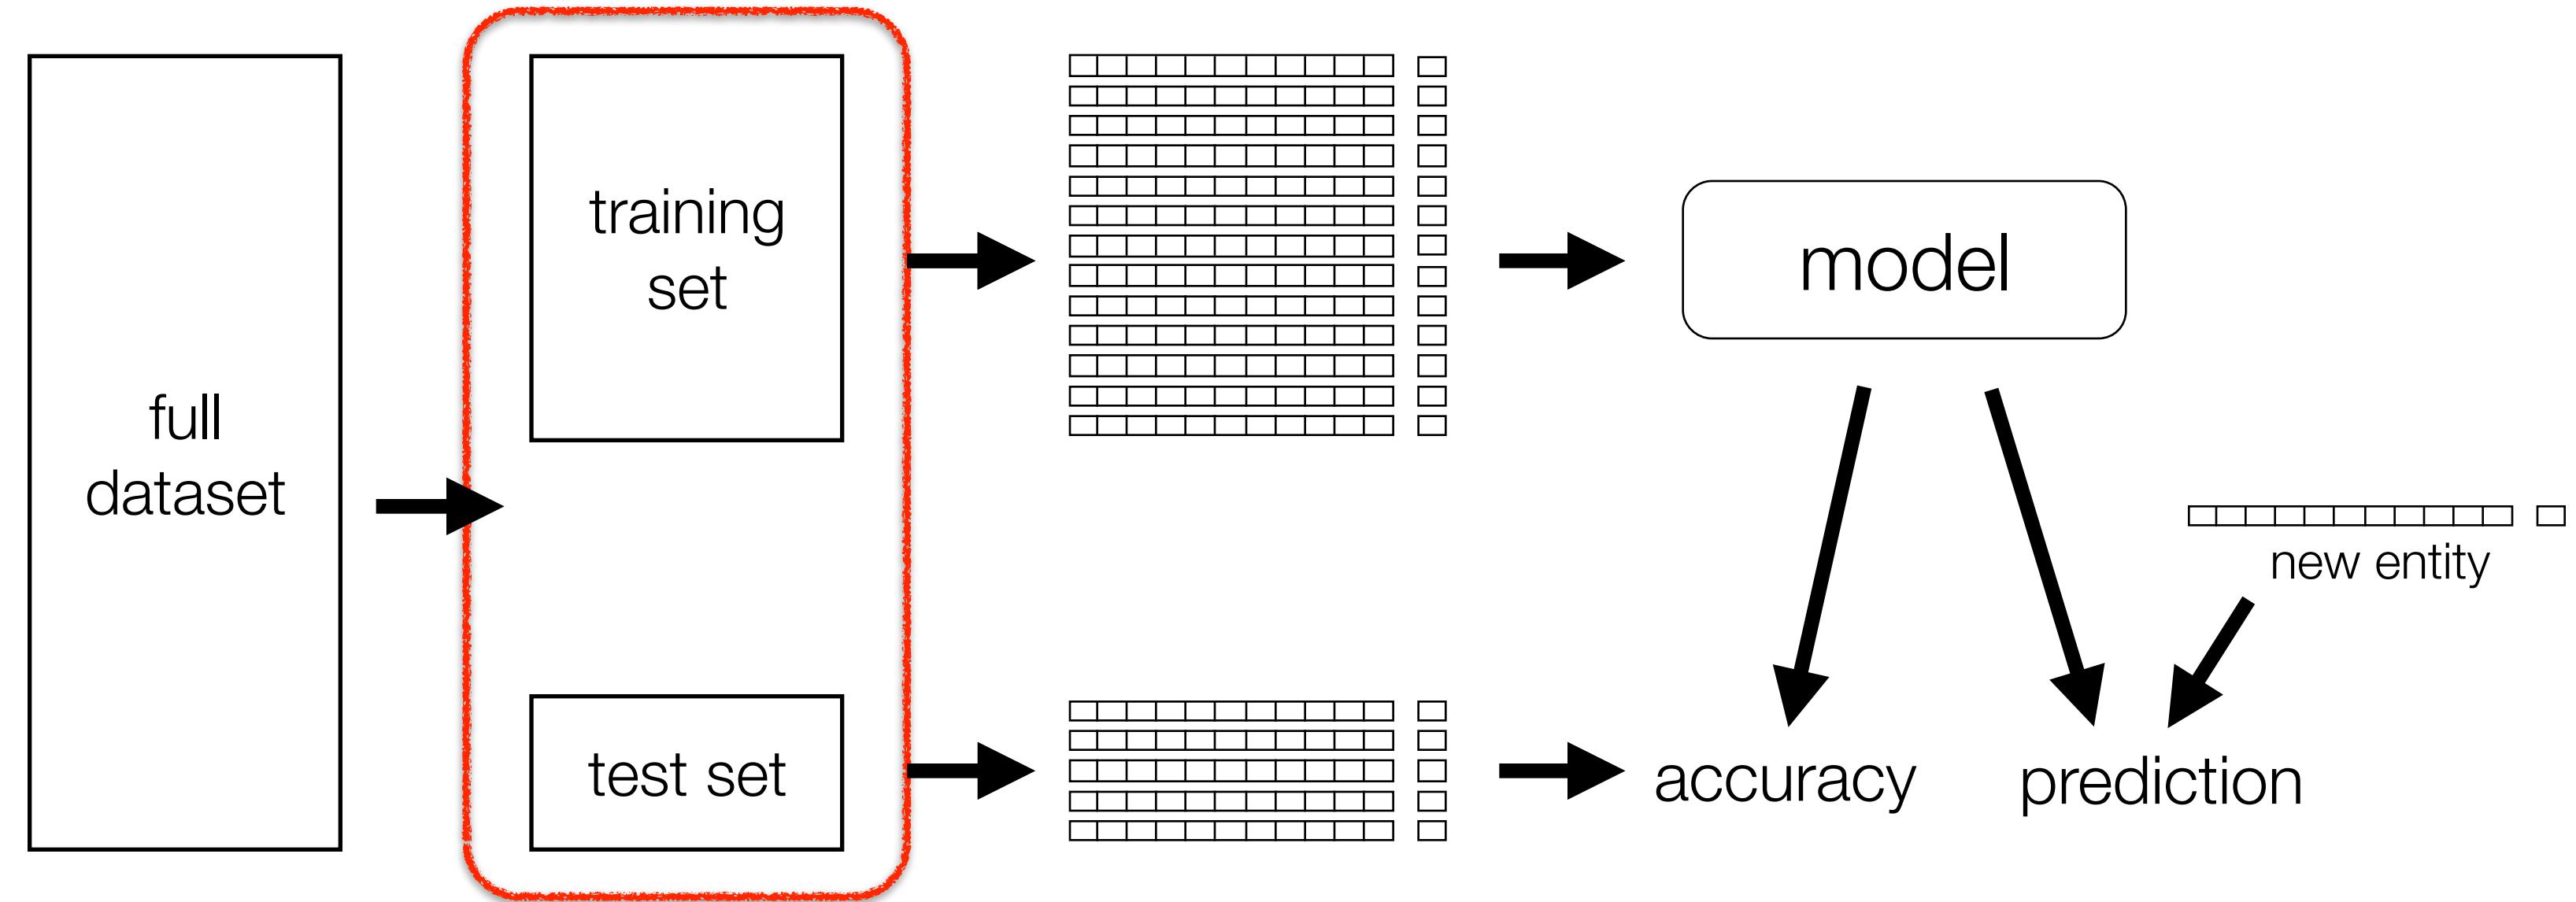


Goal: Predict song's release year from audio features

Raw Data: Millionsong Dataset from UCI ML Repository

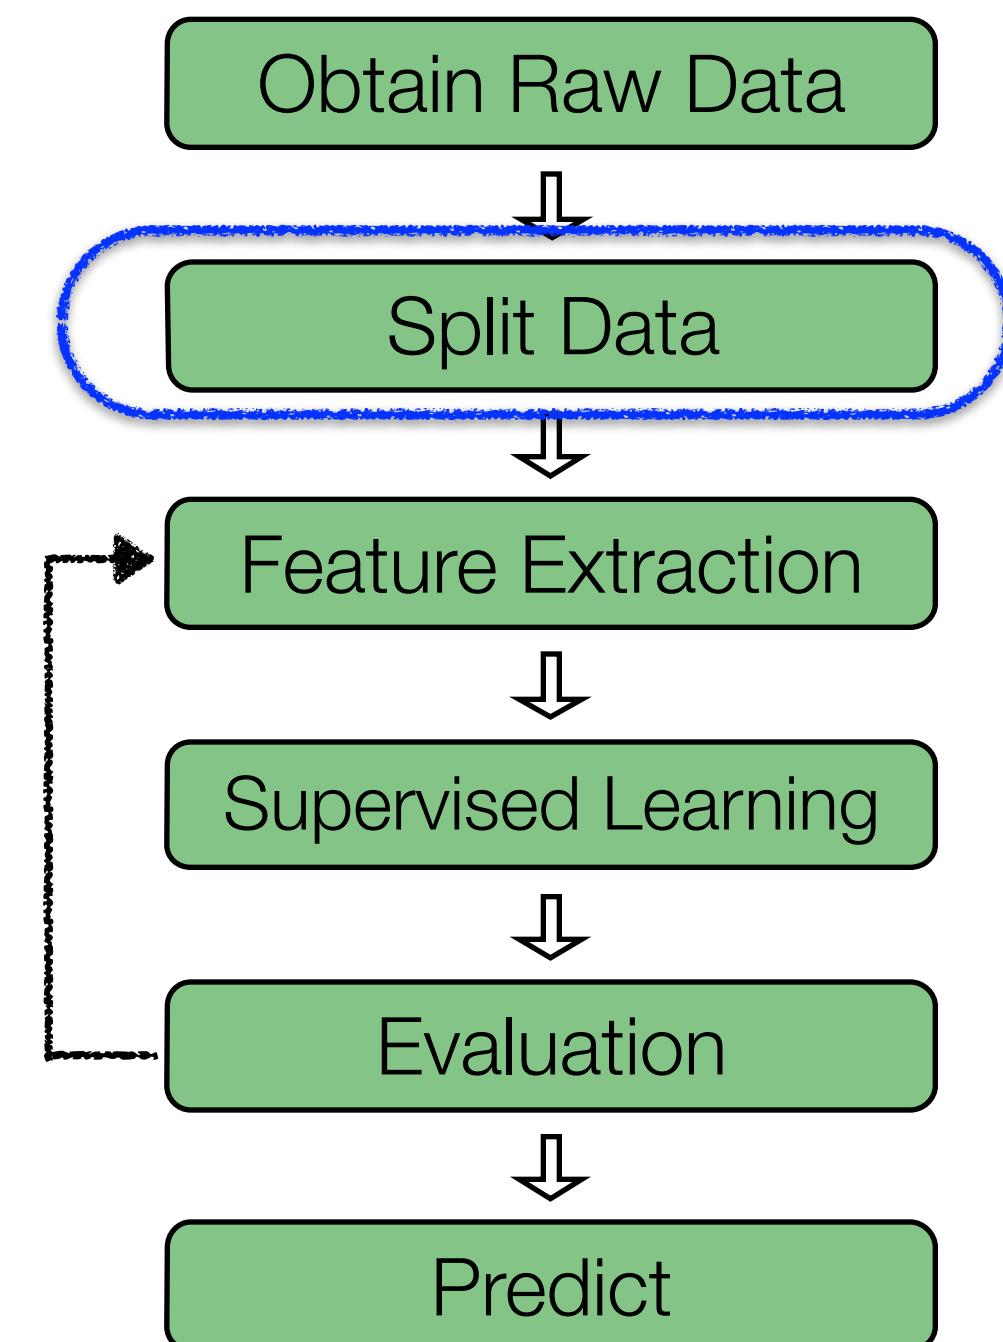
- Western, commercial tracks from 1980-2014
- 12 timbre averages (features) and release year (label)

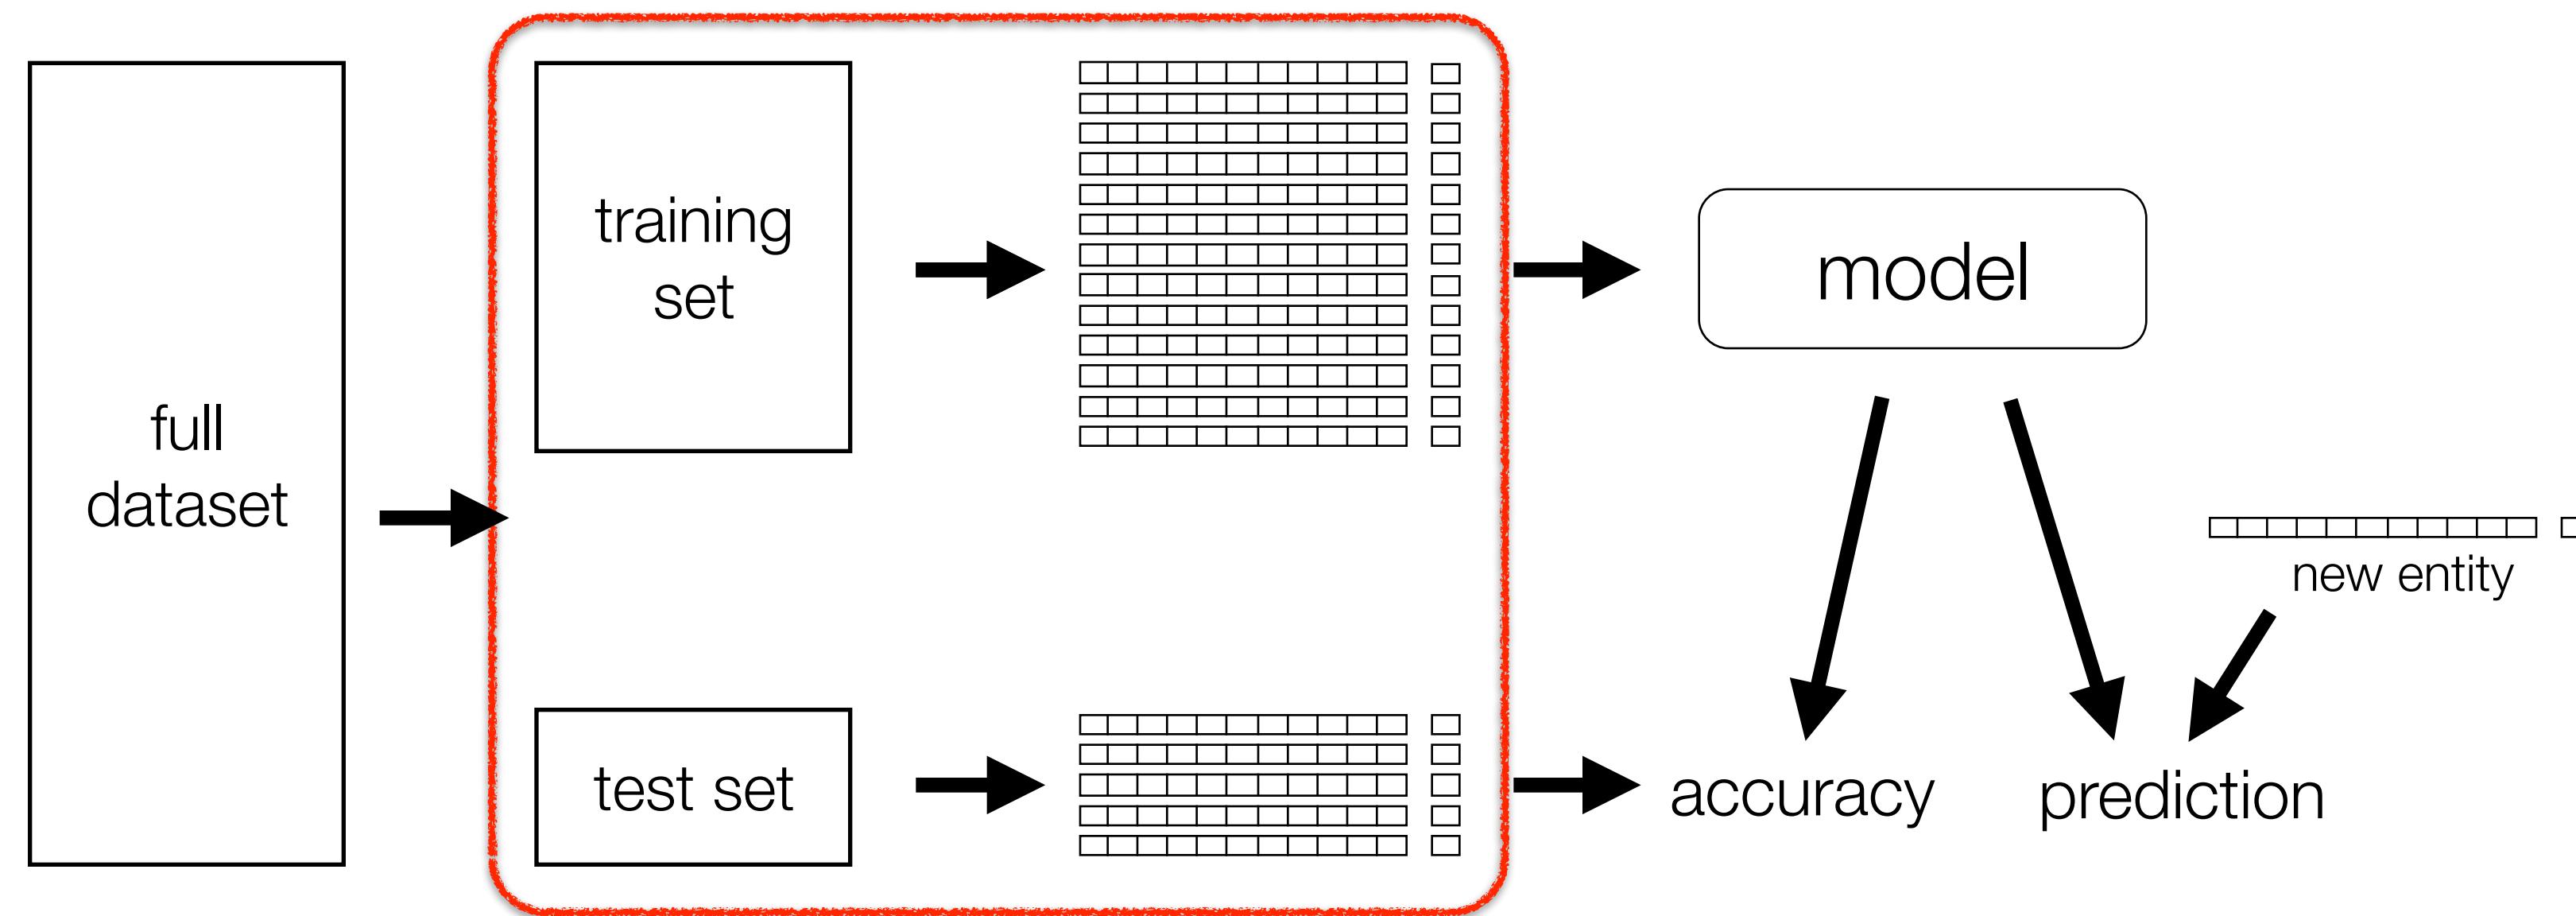




Split Data: Train on training set, evaluate with test set

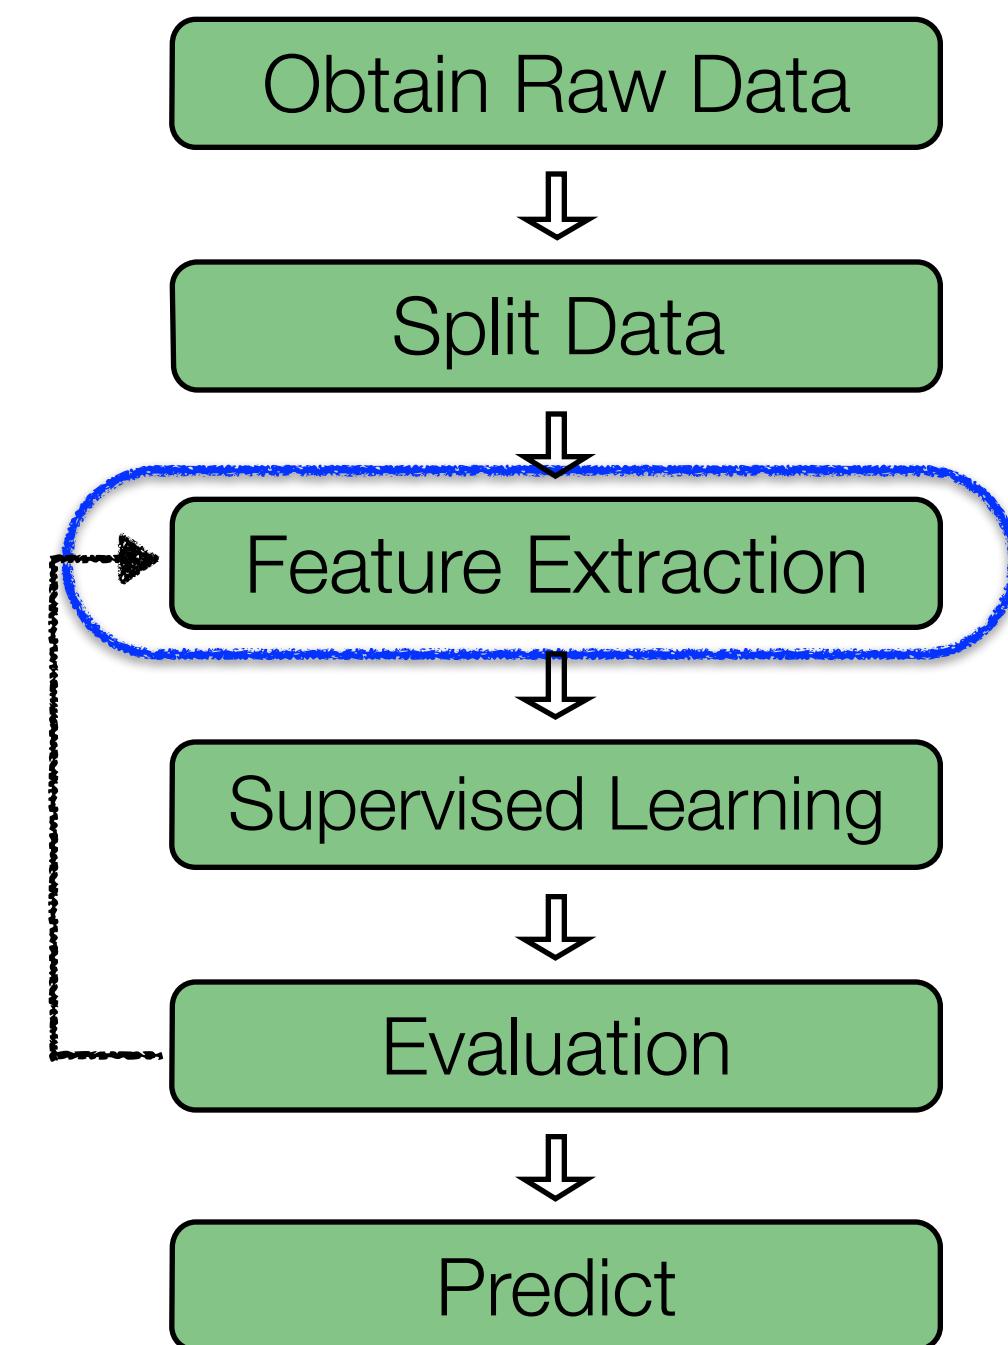
- Test set simulates unobserved data
- Test error tells us whether we've generalized well





Feature Extraction: Quadratic features

- Compute pairwise feature interactions
- Captures covariance of initial timbre features
- Leads to a non-linear model relative to raw features



Given 2 dimensional data, quadratic features are:

$$\mathbf{x} = [x_1 \quad x_2]^\top \implies \Phi(\mathbf{x}) = [x_1^2 \quad x_1x_2 \quad x_2x_1 \quad x_2^2]^\top$$

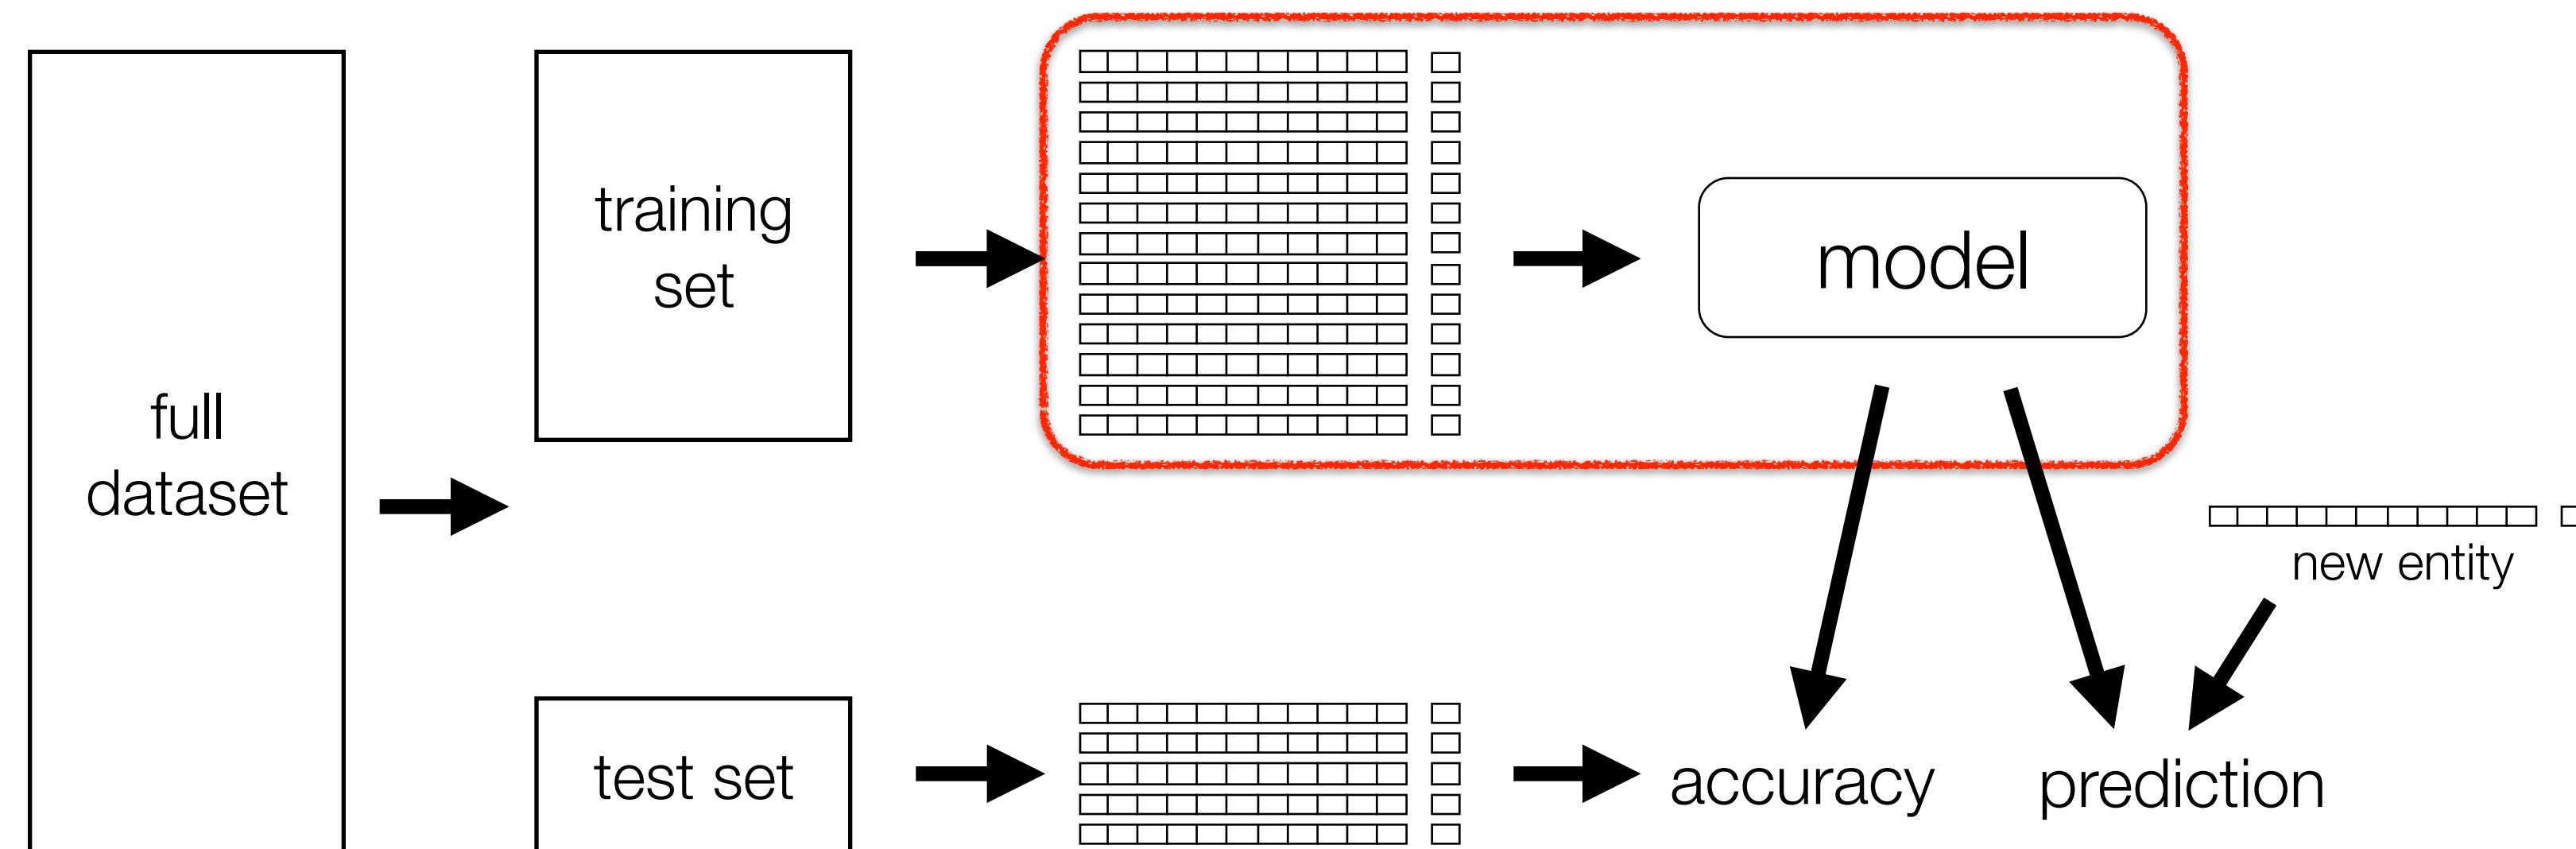
$$\mathbf{z} = [z_1 \quad z_2]^\top \implies \Phi(\mathbf{z}) = [z_1^2 \quad z_1z_2 \quad z_2z_1 \quad z_2^2]^\top$$

More succinctly:

$$\Phi'(\mathbf{x}) = [x_1^2 \quad \sqrt{2}x_1x_2 \quad x_2^2]^\top \quad \Phi'(\mathbf{z}) = [z_1^2 \quad \sqrt{2}z_1z_2 \quad z_2^2]^\top$$

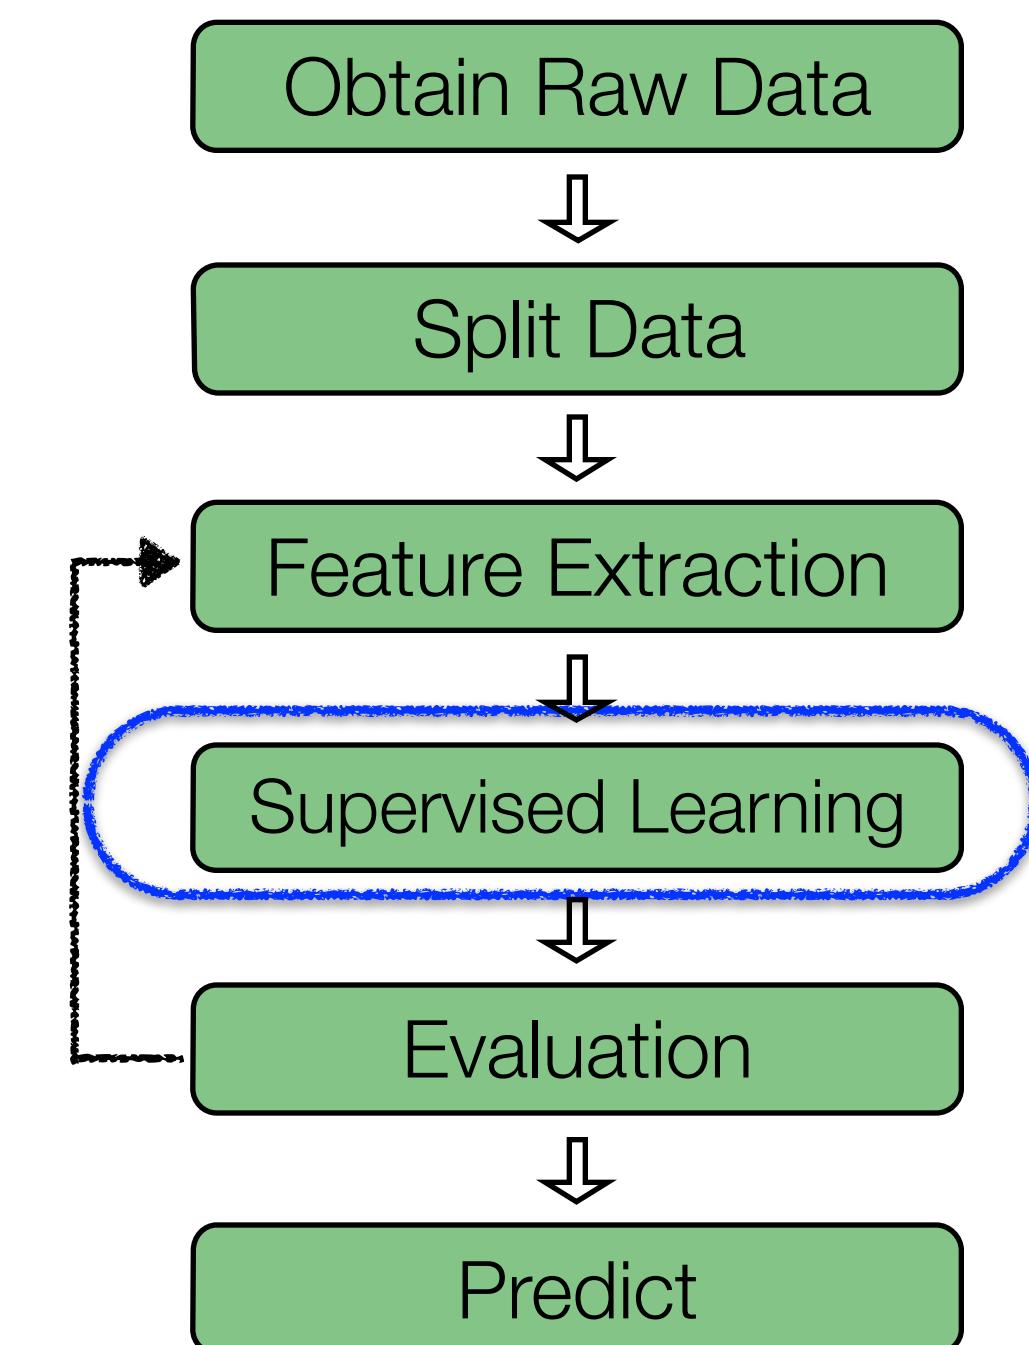
Equivalent inner products:

$$\Phi(\mathbf{x})^\top \Phi(\mathbf{z}) = x_1^2 z_1^2 + 2x_1 x_2 z_1 z_2 + x_2^2 z_2^2 = \Phi'(\mathbf{x})^\top \Phi'(\mathbf{z})$$



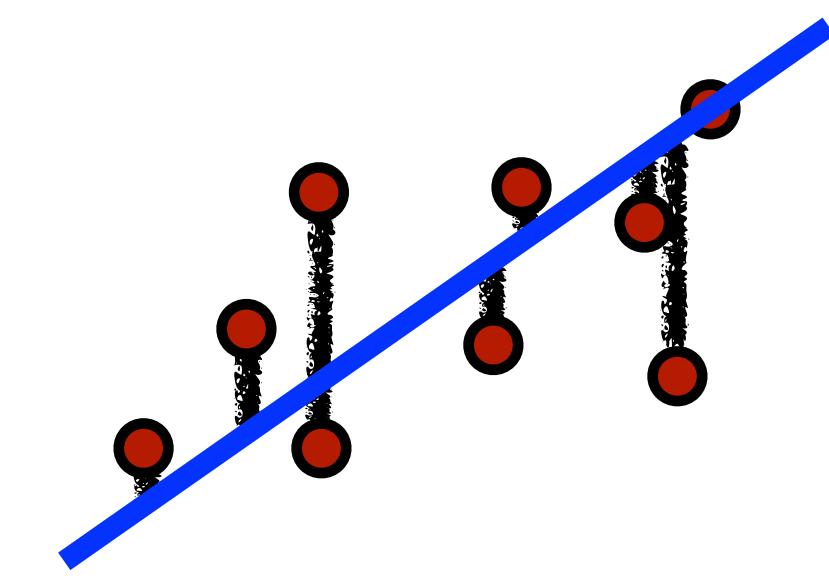
Supervised Learning: Least Squares Regression

- Learn a mapping from entities to continuous labels given a training set
- Audio features → Song year



Given n training points with d features, we define:

- $\mathbf{X} \in \mathbb{R}^{n \times d}$: matrix storing points
- $\mathbf{y} \in \mathbb{R}^n$: real-valued labels
- $\hat{\mathbf{y}} \in \mathbb{R}^n$: predicted labels, where $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$
- $\mathbf{w} \in \mathbb{R}^d$: regression parameters / model to learn



Ridge Regression: Learn mapping (\mathbf{w}) that minimizes residual sum of squares along with a regularization term:

$$\min_{\mathbf{w}} \frac{\text{Training Error}}{\|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2} + \frac{\text{Model Complexity}}{\lambda \|\mathbf{w}\|_2^2}$$

Closed-form solution: $\mathbf{w} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_d)^{-1} \mathbf{X}^\top \mathbf{y}$

Ridge Regression: Learn mapping (\mathbf{w}) that minimizes residual sum of squares along with a regularization term:

$$\min_{\mathbf{w}} \frac{\text{Training Error}}{\|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2} + \frac{\text{Model Complexity}}{\lambda \|\mathbf{w}\|_2^2}$$

free parameter trades off between training error and model complexity

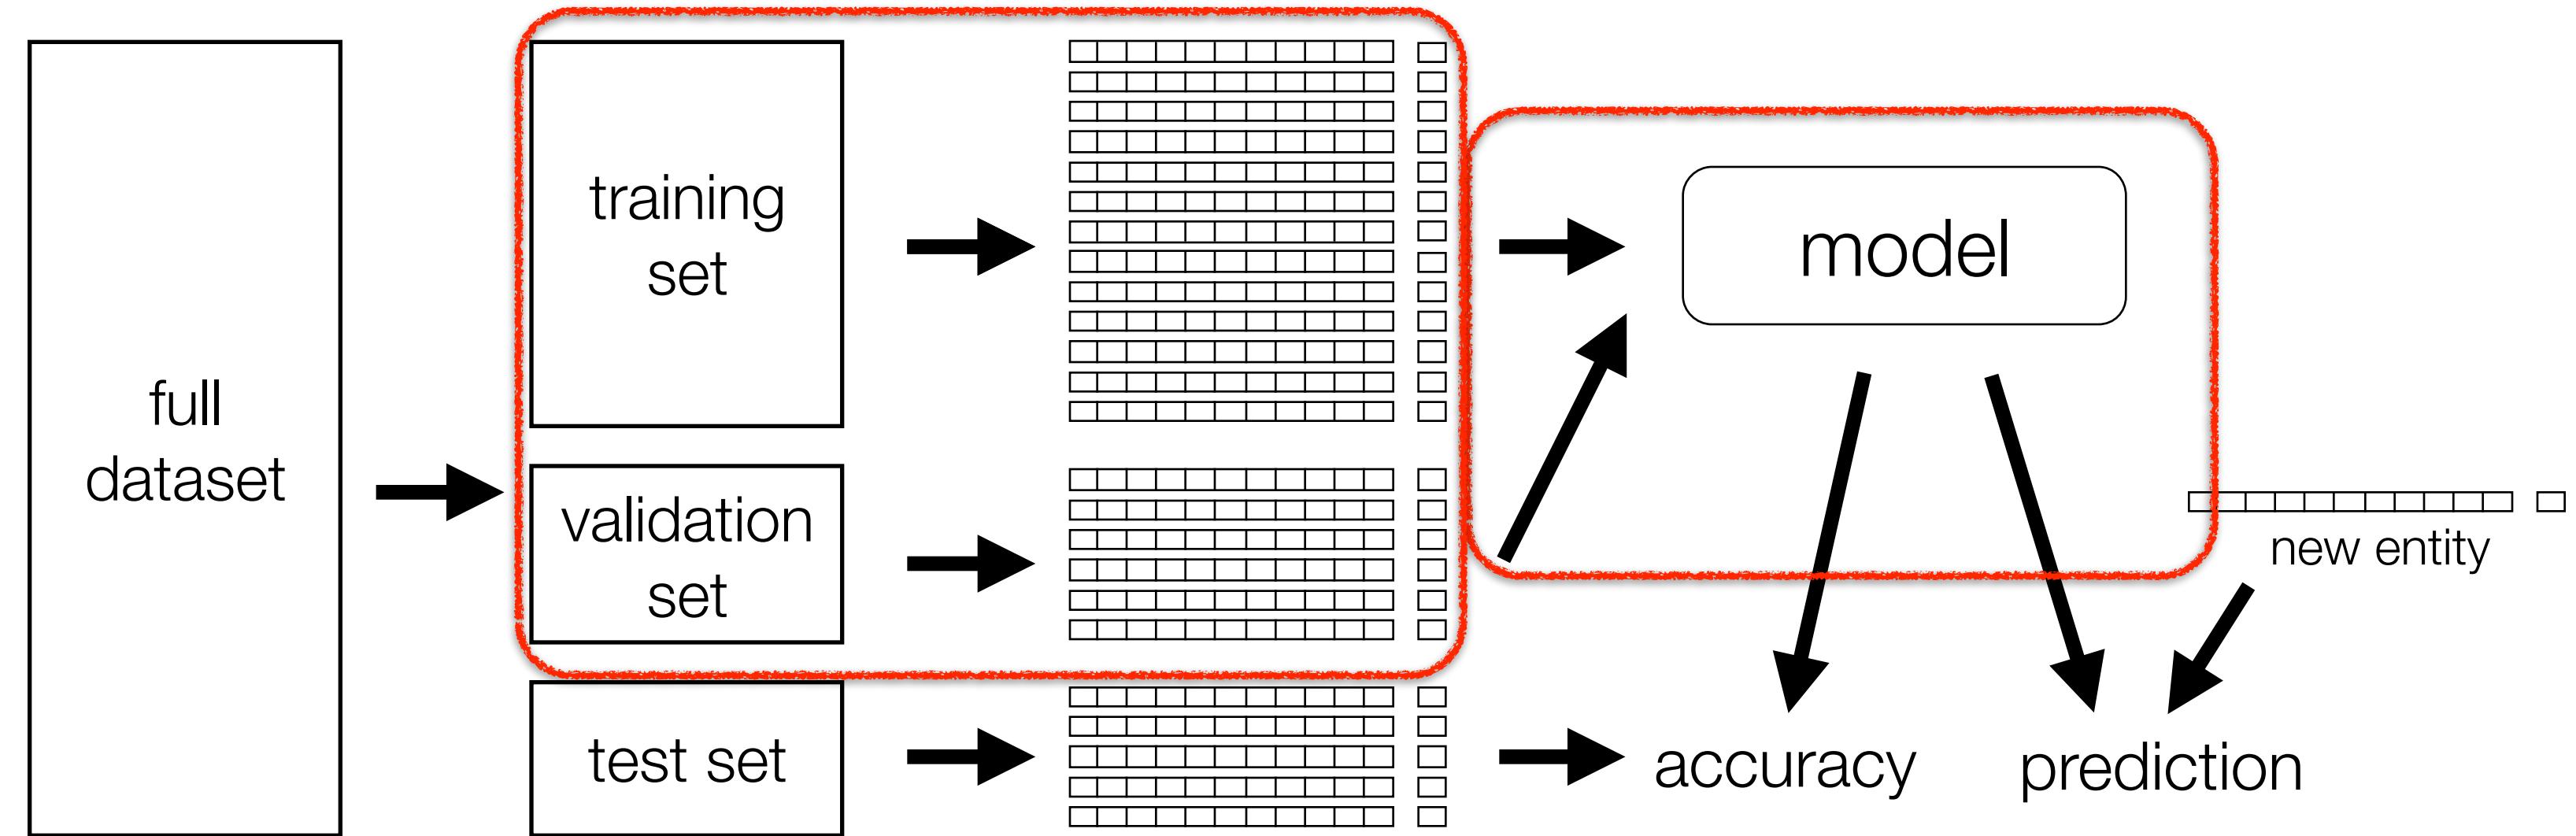
How do we choose a good value for this free parameter?

- Most methods have free parameters / ‘hyperparameters’ to tune

First thought: Search over multiple values, evaluate each on test set

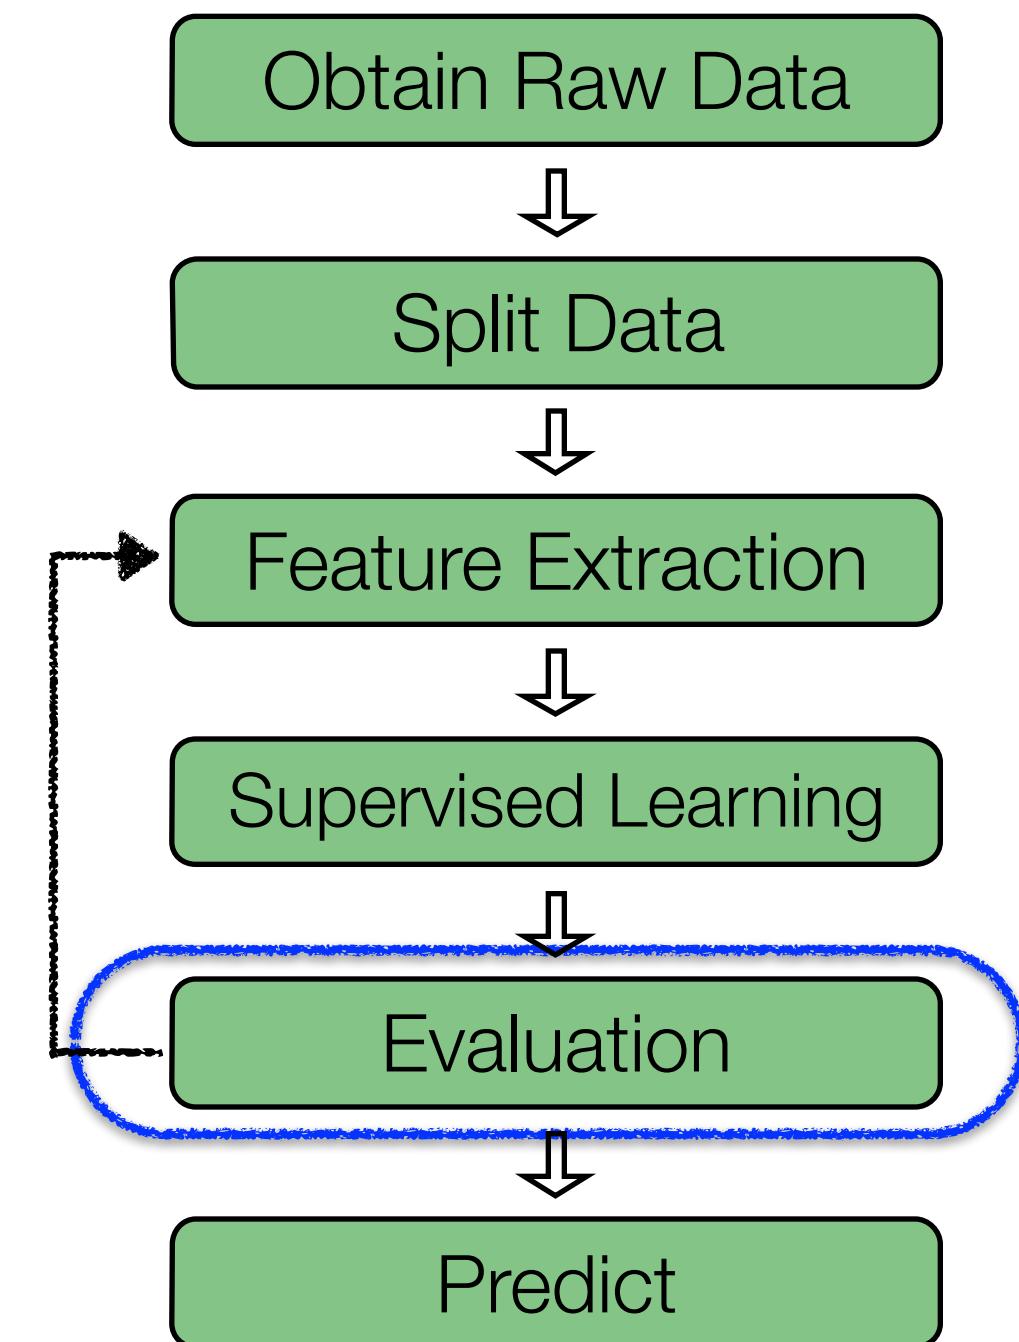
- But, goal of test set is to simulate unobserved data
- We may overfit if we use it to choose hyperparameters

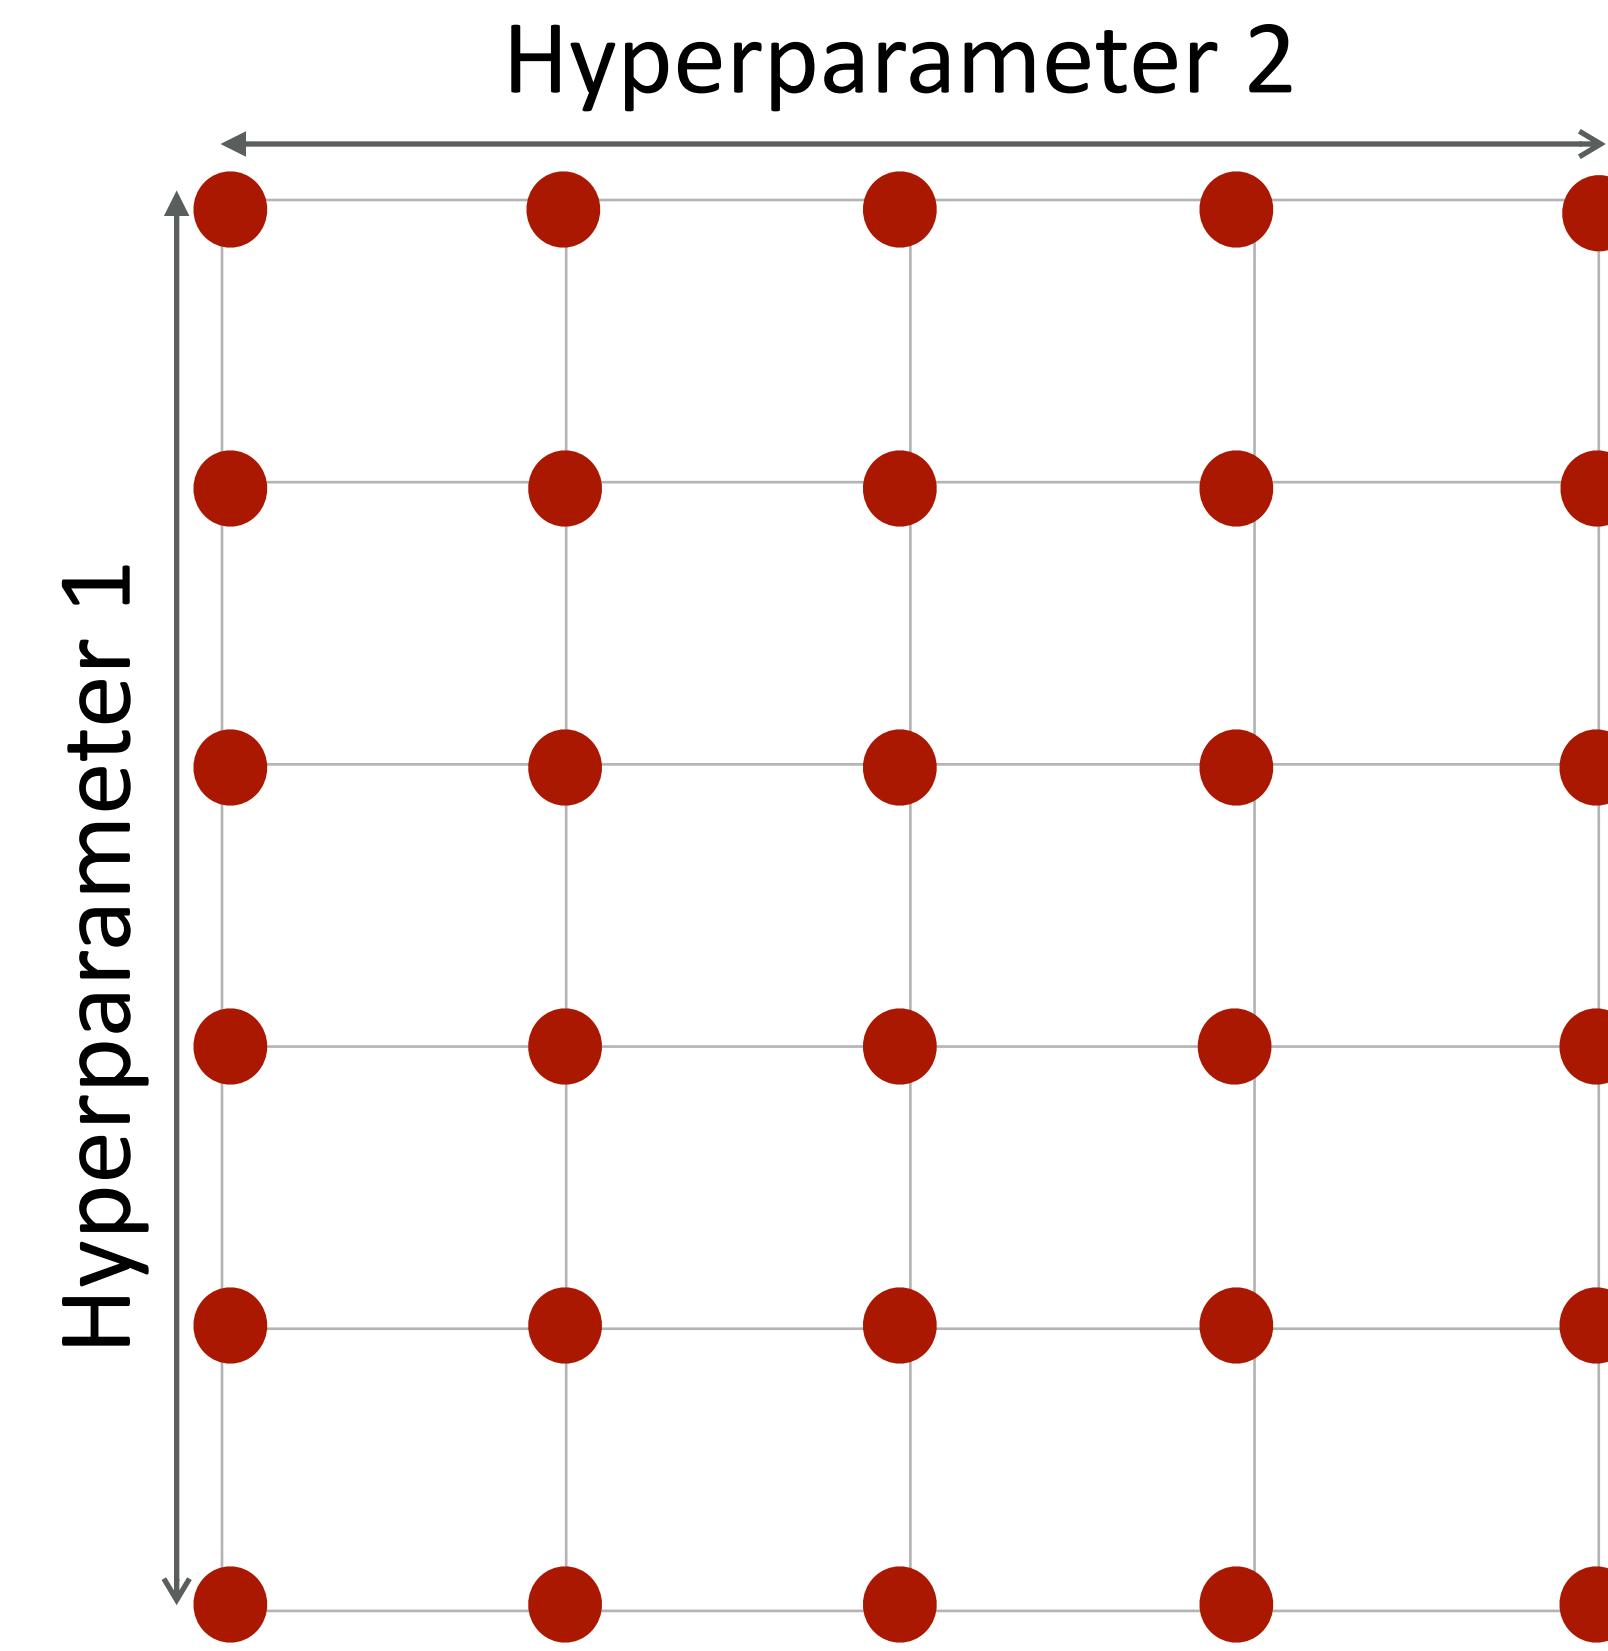
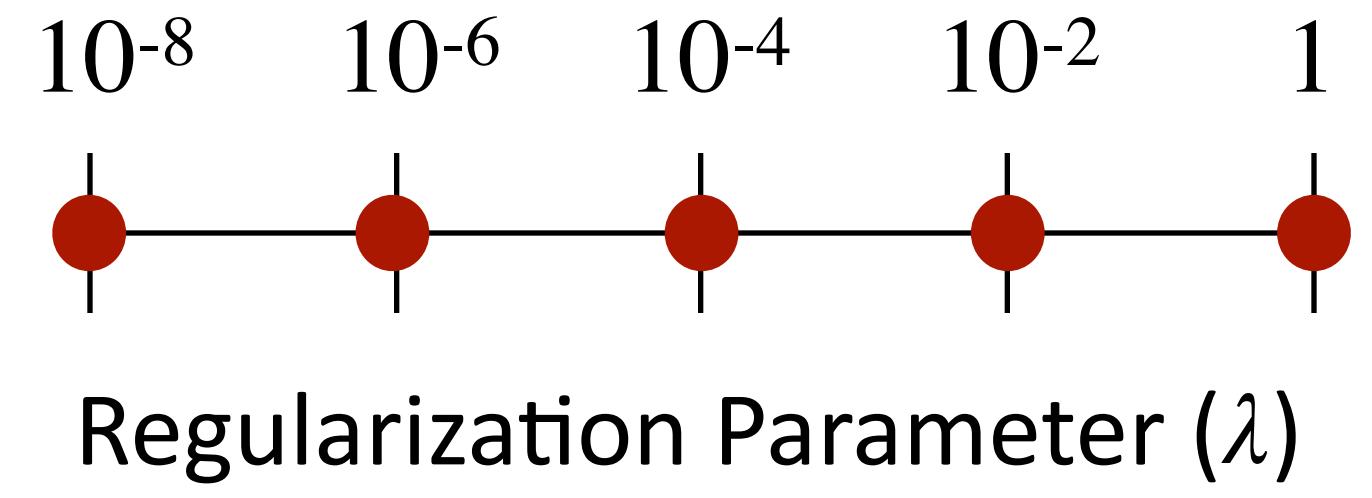
Second thought: **Create another hold out dataset for this search**



Evaluation (Part 1): Hyperparameter tuning

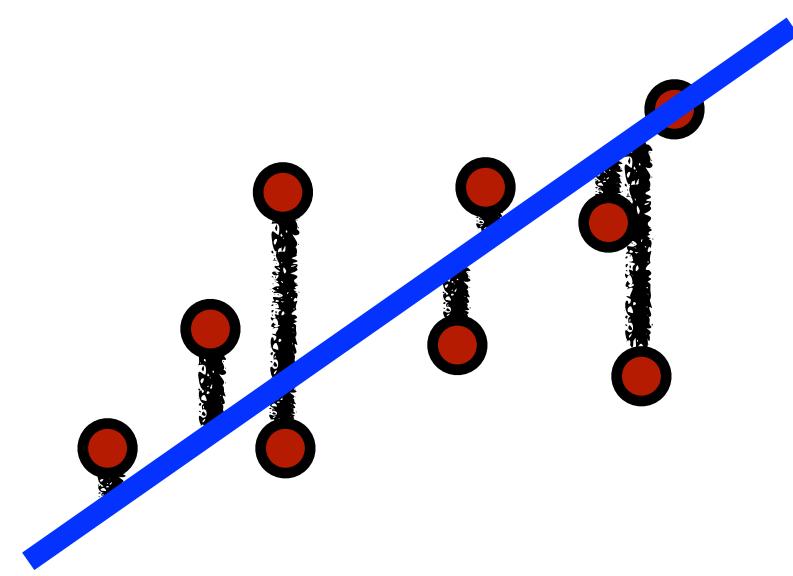
- *Training*: train various models
- *Validation*: evaluate various models (e.g., Grid Search)
- *Test*: evaluate final model's accuracy





- Grid Search:** Exhaustively search through hyperparameter space
- Define and discretize search space (linear or log scale)
 - Evaluate points via validation error

Evaluating Predictions



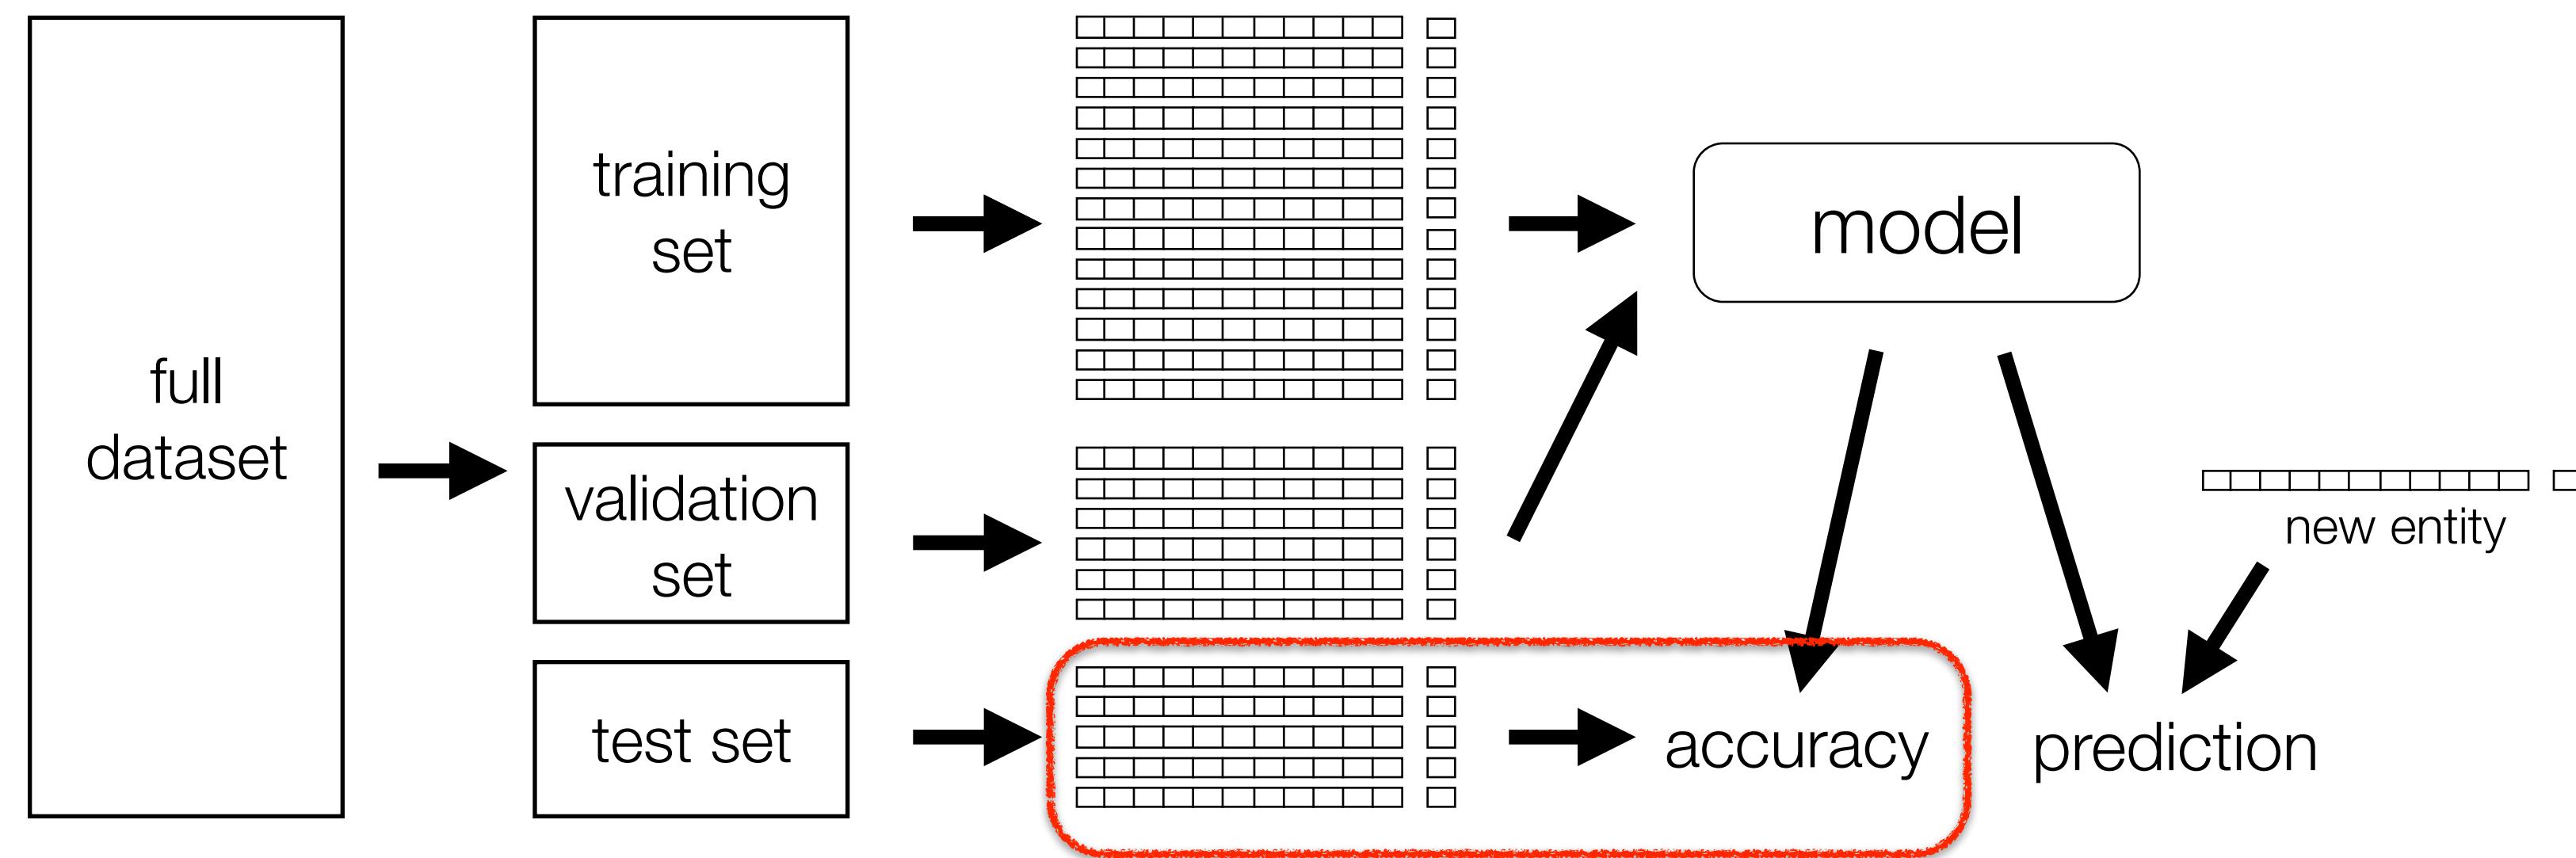
How can we compare labels and predictions for n validation points?

Least squares optimization involves squared loss, $(y - \hat{y})^2$, so it seems reasonable to use mean squared error (**MSE**):

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2$$

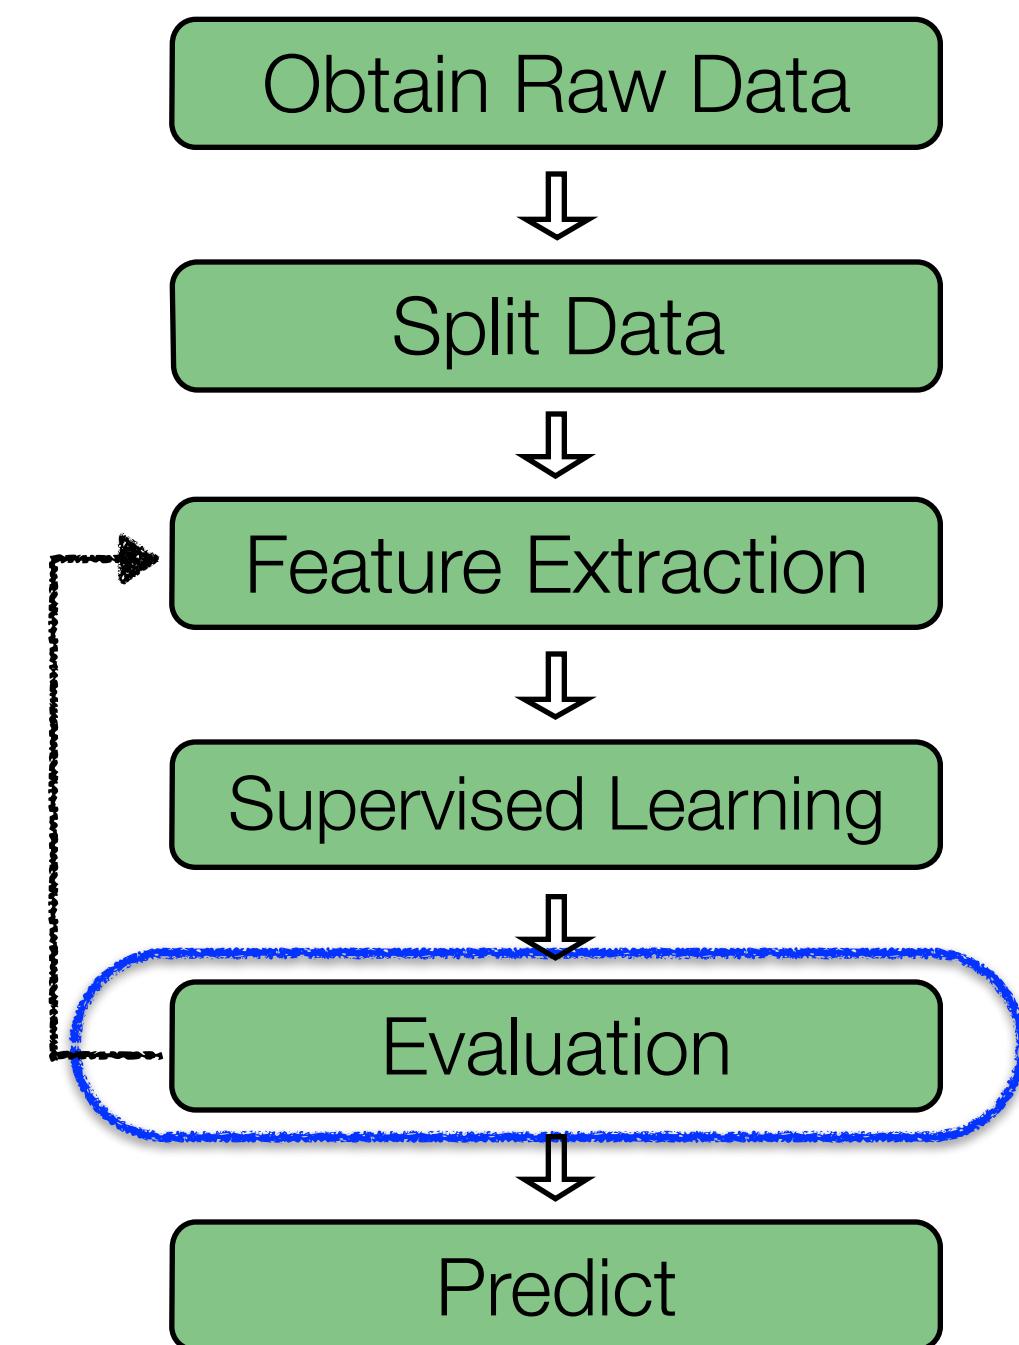
But MSE's unit of measurement is square of quantity being measured, e.g., "squared years" for song prediction

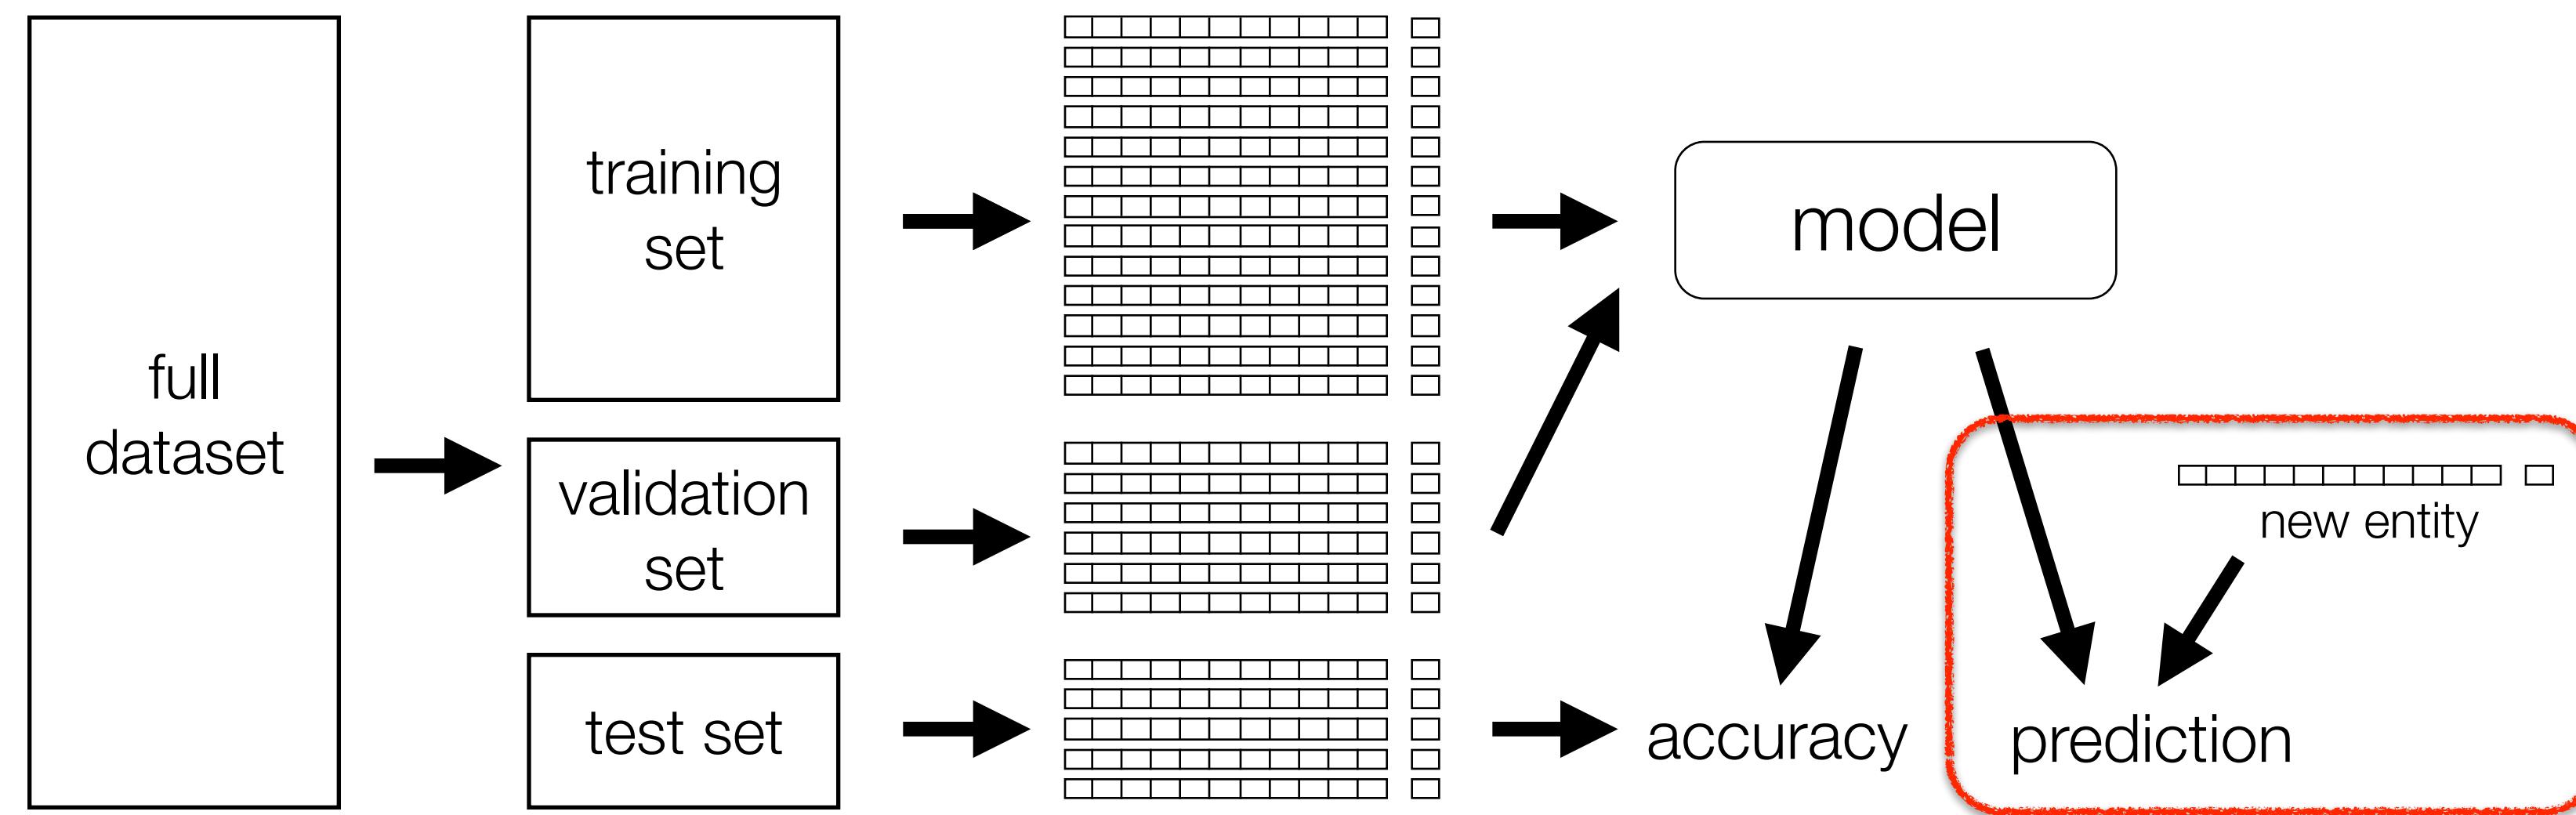
More natural to use root-mean-square error (**RMSE**), i.e., $\sqrt{\text{MSE}}$



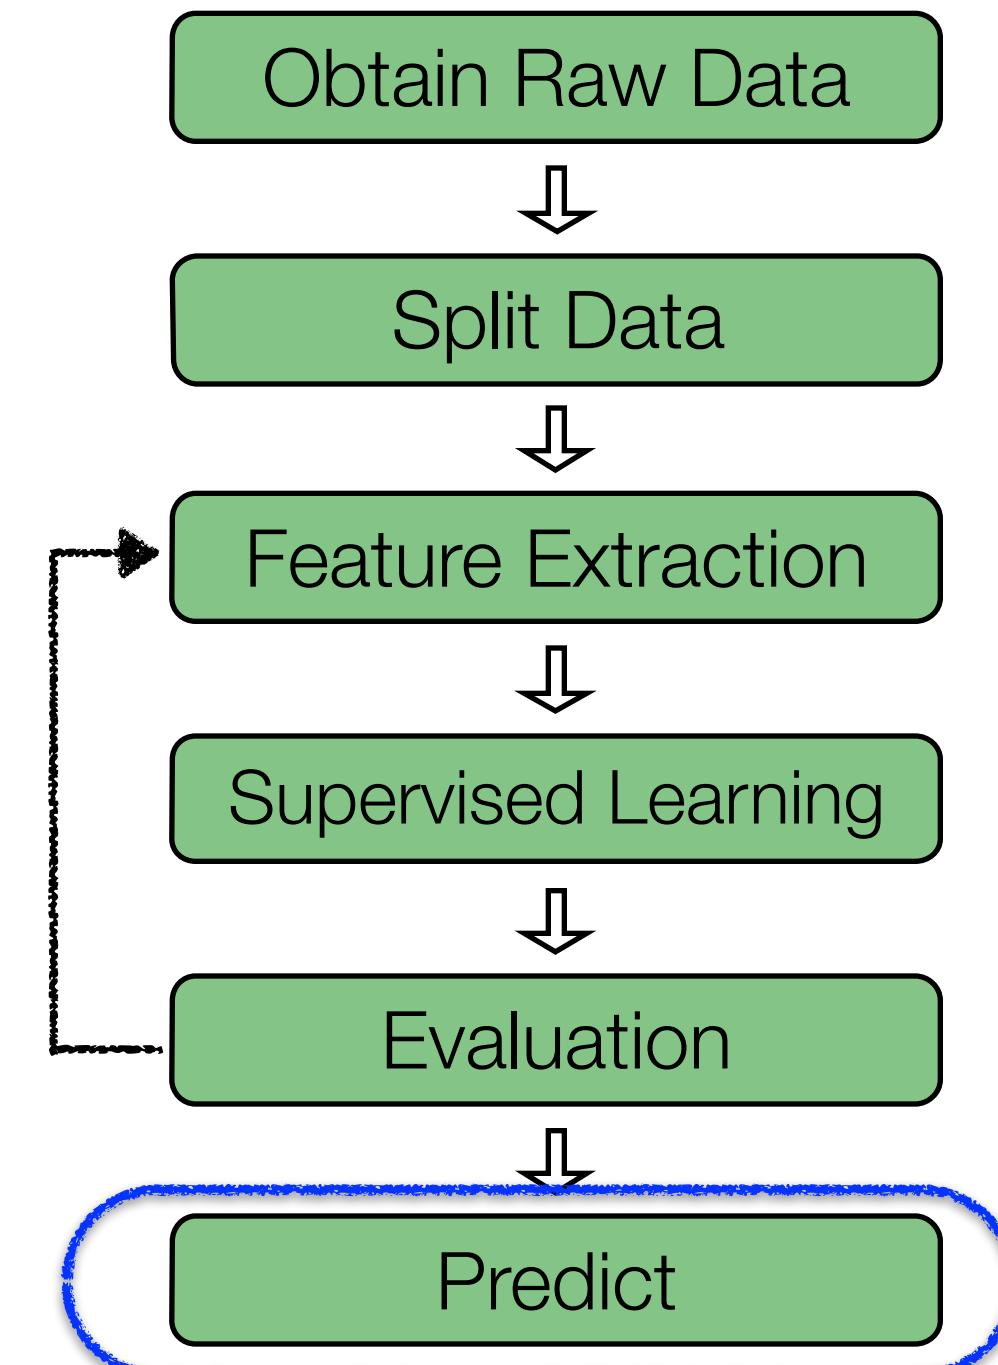
Evaluation (Part 2): Evaluate final model

- Training set: train various models
- Validation set: evaluate various models
- *Test set*: evaluate final model's accuracy





Predict: Final model can then be used to make predictions on future observations, e.g., new songs

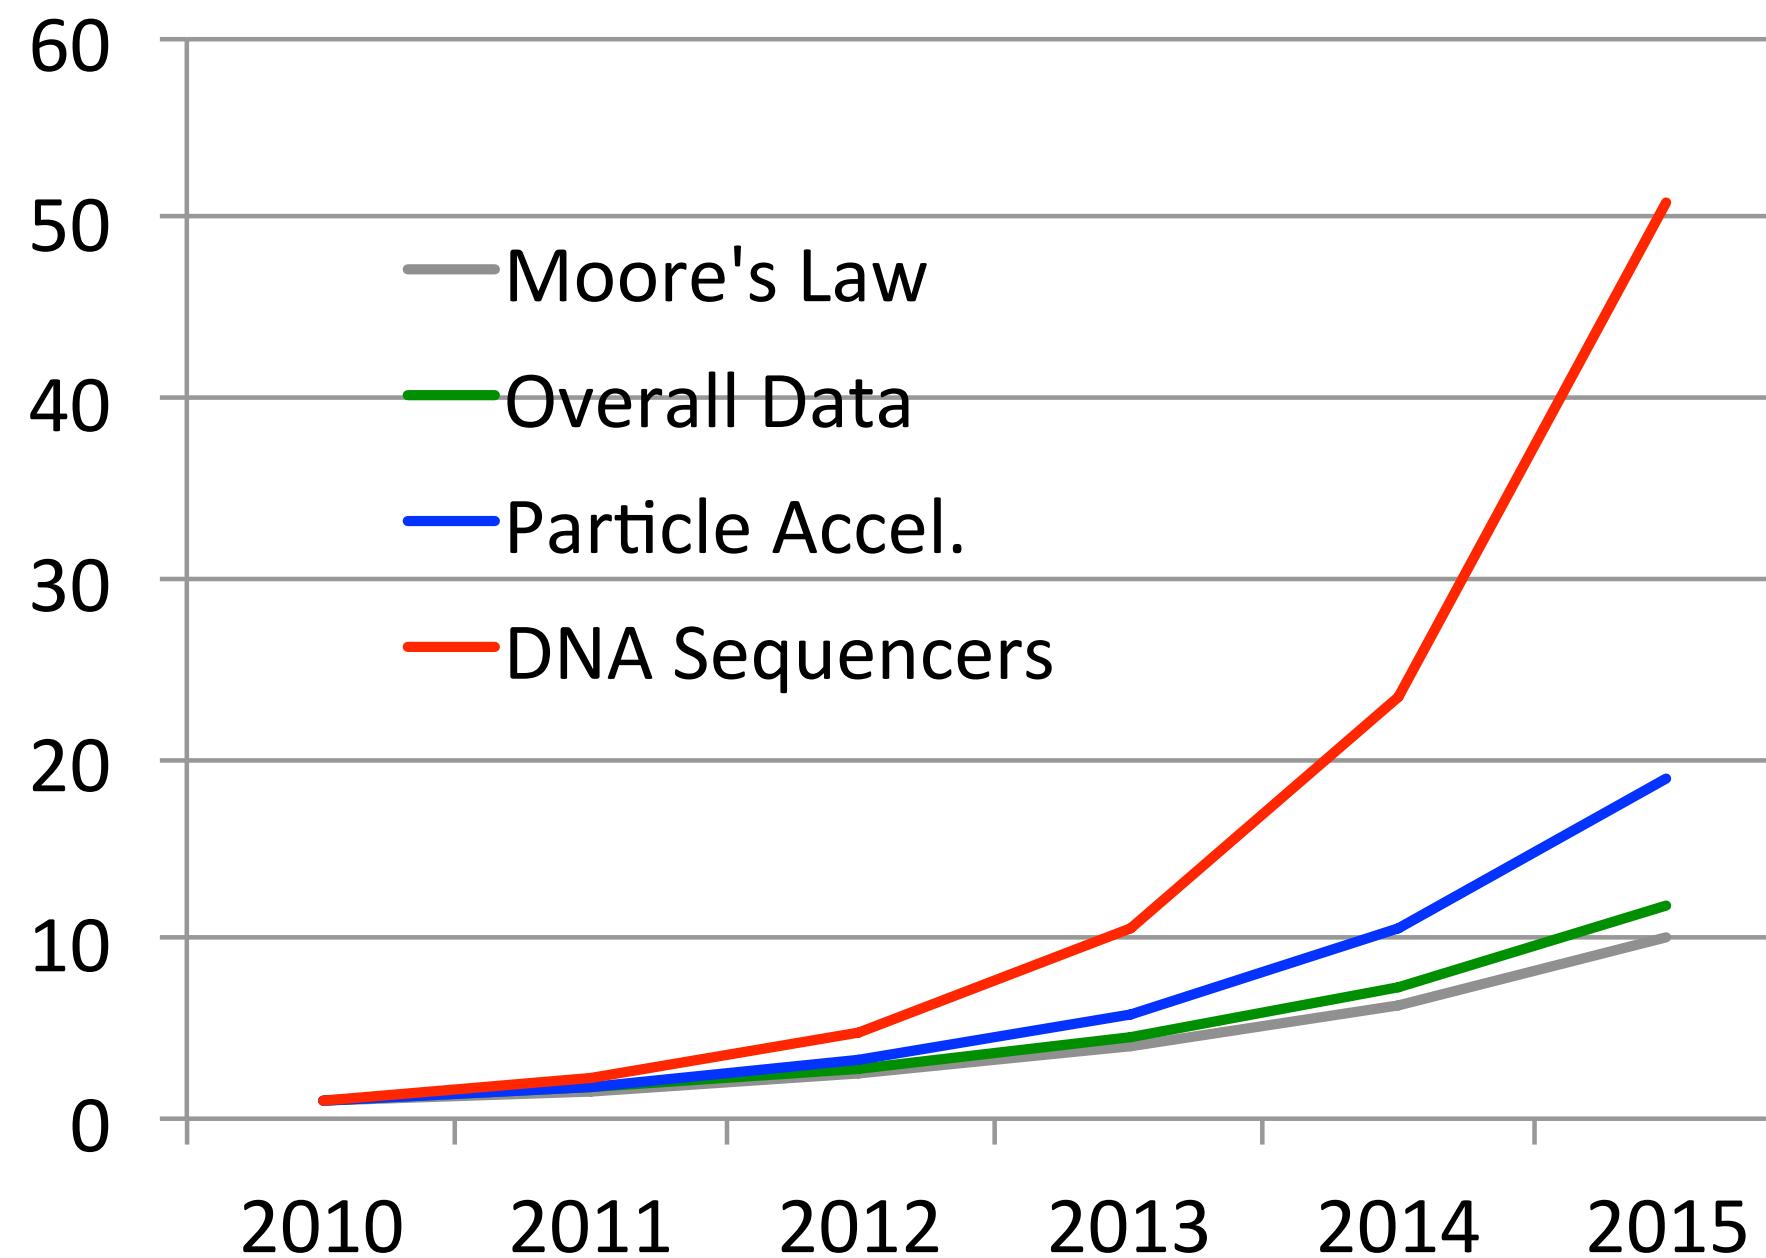


Distributed ML: Computation and Storage

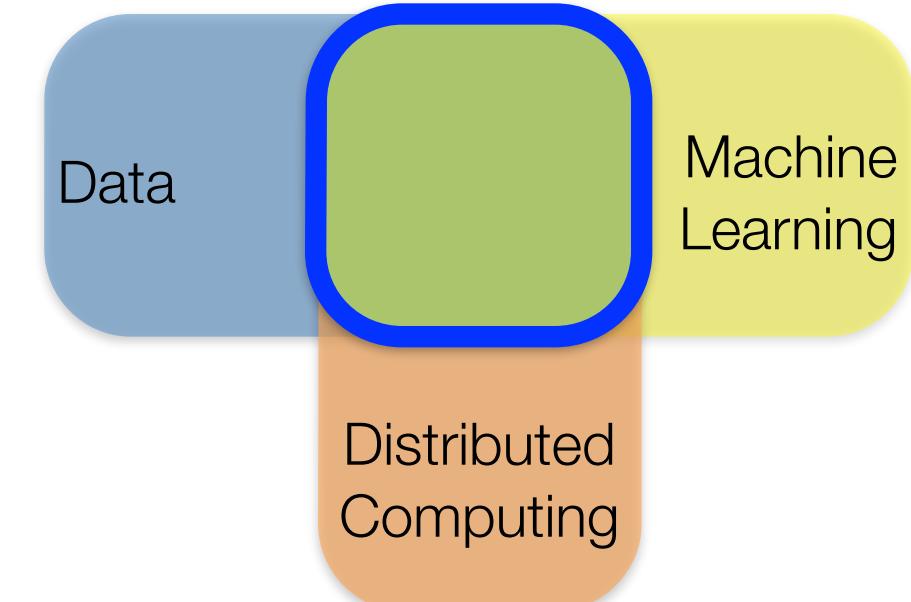


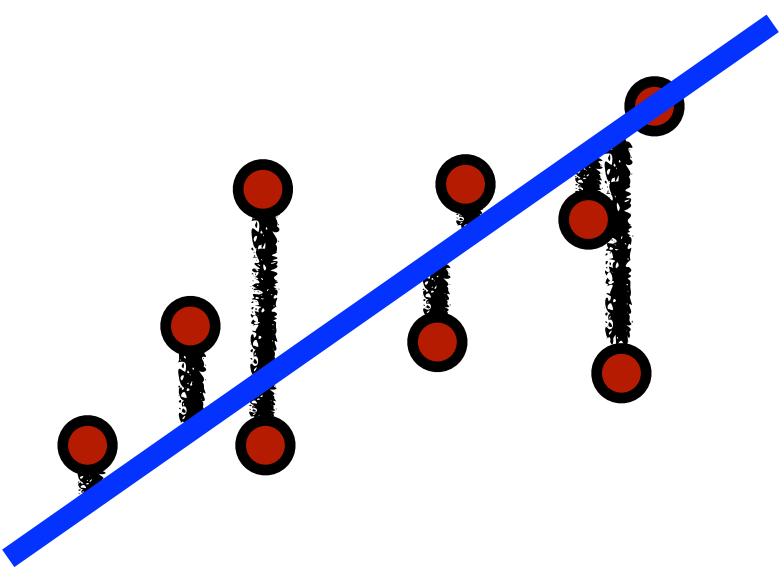
Challenge: Scalability

Classic ML techniques are not always suitable for modern datasets



Data Grows Faster
than Moore's Law
[IDC report, Kathy Yelick, LBNL]





Least Squares Regression: Learn mapping (\mathbf{w}) from features to labels that minimizes residual sum of squares:

$$\min_{\mathbf{w}} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

Closed form solution: $\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ (if inverse exists)

How do we solve this computationally?

- Computational profile similar for Ridge Regression

Computing Closed Form Solution

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Computation: $O(nd^2 + d^3)$ operations

Consider number of arithmetic operations (+, -, ×, /)

Computational bottlenecks:

- Matrix multiply of $\mathbf{X}^\top \mathbf{X}$: $O(nd^2)$ operations
- Matrix inverse: $O(d^3)$ operations

Other methods (Cholesky, QR, SVD) have same complexity

Storage Requirements

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Computation: $O(nd^2 + d^3)$ operations

Storage: $O(nd + d^2)$ floats

Consider storing values as floats (8 bytes)

Storage bottlenecks:

- $\mathbf{X}^\top \mathbf{X}$ and its inverse: $O(d^2)$ floats
- \mathbf{X} : $O(nd)$ floats

Big n and Small d

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Computation: $O(\underline{nd^2} + d^3)$ operations

Storage: $O(\underline{nd} + d^2)$ floats

Assume $O(d^3)$ computation and $O(d^2)$ storage feasible on single machine

Storing \mathbf{X} and computing $\mathbf{X}^\top \mathbf{X}$ are the bottlenecks

Can distribute storage and computation!

- Store data points (rows of \mathbf{X}) across machines
- Compute $\mathbf{X}^\top \mathbf{X}$ as a sum of outer products

Matrix Multiplication via Inner Products

Each entry of output matrix is result of inner product of inputs matrices

$$\begin{bmatrix} 9 & 3 & 5 \\ 4 & 1 & 2 \end{bmatrix} \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 28 \end{bmatrix}$$

$$9 \times 1 + 3 \times 3 + 5 \times 2 = 28$$

Matrix Multiplication via Inner Products

Each entry of output matrix is result of inner product of inputs matrices

$$\begin{bmatrix} 9 & 3 & 5 \\ 4 & 1 & 2 \end{bmatrix} \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 28 & 18 \end{bmatrix}$$

Matrix Multiplication via Inner Products

Each entry of output matrix is result of inner product of inputs matrices

$$\begin{bmatrix} 9 & 3 & 5 \\ 4 & 1 & 2 \end{bmatrix} \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 28 & 18 \\ 11 & 9 \end{bmatrix}$$

The matrix multiplication is shown. The first matrix has columns $\begin{bmatrix} 9 \\ 4 \end{bmatrix}$, $\begin{bmatrix} 3 \\ 1 \end{bmatrix}$, and $\begin{bmatrix} 5 \\ 2 \end{bmatrix}$. The second matrix has rows $\begin{bmatrix} 1 \end{bmatrix}$, $\begin{bmatrix} 3 \end{bmatrix}$, and $\begin{bmatrix} 2 \\ -5 \\ 3 \end{bmatrix}$. The resulting matrix is $\begin{bmatrix} 28 & 18 \\ 11 & 9 \end{bmatrix}$.

Matrix Multiplication via Outer Products

Output matrix is **sum of outer products** between corresponding rows and columns of input matrices

$$\begin{bmatrix} 9 \\ 4 \end{bmatrix} \begin{bmatrix} 3 & 5 \\ 1 & 2 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & -5 \\ 2 & 3 \end{bmatrix} = \begin{bmatrix} \quad & \quad \\ \quad & \quad \end{bmatrix}$$

$$\begin{bmatrix} 9 & 18 \\ 4 & 8 \end{bmatrix}$$

Matrix Multiplication via Outer Products

Output matrix is **sum of outer products** between corresponding rows and columns of input matrices

$$\begin{bmatrix} 9 & \boxed{3} & 5 \\ 4 & \boxed{1} & 2 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & -5 \\ 2 & 3 \end{bmatrix} = \begin{bmatrix} & \\ & \end{bmatrix}$$

$$\begin{bmatrix} 9 & 18 \\ 4 & 8 \end{bmatrix} + \begin{bmatrix} 9 & -15 \\ 3 & -5 \end{bmatrix}$$

Matrix Multiplication via Outer Products

Output matrix is **sum of outer products** between corresponding rows and columns of input matrices

$$\begin{bmatrix} 9 & 3 & \boxed{5} \\ 4 & 1 & \boxed{2} \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & -5 \\ \boxed{2} & 3 \end{bmatrix} = \begin{bmatrix} & \\ & \end{bmatrix}$$

$$\begin{bmatrix} 9 & 18 \\ 4 & 8 \end{bmatrix} + \begin{bmatrix} 9 & -15 \\ 3 & -5 \end{bmatrix} + \begin{bmatrix} 10 & 15 \\ 4 & 6 \end{bmatrix}$$

Matrix Multiplication via Outer Products

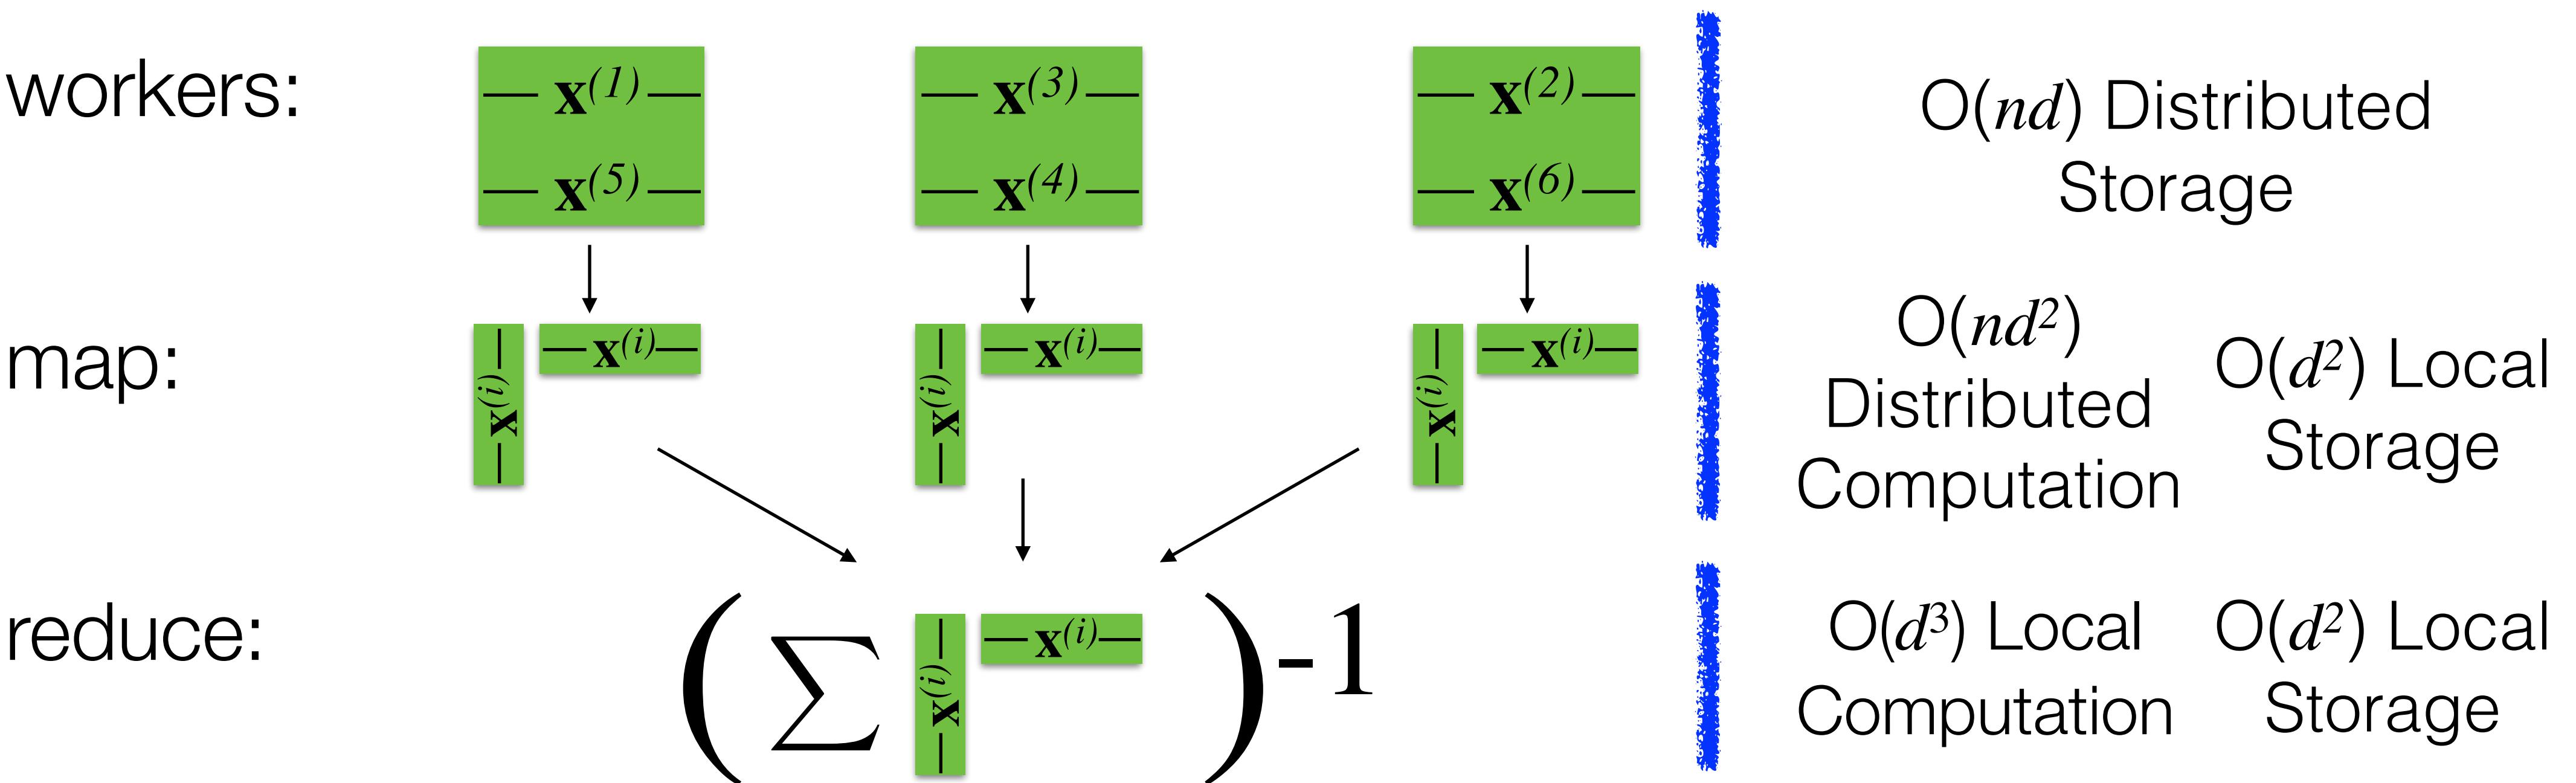
Output matrix is **sum of outer products** between corresponding rows and columns of input matrices

$$\begin{bmatrix} 9 & 3 & \boxed{5} \\ 4 & 1 & \boxed{2} \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & -5 \\ \boxed{2} & \boxed{3} \end{bmatrix} = \begin{bmatrix} 28 & 18 \\ 11 & 9 \end{bmatrix}$$

$$\begin{bmatrix} 9 & 18 \\ 4 & 8 \end{bmatrix} + \begin{bmatrix} 9 & -15 \\ 3 & -5 \end{bmatrix} + \begin{bmatrix} 10 & 15 \\ 4 & 6 \end{bmatrix}$$

$$\mathbf{X}^\top \mathbf{X} = \begin{matrix} & n \\ d & \left| \begin{array}{c} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(2)} \\ \cdots \\ \mathbf{x}^{(n)} \end{array} \right| \end{matrix} = \sum_{i=1}^n \begin{matrix} & n \\ d & \left| \begin{array}{c} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(2)} \\ \vdots \\ \mathbf{x}^{(i)} \\ \vdots \\ \mathbf{x}^{(n)} \end{array} \right| \end{matrix}$$

Example: $n = 6$; 3 workers



> `trainData.map(computeOuterProduct)
.reduce(sumAndInvert)`

workers:

$$\begin{array}{|c|} \hline - \mathbf{x}^{(1)} - \\ \hline - \mathbf{x}^{(5)} - \\ \hline \end{array}$$

$$\begin{array}{|c|} \hline - \mathbf{x}^{(3)} - \\ \hline - \mathbf{x}^{(4)} - \\ \hline \end{array}$$

$$\begin{array}{|c|} \hline - \mathbf{x}^{(2)} - \\ \hline - \mathbf{x}^{(6)} - \\ \hline \end{array}$$

map:

$$\begin{array}{|c|c|} \hline - \mathbf{x}^{(i)} - & - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - & \\ \hline \end{array}$$

$$\begin{array}{|c|c|} \hline - \mathbf{x}^{(i)} - & - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - & \\ \hline \end{array}$$

$$\begin{array}{|c|c|} \hline - \mathbf{x}^{(i)} - & - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - & \\ \hline \end{array}$$

reduce:

$$\left(\sum \begin{array}{|c|c|} \hline - \mathbf{x}^{(i)} - & - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - & \\ \hline \end{array} \right) - 1$$

$O(nd)$ Distributed Storage

$O(nd^2)$ Distributed Computation $O(d^2)$ Local Storage

$O(d^3)$ Local Computation $O(d^2)$ Local Storage

Distributed ML: Computation and Storage, Part II



Big n and Small d

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Computation: $O(\underline{nd^2} + d^3)$ operations

Storage: $O(\underline{nd} + d^2)$ floats

Assume $O(d^3)$ computation and $O(d^2)$ storage feasible on single machine

Can distribute storage and computation!

- Store data points (rows of \mathbf{X}) across machines
- Compute $\mathbf{X}^\top \mathbf{X}$ as a sum of outer products

Big n and Small d

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Computation: $O(\cancel{nd^2} + d^3)$ operations

Storage: $O(\cancel{nd} + d^2)$ floats

```
> trainData.map(computeOuterProduct)
    .reduce(sumAndInvert)
```

Big n and Big d

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Computation: $O(\underline{nd^2} + \underline{d^3})$ operations

Storage: $O(\underline{nd} + \underline{d^2})$ floats

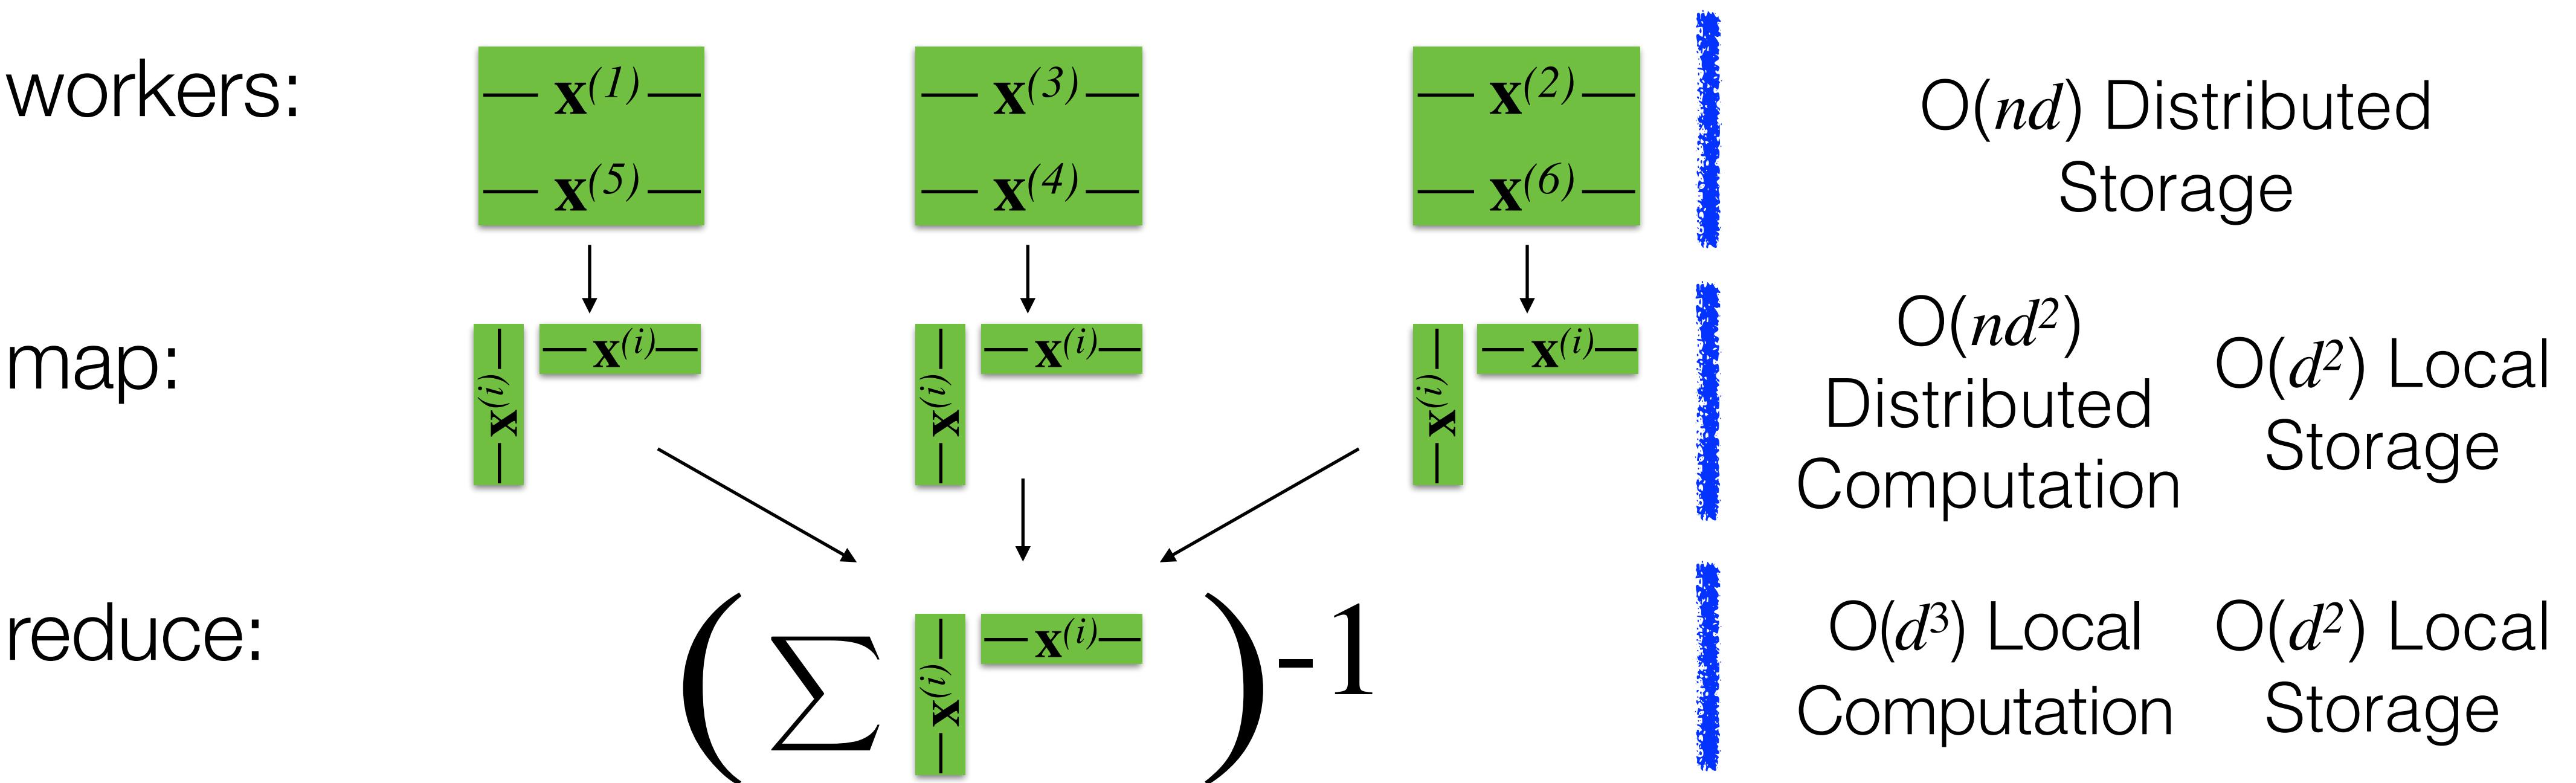
As before, storing \mathbf{X} and computing $\mathbf{X}^\top \mathbf{X}$ are bottlenecks

Now, storing and operating on $\mathbf{X}^\top \mathbf{X}$ is also a bottleneck

- Can't easily distribute!

$$\mathbf{X}^\top \mathbf{X} = \begin{matrix} & n \\ d & \left| \begin{array}{c} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(2)} \\ \cdots \\ \mathbf{x}^{(n)} \end{array} \right| \end{matrix} = \sum_{i=1}^n \begin{matrix} & n \\ d & \left| \begin{array}{c} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(2)} \\ \vdots \\ \mathbf{x}^{(i)} \\ \vdots \\ \mathbf{x}^{(n)} \end{array} \right| \end{matrix}$$

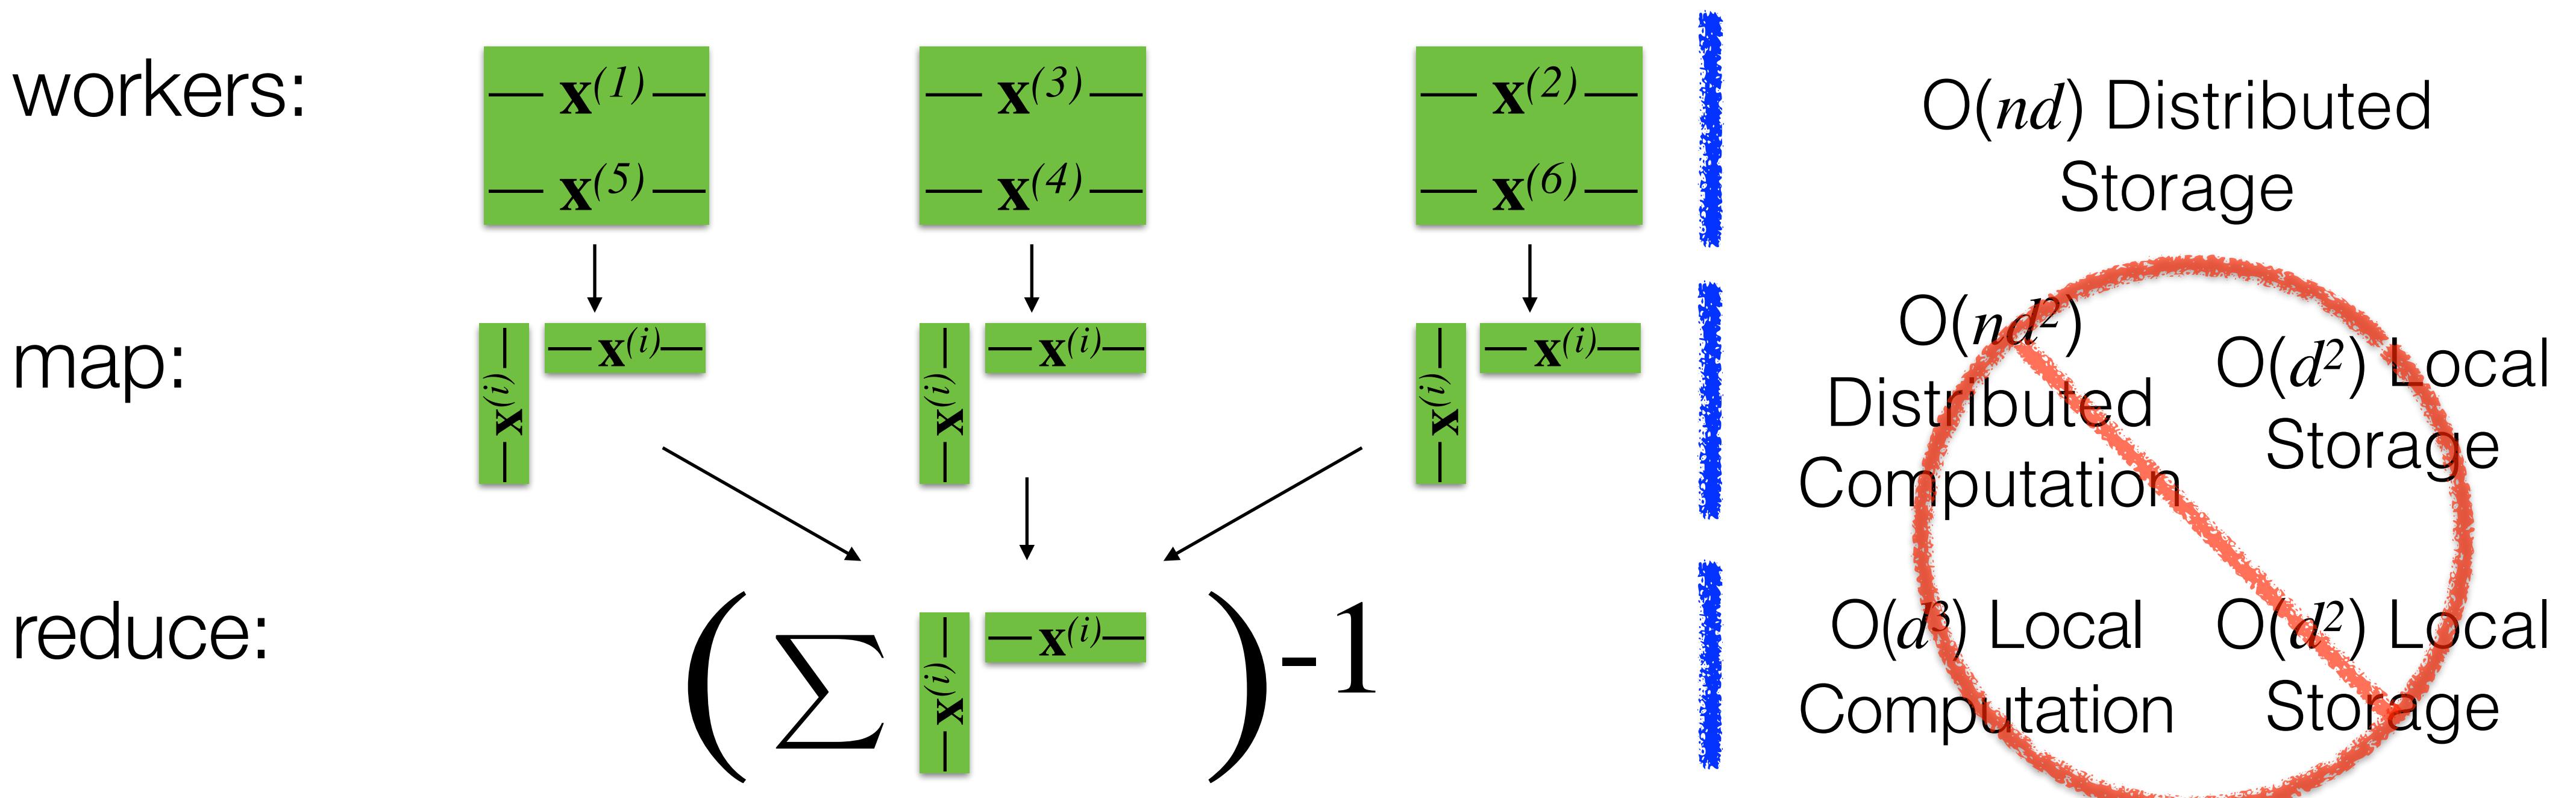
Example: $n = 6$; 3 workers



$$\mathbf{X}^\top \mathbf{X} =$$

$$= \sum_{i=1}^n \mathbf{x}^{(i)} \mathbf{x}^{(i)\top}$$

Example: $n = 6$; 3 workers



Big n and Big d

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Computation: $O(nd^2 + d^3)$ operations

Storage: $O(nd + d^2)$ floats

As before, storing \mathbf{X} and computing $\mathbf{X}^\top \mathbf{X}$ are bottlenecks

Now, storing and operating on $\mathbf{X}^\top \mathbf{X}$ is also a bottleneck

- Can't easily distribute!

1st Rule of thumb

Computation and storage should be linear (in n, d)

Big n and Big d

We need methods that are linear in time and space

One idea: **Exploit sparsity**

- Explicit sparsity can provide orders of magnitude storage and computational gains

Sparse data is prevalent

- Text processing: bag-of-words, n-grams
- Collaborative filtering: ratings matrix
- Graphs: adjacency matrix
- Categorical features: one-hot-encoding
- Genomics: SNPs, variant calling

dense : $\underline{1.} \underline{0.} \underline{0.} \underline{0.} \underline{0.} \underline{0.} \underline{3.}$

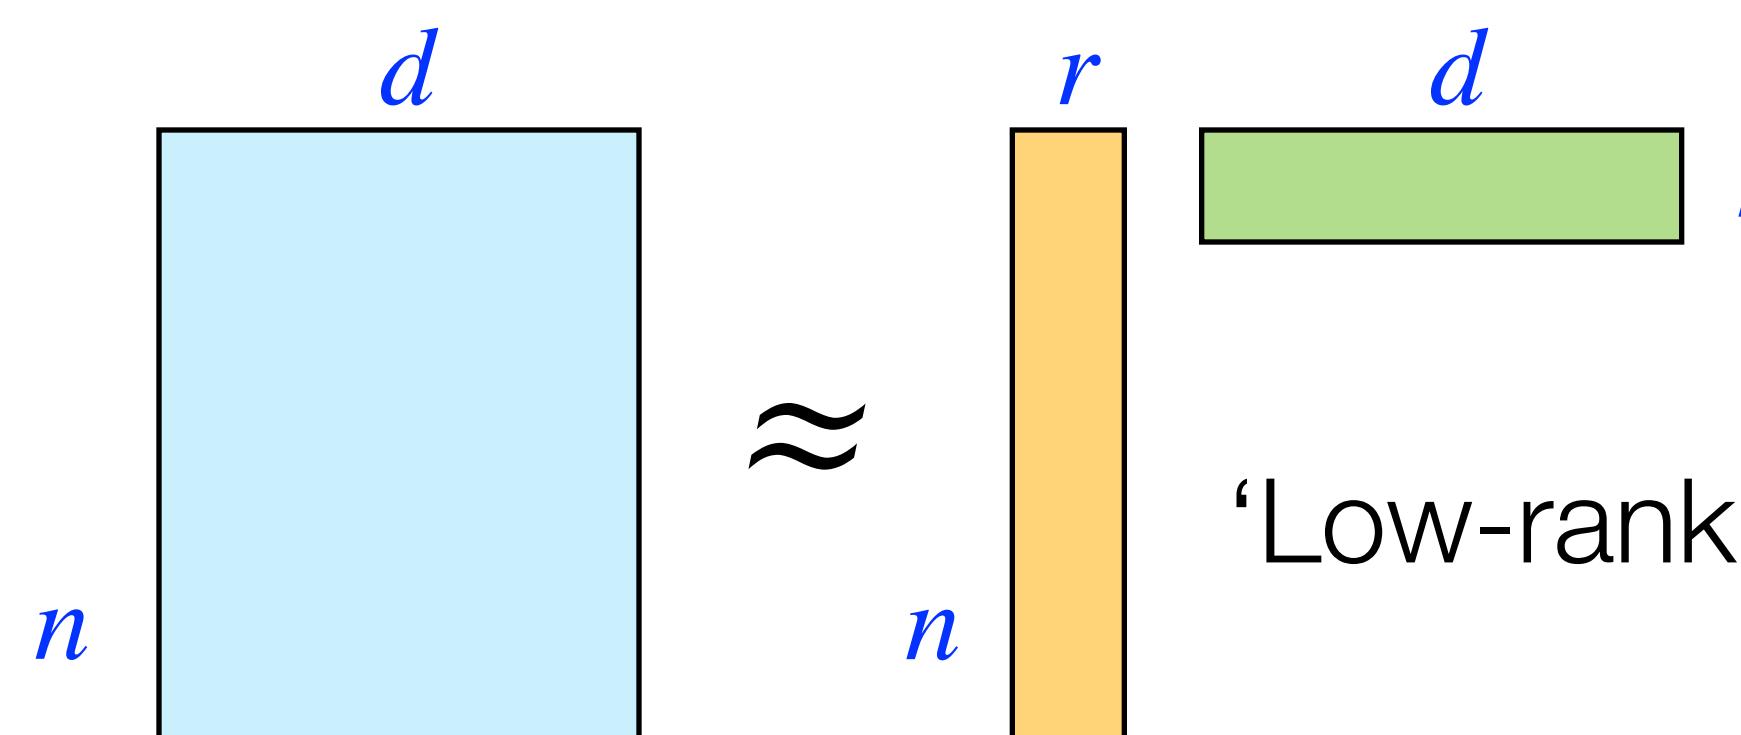
sparse : $\begin{cases} \text{size : 7} \\ \text{indices : } \underline{0} \underline{6} \\ \text{values : } \underline{1.} \underline{3.} \end{cases}$

Big n and Big d

We need methods that are linear in time and space

One idea: **Exploit sparsity**

- Explicit sparsity can provide orders of magnitude storage and computational gains
- Latent sparsity assumption can be used to reduce dimension, e.g., PCA, low-rank approximation (unsupervised learning)



Big n and Big d

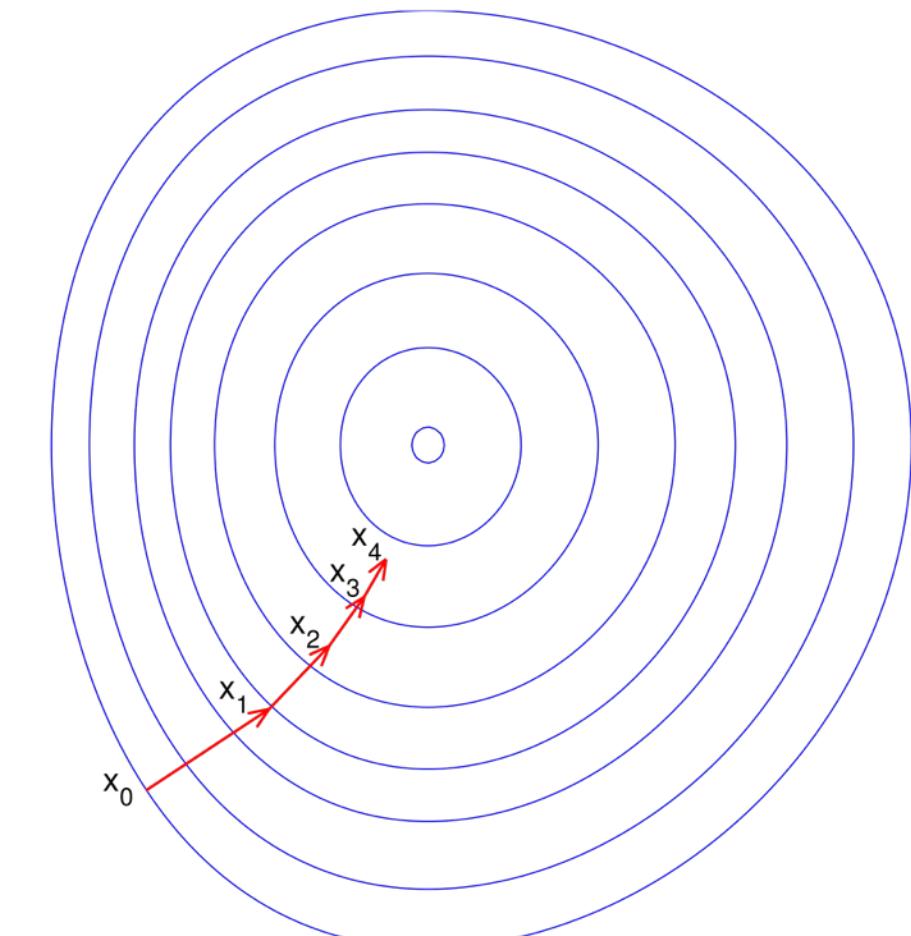
We need methods that are linear in time and space

One idea: **Exploit sparsity**

- Explicit sparsity can provide orders of magnitude storage and computational gains
- Latent sparsity assumption can be used to reduce dimension, e.g., PCA, low-rank approximation (unsupervised learning)

Another idea: **Use different algorithms**

- Gradient descent is an iterative algorithm that requires $O(nd)$ computation and $O(d)$ local storage per iteration



Closed Form Solution for Big n and Big d

Example: $n = 6$; 3 workers

workers:

$$\begin{array}{|c|} \hline - \mathbf{x}^{(1)} - \\ \hline - \mathbf{x}^{(5)} - \\ \hline \end{array}$$

$$\begin{array}{|c|} \hline - \mathbf{x}^{(3)} - \\ \hline - \mathbf{x}^{(4)} - \\ \hline \end{array}$$

$$\begin{array}{|c|} \hline - \mathbf{x}^{(2)} - \\ \hline - \mathbf{x}^{(6)} - \\ \hline \end{array}$$

$O(nd)$ Distributed Storage

map:

$$\begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline \end{array}$$

$$\begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline \end{array}$$

$$\begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline \end{array}$$

$O(nd^2)$ Distributed Computation

$O(d^2)$ Local Storage

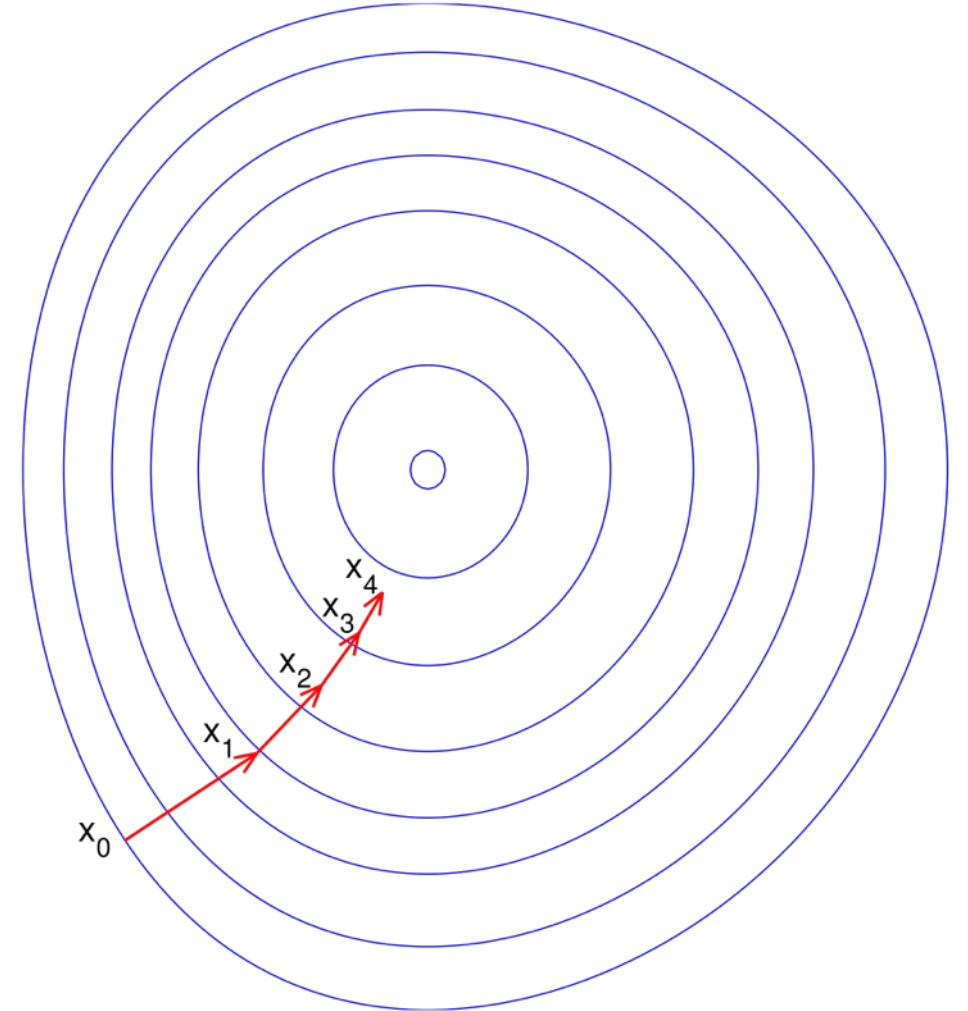
reduce:

$$\left(\sum \begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline \end{array} \right) - 1$$

$O(d^3)$ Local Computation

$O(d^2)$ Local Storage

Gradient Descent for Big n and Big d



Example: $n = 6$; 3 workers

workers:

$\begin{array}{|c|} \hline - \mathbf{x}^{(1)} - \\ \hline - \mathbf{x}^{(5)} - \\ \hline \end{array}$

$\begin{array}{|c|} \hline - \mathbf{x}^{(3)} - \\ \hline - \mathbf{x}^{(4)} - \\ \hline \end{array}$

$\begin{array}{|c|} \hline - \mathbf{x}^{(2)} - \\ \hline - \mathbf{x}^{(6)} - \\ \hline \end{array}$

map:

$\begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - \\ \hline \end{array}$

$\begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - \\ \hline \end{array}$

$\begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - \\ \hline \end{array}$

reduce:

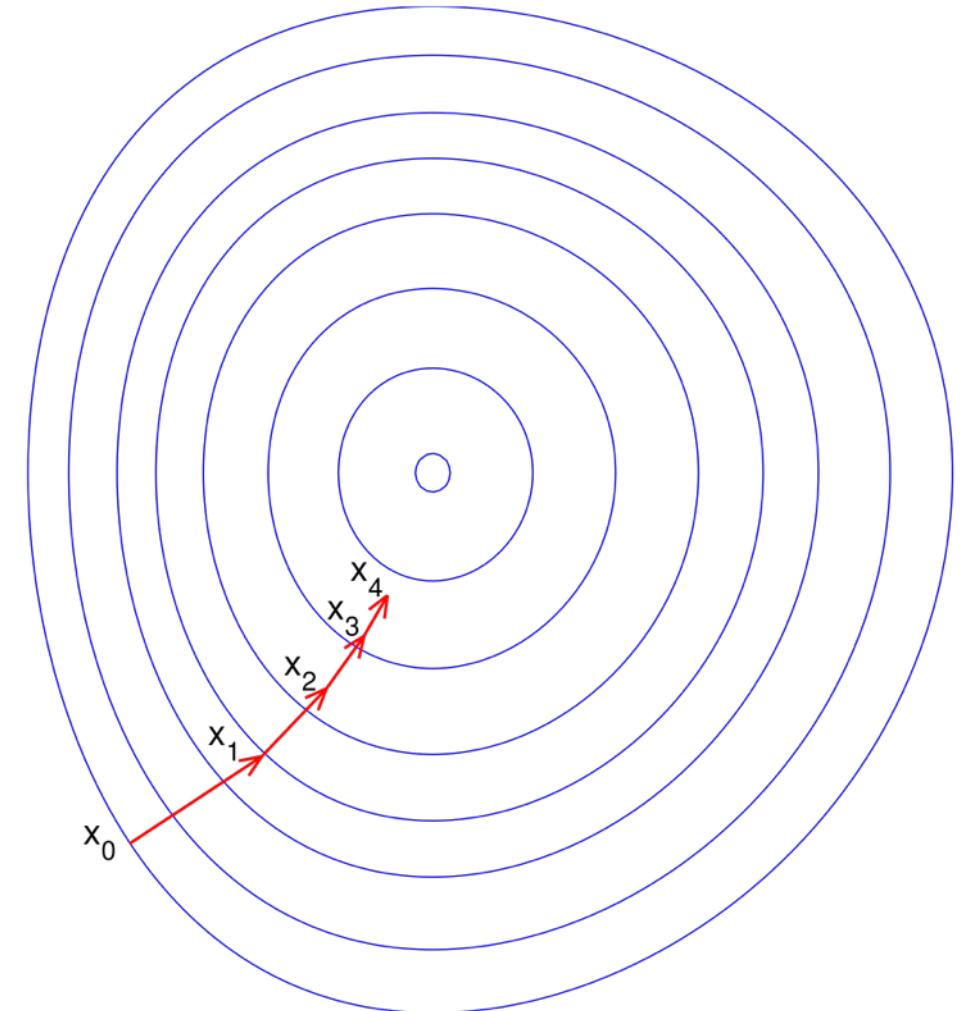
$$\left(\sum \begin{array}{|c|} \hline - \mathbf{x}^{(i)} - \\ \hline - \mathbf{x}^{(i)} - \\ \hline \end{array} \right) - 1$$

$O(nd)$ Distributed Storage

$O(nd^2)$ Distributed Computation $O(d^2)$ Local Storage

$O(d^3)$ Local Computation $O(d^2)$ Local Storage

Gradient Descent for Big n and Big d



Example: $n = 6$; 3 workers

workers:

$$\begin{array}{c} \text{--- } \mathbf{x}^{(1)} \text{ ---} \\ \text{--- } \mathbf{x}^{(5)} \text{ ---} \end{array}$$

$$\begin{array}{c} \text{--- } \mathbf{x}^{(3)} \text{ ---} \\ \text{--- } \mathbf{x}^{(4)} \text{ ---} \end{array}$$

$$\begin{array}{c} \text{--- } \mathbf{x}^{(2)} \text{ ---} \\ \text{--- } \mathbf{x}^{(6)} \text{ ---} \end{array}$$

map:

?

?

?

reduce:

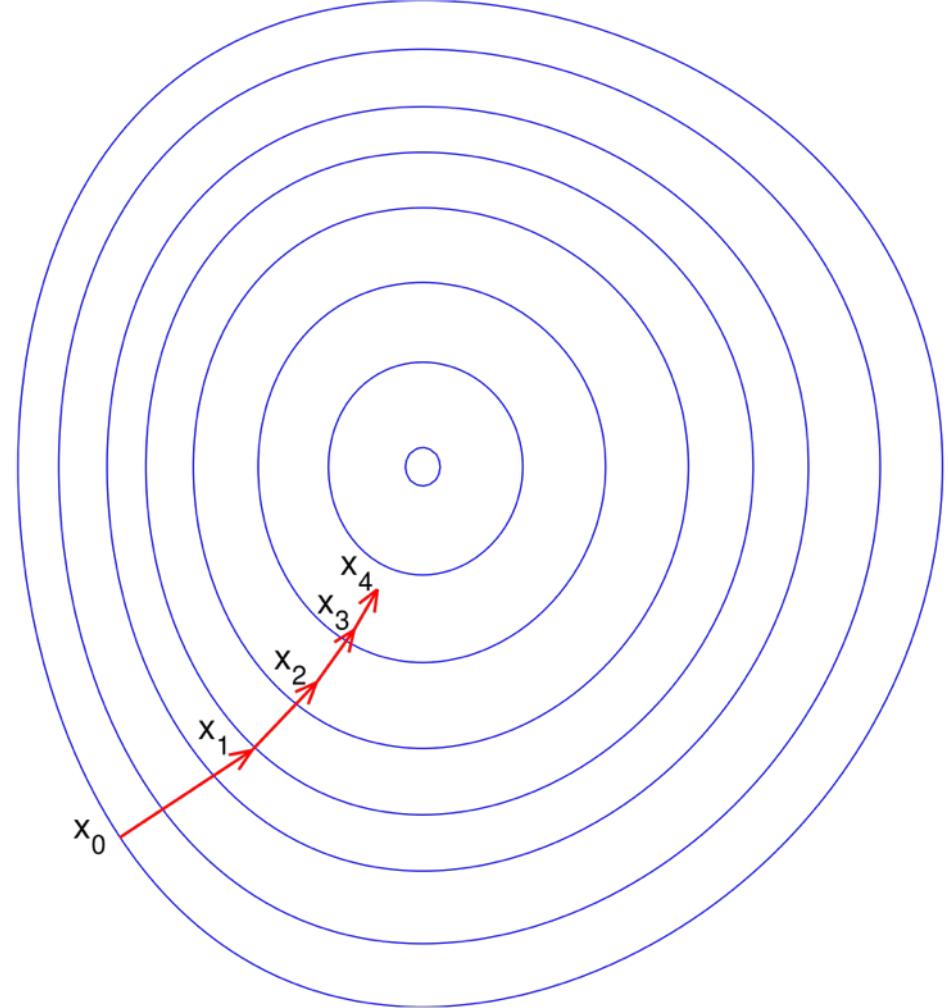
$$\left(\sum \frac{1}{\mathbf{x}^{(i)}} \text{--- } \mathbf{x}^{(i)} \text{ ---} \right) - 1$$

$O(nd)$ Distributed Storage

$O(nd)$
 ~~$O(nd^2)$~~ $O(d)$
Distributed Computation ~~$O(d^2)$~~ Local Storage

$O(d^3)$ Local Computation $O(d^2)$ Local Storage

Gradient Descent for Big n and Big d



Example: $n = 6$; 3 workers

workers:

$\begin{array}{|c|} \hline \mathbf{x}^{(1)} \\ \hline \mathbf{x}^{(5)} \\ \hline \end{array}$

$\begin{array}{|c|} \hline \mathbf{x}^{(3)} \\ \hline \mathbf{x}^{(4)} \\ \hline \end{array}$

$\begin{array}{|c|} \hline \mathbf{x}^{(2)} \\ \hline \mathbf{x}^{(6)} \\ \hline \end{array}$

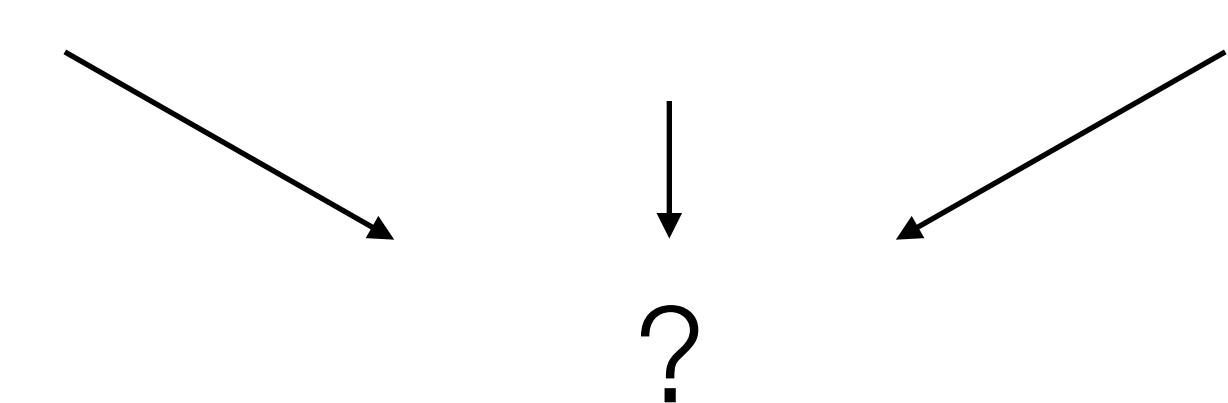
map:

?

?

?

reduce:



$O(nd)$ Distributed Storage

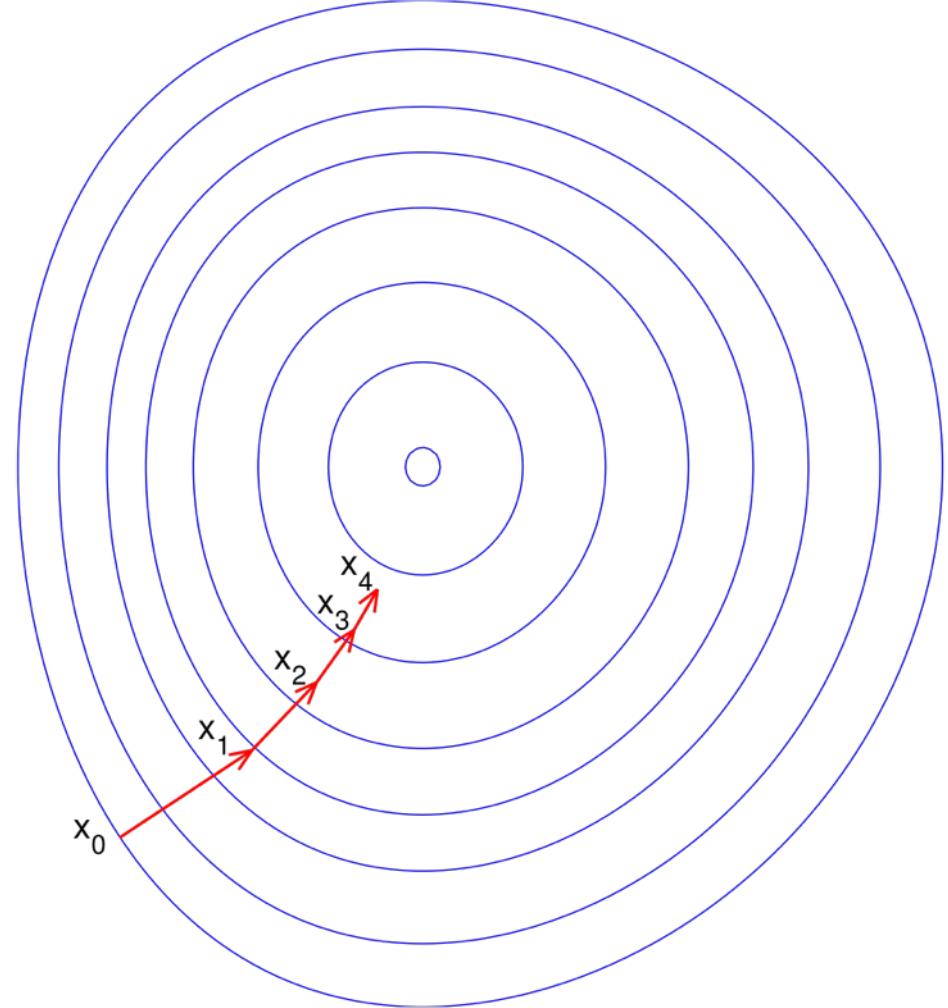
$O(nd)$
 $O(nd^2)$ Distributed Computation

$O(d)$
 $O(d^2)$ Local Storage

$O(d)$
 $O(d^3)$ Local Computation

$O(d)$
 $O(d^2)$ Local Storage

Gradient Descent for Big n and Big d



Example: $n = 6$; 3 workers

